Two AI Agents, One Stellarator

Building an engineering simulation pipeline from an open dataset – including the dead ends

Thomas Schmiegelt | hobby project `stellarator-design`

25 October 2026, updated 27 October 2026

Hobby project, driven by curiosity. Not affiliated with any research institute or company, not peer reviewed. All numbers are estimates from simplified models and may contain errors. The W7-X model is *inspired by* Wendelstein 7-X: coil paths come from public data [7]; casings, connectors, potting, supports and wall thicknesses are my own simplified assumptions and do not describe the real device. Code, computations and the draft text were produced with the help of AI agents; all figures are my own computations and renderings.

Introduction and motivation

I am not a fusion researcher. This is a private project, driven by curiosity: how does a stellarator work, and how far can one person get with open software and AI assistance on such a demanding part? Stellarator coils are a good example because design, manufacturing and strength meet here. Physics dictates the path, current and field create the forces, and the structure has to fit in between. Later the plasma vessel was added, with a cooling core of gyroid filaments.

A stellarator confines a fusion plasma with twisted 3D superconducting coils that carry large magnetic forces. In four days, a human and two AI agents built an end-to-end pipeline from a public coil dataset to a structural analysis of the coil casings: coil path → magnetic field → Lorentz force → casing geometry → stress. This report is mostly about *how* it was built: the division of labour between a local agent (Hermes) and a coding agent (Claude Code), a first model that was wrong while looking right, a fully parametric CAD model that broke the mesher, and a second, CAD-free code path that finally worked. For a W7-X-like system of 70 coils the result is 99 % of casing stresses below 43 MPa – under strongly simplified supports. A follow-up (Section 9) computes the flow through the gyroid cooling core of the 1 GW plasma vessel with FluidX3D and suggests improvements.

Contents

1 Setup **3**
- 1.1 Agents and responsibilities . . . 3
- 1.2 Tool chain . . . 3

2 Research phase **3**

3 First results – and why they were wrong **4**
- 3.1 Validating the solver . . . 4
- 3.2 The surrogate model (report v0.1) . . . 4

4 Restart: mesh directly from the data **6**

5 The parametric CAD branch **10**

6 The parallel branch: no CAD kernel **12**

7 Visualisation: from mesh to Blender **16**

8 Takeaways for working with coding agents **20**

9 Follow-up: cooling the plasma vessel **21**
- 9.1 Task and layout . . . 21
- 9.2 History . . . 21
- 9.3 An error in the lattice . . . 21
- 9.4 Is it enough? A rough heat estimate . . . 25
- 9.5 Suggested improvements . . . 26
- 9.6 Collaboration . . . 26

10 Follow-up: double filament, clearance and liquid lead **27**
- 10.1 Wall layout . . . 27
- 10.2 Water at 60 to 150 °C . . . 27
- 10.3 Clearance to the coils . . . 28
- 10.4 Liquid lead instead of water? . . . 28
- 10.5 Collaboration . . . 30

11 Coolant, strength and heat balance up to the turbine **31**
- 11.1 Variant f12: 5 mm inner skin . . . 31
- 11.2 Ranking by pairwise comparison . . . 31
- 11.3 Does the thin skin hold the pressure? . . . 32
- 11.4 Sourced reference values . . . 33
- 11.5 Decision: pressurised water as in a nuclear plant . . . 33
- 11.6 Derivation of the heat balance . . . 33
- 11.7 Do the pipes fit through the gaps between the coils? . . . 35
- 11.8 Collaboration . . . 36

1 Setup

Stellarator in five sentences. A fusion plasma (over 100 million degrees) is held by a magnetic field inside a ring-shaped vacuum vessel. To keep particles from drifting out, field lines must wind helically around the ring. A tokamak gets this twist from a large current in the plasma itself; a stellarator gets it entirely from external coils, which therefore have to be twisted in 3D. This makes steady-state operation easier but turns each coil into a demanding 3D part with tight tolerances. The reference device is Wendelstein 7-X (IPP Greifswald) with 50 non-planar and 20 planar superconducting coils [1].

1.1 Agents and responsibilities

Human	Goal and direction, reviews results *as images*, decides, supplies the key design ideas when stuck.
Hermes	Hermes Agent [15], run locally with open-weight models: literature research, project bookkeeping, writing reusable *skills*, long-running jobs overnight.
Claude Code	Coding agent [16]: simulation pipeline, debugging, post-processing, figures.

The agents never talk directly. They coordinate through files in the repository: an append-only findings log (`BEFUNDE.md`; every entry has observation, measurement, location, status), a handover file for decisions and running jobs (`ERINNERUNGEN.md`), and skills – short Markdown playbooks an agent loads before a task. Long simulations are started detached (`setsid nohup`) so they survive a session ending or a context compaction. The very first human intervention was a stop: Hermes started building tooling right away and was told to do the research first.

1.2 Tool chain

Physics data	simsopt [7] (coil Fourier coefficients), DESC [8] (plasma shape optimisation)
Geometry	Python + gmsh/OpenCASCADE [9] (direct meshing), FreeCAD 1.1 headless [11] (parametric CAD)
Solver	ElmerFEM 9.0 [10]: magnetostatics and linear elasticity in one coupled run
Own code	Biot-Savart (numba), FFT field solver, field-line tracer, collision checks
Visualisation	PyVista/VTK [12], ParaView [13], Blender 4.0 [14] via glTF
Hardware	one workstation with one GPU (reserved for the local agent)

2 Research phase

Hermes collected the basics: W7-X has a major radius of 5.5 m, five field periods and 2.5 T on axis [1]. Each coil is wound from a cable-in-conduit conductor: 243 NbTi strands of 0.57 mm, each with 144 filaments in copper, inside a co-extruded 16 mm × 16 mm aluminium jacket with 37 % void for helium cooling [2, 3]. From a set of 33 PDFs supplied by the user, four were relevant; the most useful one sizes coil structures for EU-DEMO with a 660 MPa stress limit and wedged inner legs [4]. A rate-limited arXiv search added work on HTS strain sensitivity [6] and on manufacturing tolerances (W7-X: 2 mm) [5].

Lesson – verify citations. The bibliography an agent produced during research contained entries without authors and with wrong venues or years. For this report every reference was checked individually against full text or publisher metadata (Crossref, arXiv API). Treat agent-written references as unverified until checked.

3 First results – and why they were wrong

3.1 Validating the solver

Before simulating coils, ElmerFEM was tested on a case with an exact answer: an infinite straight wire, $B = \mu_0 I / 2\pi r$. The first runs returned zero everywhere – silently. Root causes: a mesh format the converter misread, and wrong keywords in the solver input file, which Elmer ignores without an error. Starting from Elmer's own regression test instead of hand-written input fixed it: 0.09 % error inside the wire, 2 % outside (Fig. 4).

3.2 The surrogate model (report v0.1)

Hermes believed the real coil paths were unavailable (the URLs it tried returned 404) and generated a parametric surrogate: ellipses and bands, current calibrated to 2.5 T. Compared with W7-X the ampere-turns per coil matched to 0.90×. A report was written.

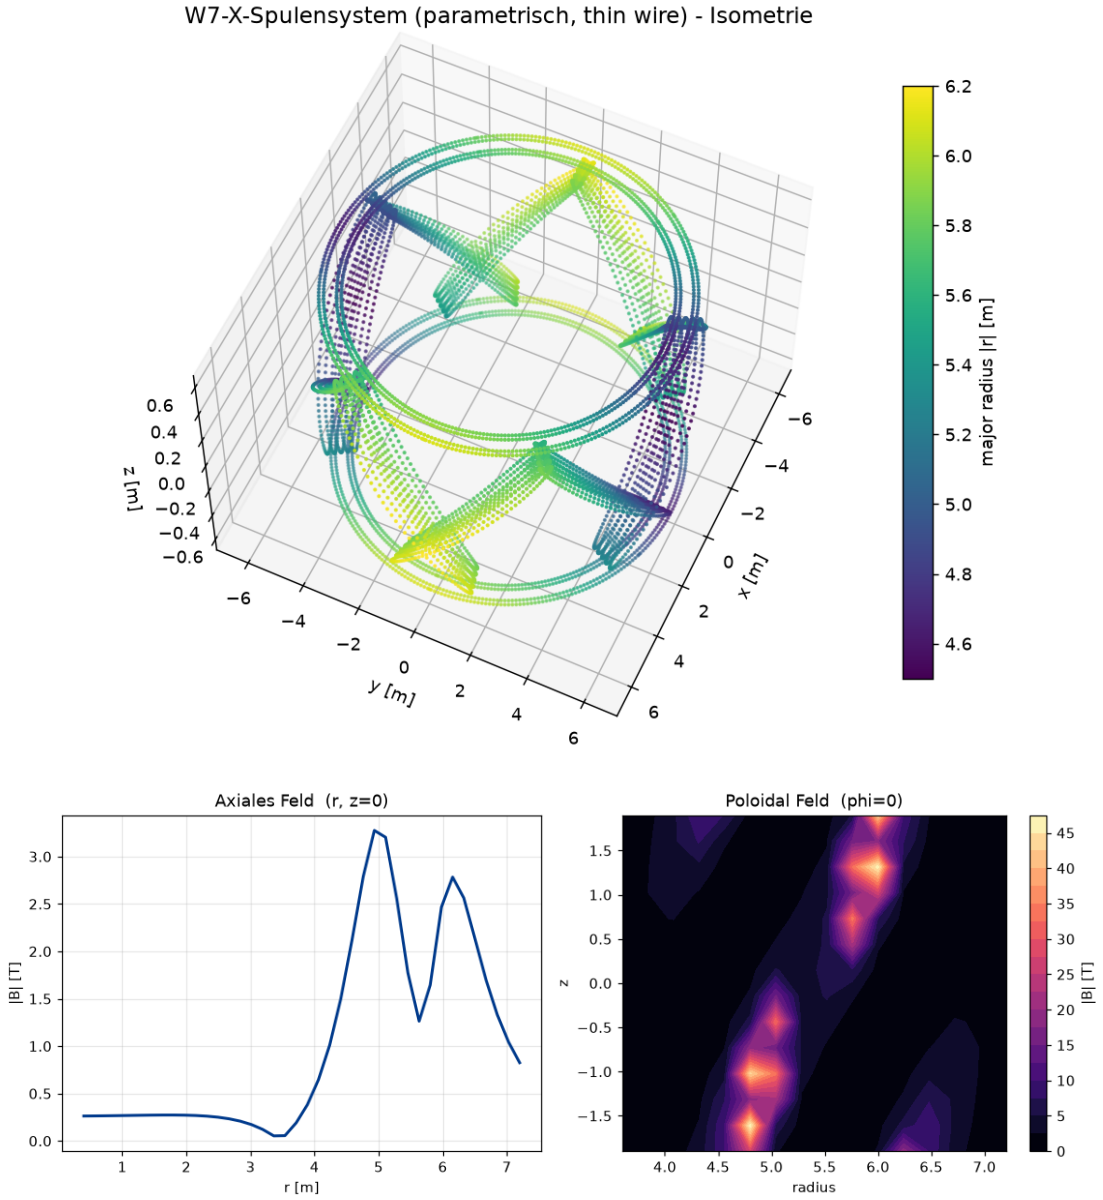


Figure 1: Report v0.1: the surrogate coil set and its field. Plausible numbers – but not a stellarator.

A topological check by the coding agent – does each coil actually encircle the plasma axis? – failed for 20 of 50 coils; the "planar" coils were horizontal bands above and below the plasma. The good match was a coincidence. And the data had been available all along: simsopt ships the Fourier coefficients of the seven

W7-X base coils in its package data [7]. Two more bugs surfaced later and changed published numbers: a sign error in the axis formula (calibration current 12.18 kA instead of 13.49 kA per turn), and Elmer choosing the current direction itself, so 25 mirrored coils ran backwards (1.38 T instead of 2.50 T on axis). The latter is now caught by an automated check.

Lesson – plausible is not correct. Agents are good at producing numbers that look right. Every geometry gets an independent invariant check (winding number around the axis, clearances), and every field or force is computed by a second, independent method before it goes into a report.

4 Restart: mesh directly from the data

With the real coil paths (7 base coils, 70 by symmetry) the expected physics appeared immediately: bean-shaped flux surfaces and the characteristic island chain (Fig. 2). A custom FFT voxel solver matched exact Biot-Savart to 1.3 % but destroyed the flux surfaces during field-line tracing, because a Cartesian grid breaks the device's symmetry – so field lines are always traced with Biot-Savart.

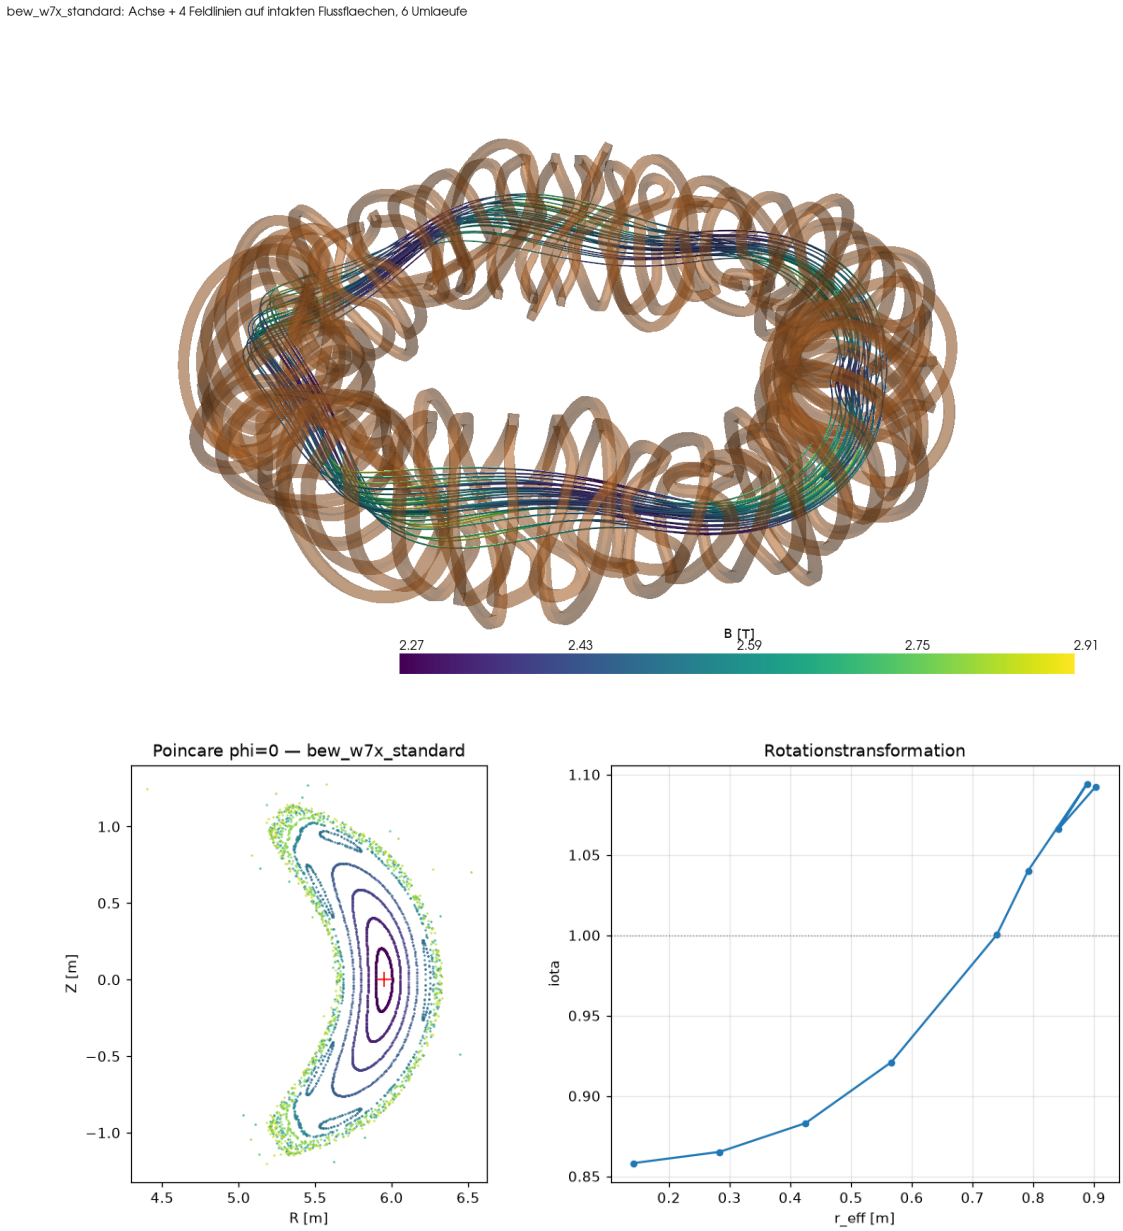


Figure 2: Field lines on flux surfaces (top), Poincaré section and rotational transform (bottom).

The plan had been to model coils in CAD. Generating the mesh directly was faster and more robust: a rectangle at every path point, two ruled half-surfaces per coil, meshed by gmsh. Sweeps and smooth lofts failed with self-intersecting surfaces. A clearance check on centre lines replaced boolean intersection tests (1.5 s instead of more than 45 min). Elmer with all 70 coils (3.3 M tetrahedra, 680 s, 10 GB) matched Biot-Savart on the axis to a median of 0.24 %.

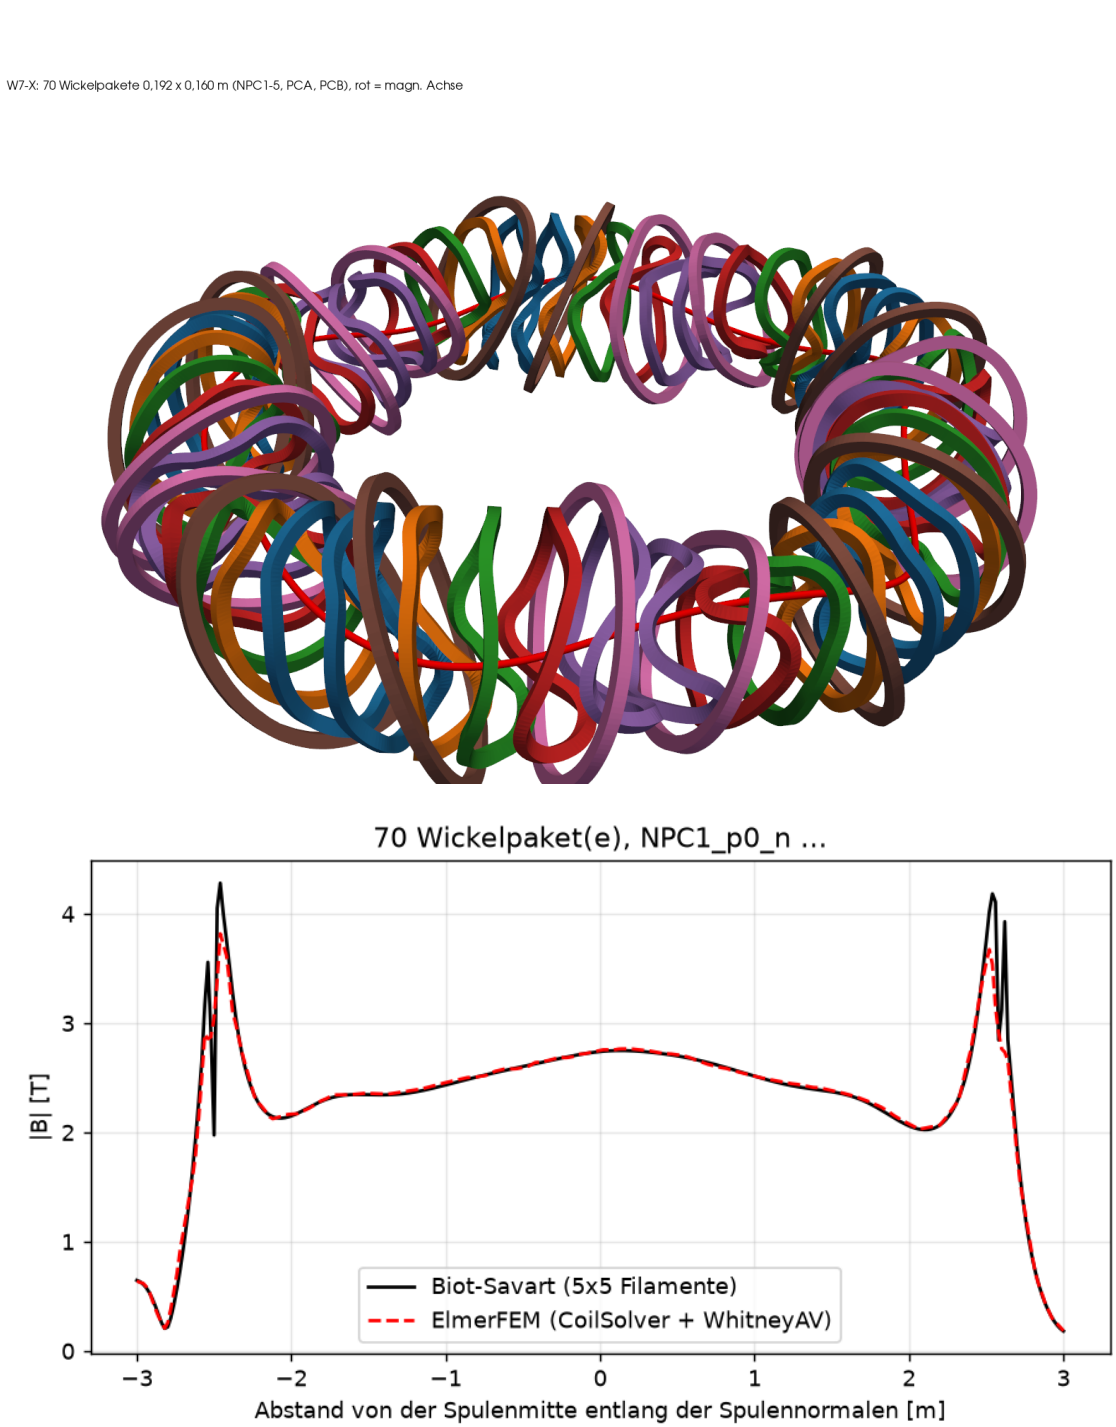


Figure 3: 70 winding packs meshed straight from the path data (top); FEM vs. Biot-Savart along a coil normal (bottom).

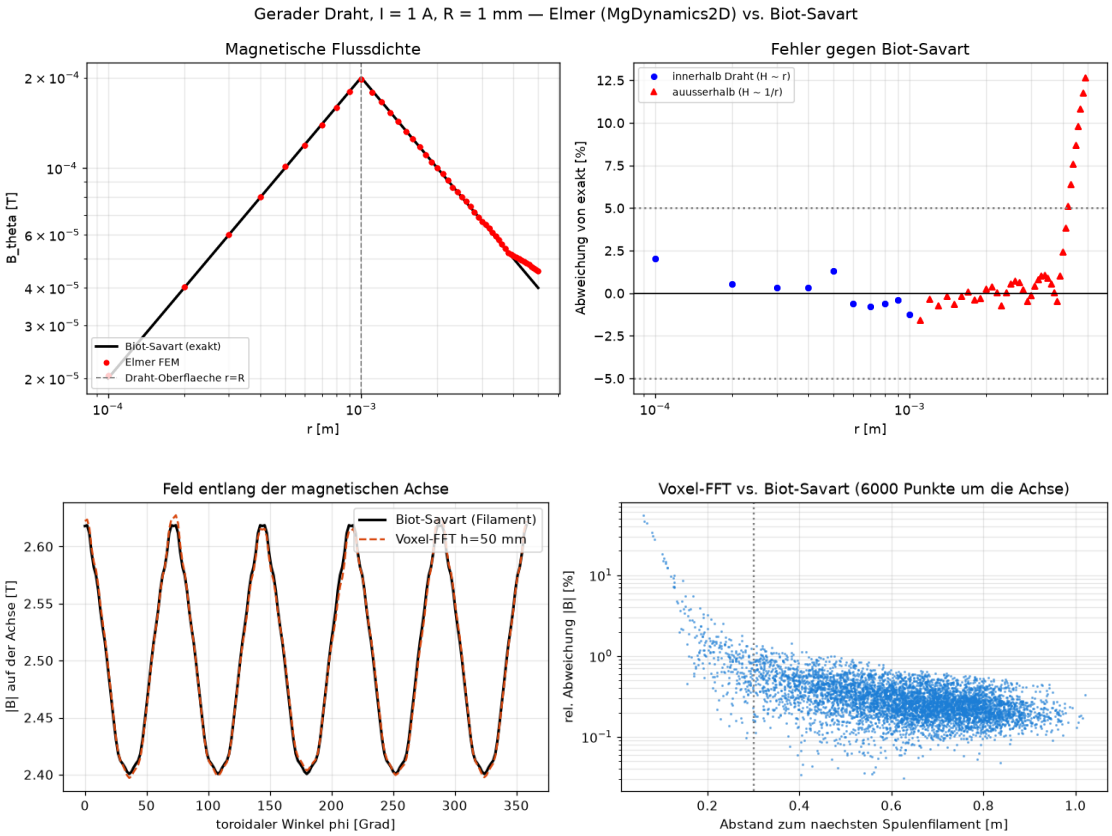


Figure 4: Validation: FEM on a straight wire (top), FFT voxel solver vs. Biot-Savart (bottom).

A side study ran the pipeline backwards from a target of 1 GW net electric power: plasma size and field → own plasma shape (DESC) → coils (simsopt). With high-temperature superconductors the device ends up at about 10 m major radius with forces of roughly 1 GN per coil – about a thousand times W7-X. From here on, the support structure is the real problem.

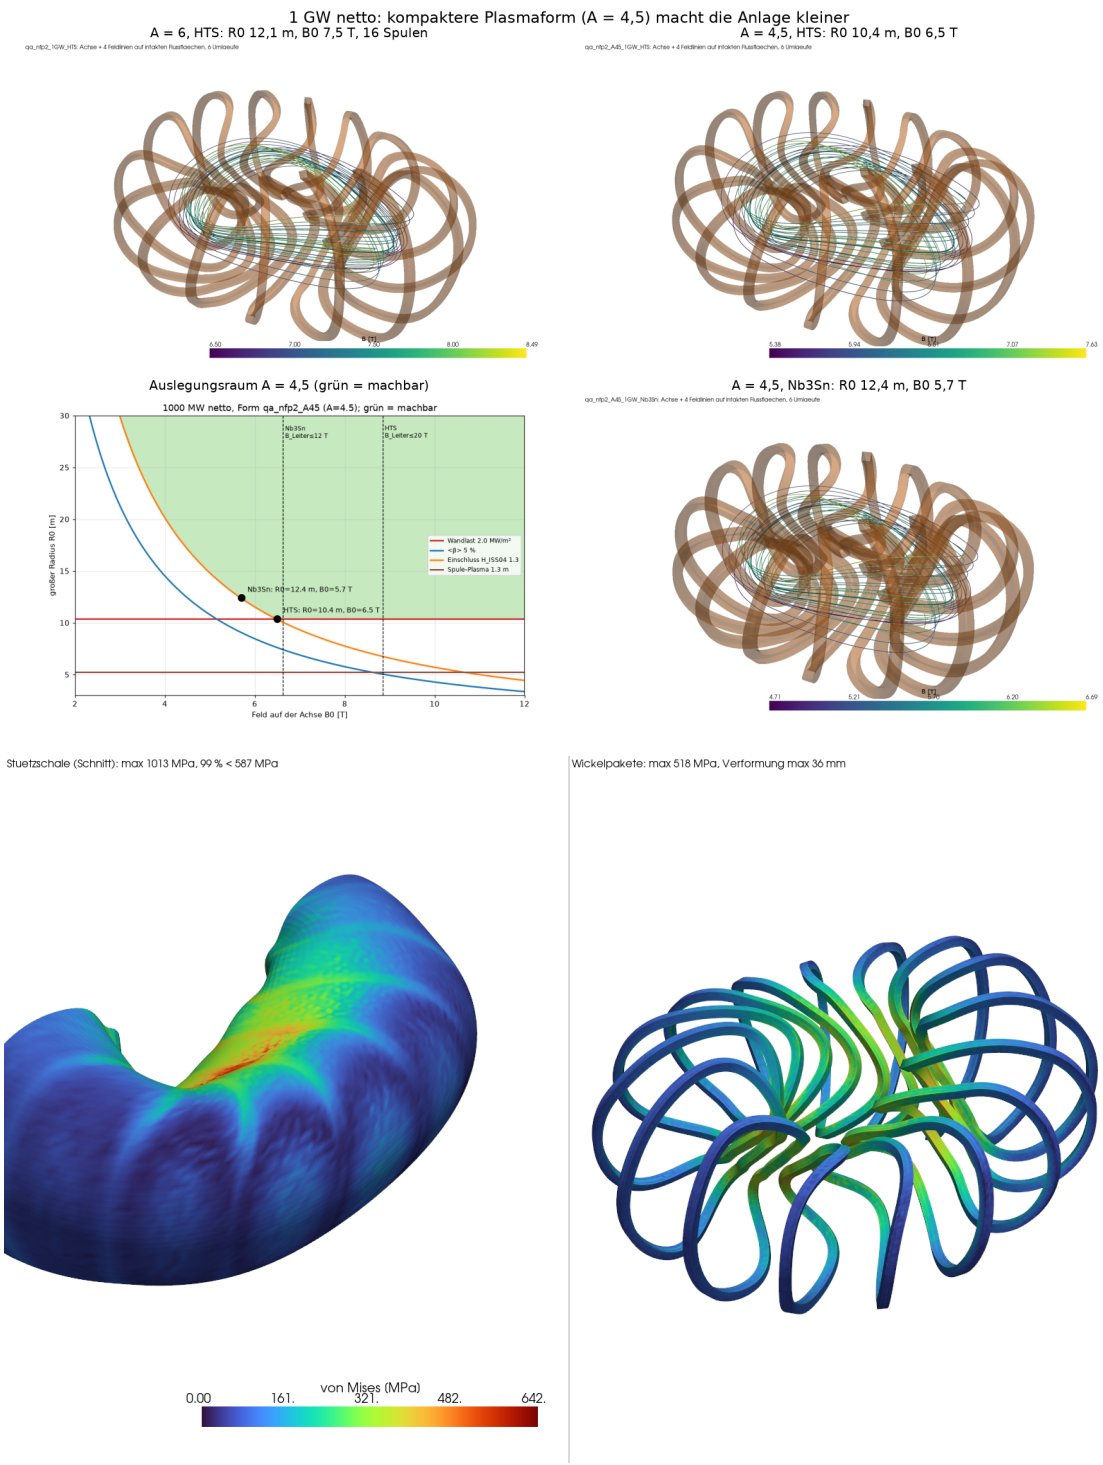


Figure 5: 1 GW sizing study (top) and a first support shell with coupled field → force → stress analysis (bottom).

5 The parametric CAD branch

A continuous support shell could not be constructed cleanly. The human proposed a technique from commercial CAD (Creo's *Variable Section Sweep*): per coil a closed guide curve as the casing's inner contour, sketches normal to it, a loft through them; all walls grow outwards. The agent implemented it in headless FreeCAD as a fully parametric model: spreadsheet parameters, sketch dimensions as expressions, symmetric coils as clones, connectors between neighbours.

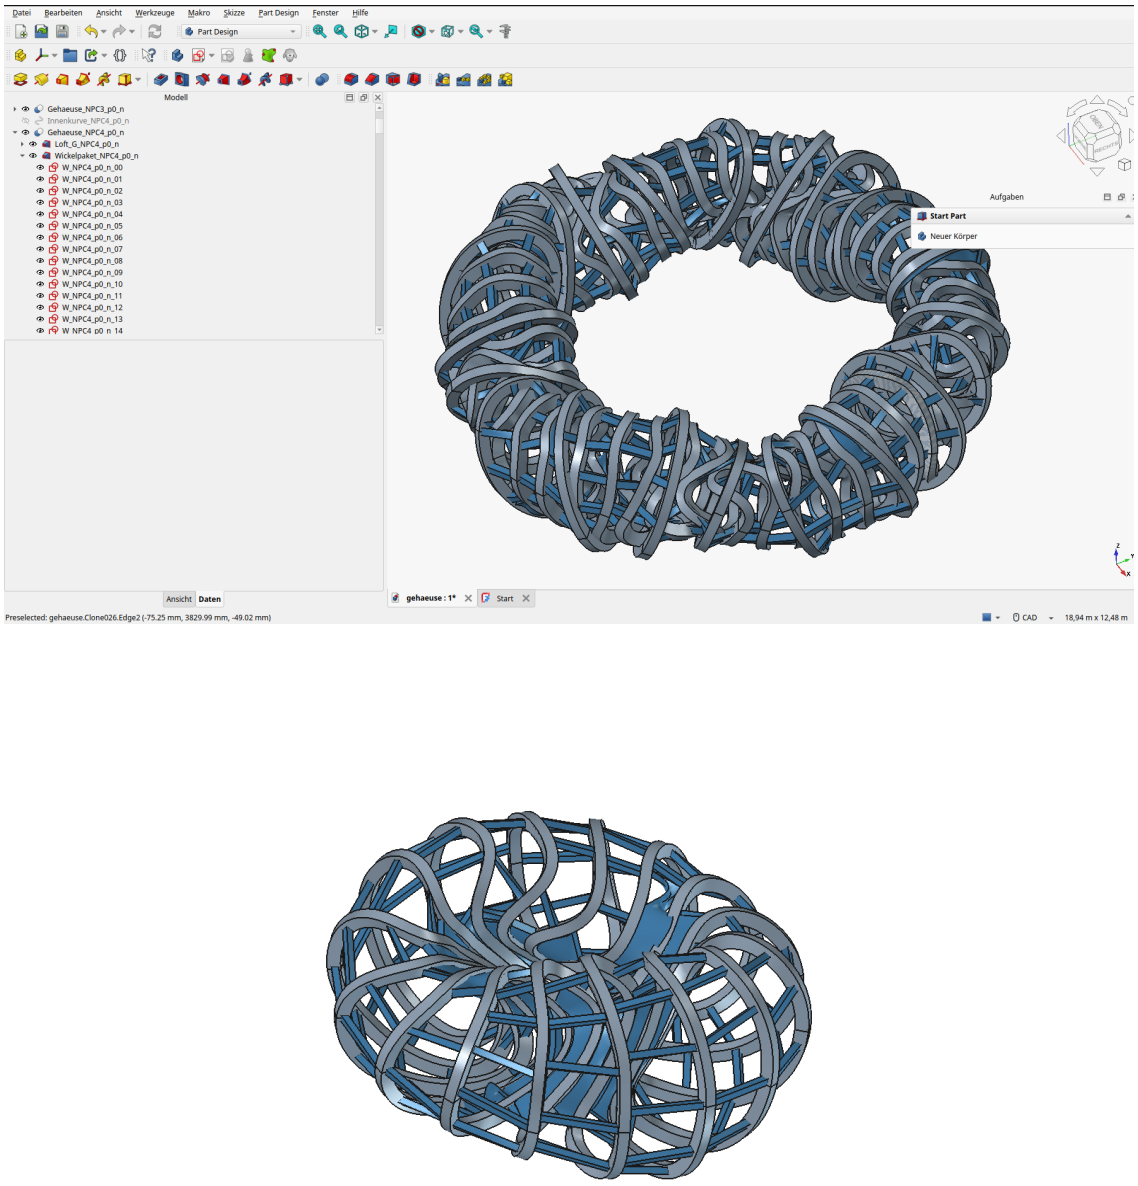


Figure 6: Parametric FreeCAD model: W7-X-like casings with connectors (top), compact 1 GW design (bottom).

What went wrong:

1. **Estimated, not designed, walls.** The first coupled run worked but gave stresses up to eight times the allowable value.
2. **Wedges break the mesher.** Inner wedges (after [4]) produced a valid CAD solid, but gmsh failed with "overlapping facets": several boolean cuts on the same winding pack created coincident faces. Six

attempts later (tolerances, fusing in FreeCAD, deliberate overlap; the last one hung gmsh for over an hour) the wedge was disabled.

3. **Collisions.** A loft bulges between sketch stations and corners of tilted neighbours overlap. An incremental assembly check (add coils one by one, count collisions, render a video) found about 100.
4. **Fix in CAD (human idea): subtract the neighbours.** Each casing minus its neighbours' envelopes plus a gap: 0 collisions in all 70 steps. Meshing the whole assembly stayed fragile.

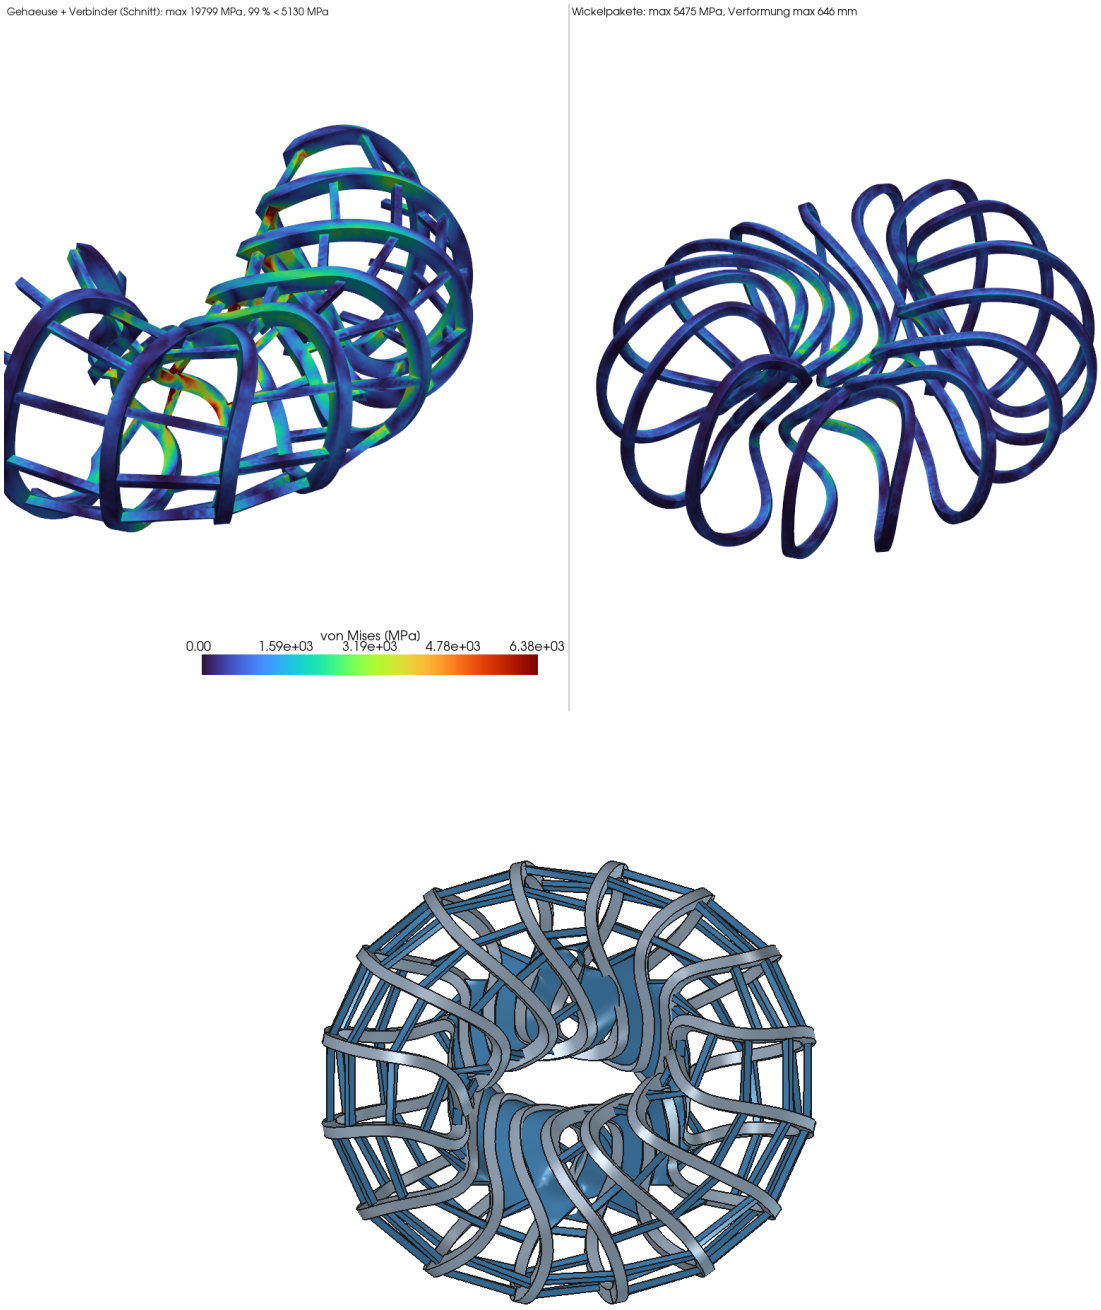


Figure 7: First stress run of the CAD casings (top); casings with inner wedges – valid in CAD, not meshable (bottom).

Lesson – valid is not meshable. The more boolean operations meet on the same faces, the more sliver faces appear, and the mesher fails far away from where the problem was created. Keep the analysis model separate from the design model.

6 The parallel branch: no CAD kernel

The human's suggestion: *"keep this morning's working approach for the analysis and run the CAD path in parallel as an add-on."* Each casing became two rectangular tubes on the exact coil path (outer contour and winding-pack cavity) as separate surface meshes – no boolean operations at all. Wall thicknesses follow from load and available space; a disentangling step tests surface points of foreign casings (including intermediate rings between stations) against the local cross-section and shrinks the cheapest wall: $4980 \to 40 \to 0$ contact points.

Model	steel casings ($E = 200\,\mathrm{GPa}$), homogenised winding packs ($E = 100\,\mathrm{GPa}$), linear elastic
Load	Lorentz force $\vec{J} \times \vec{B}$ from the magnetostatic solution in the same run
Supports	one clamped patch per casing (simplification; no connectors, no support ring)
Mesh / run time	3.02 M tetrahedra in 78 s; solve 55 min, about 14 GB RAM
Force per coil	FEM 2.42 MN vs. Biot-Savart 2.39 MN (+1 %)
Casing von Mises	mean 11 MPa, 99 % below 43 MPa, max 110 MPa; deformation below 1 mm

Because of the simplified supports and walls as thin as 3 mm in places this is *not* a strength proof – but it is a reproducible end-to-end pipeline.

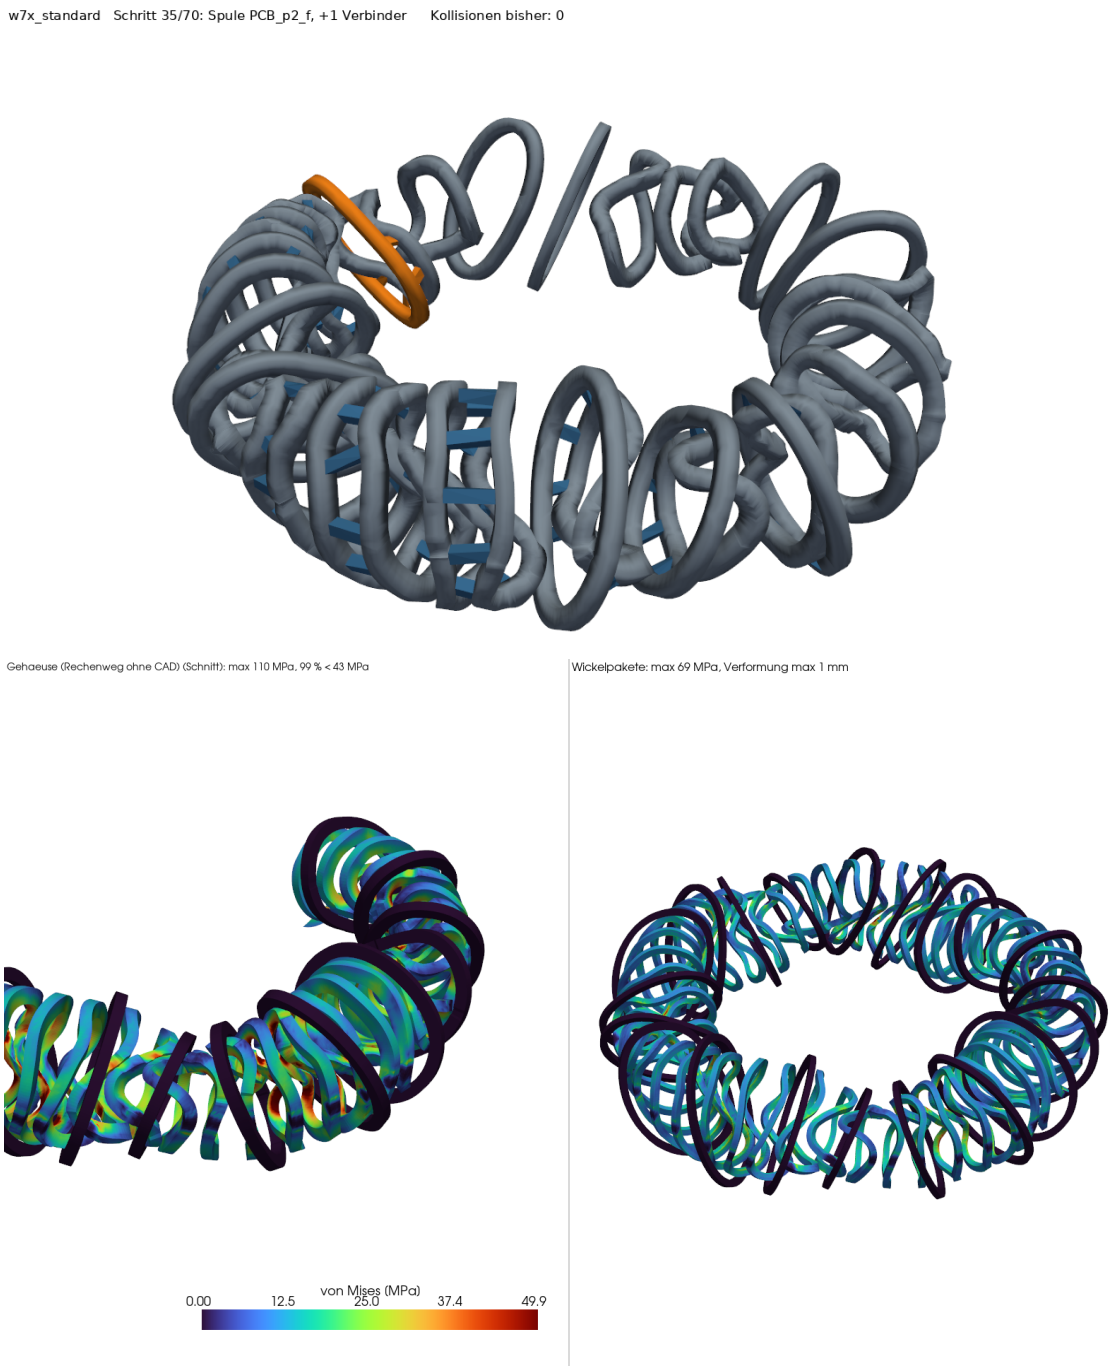


Figure 8: Incremental collision check in the CAD branch, step 35/70 (top); von Mises stress from the CAD-free branch (bottom).

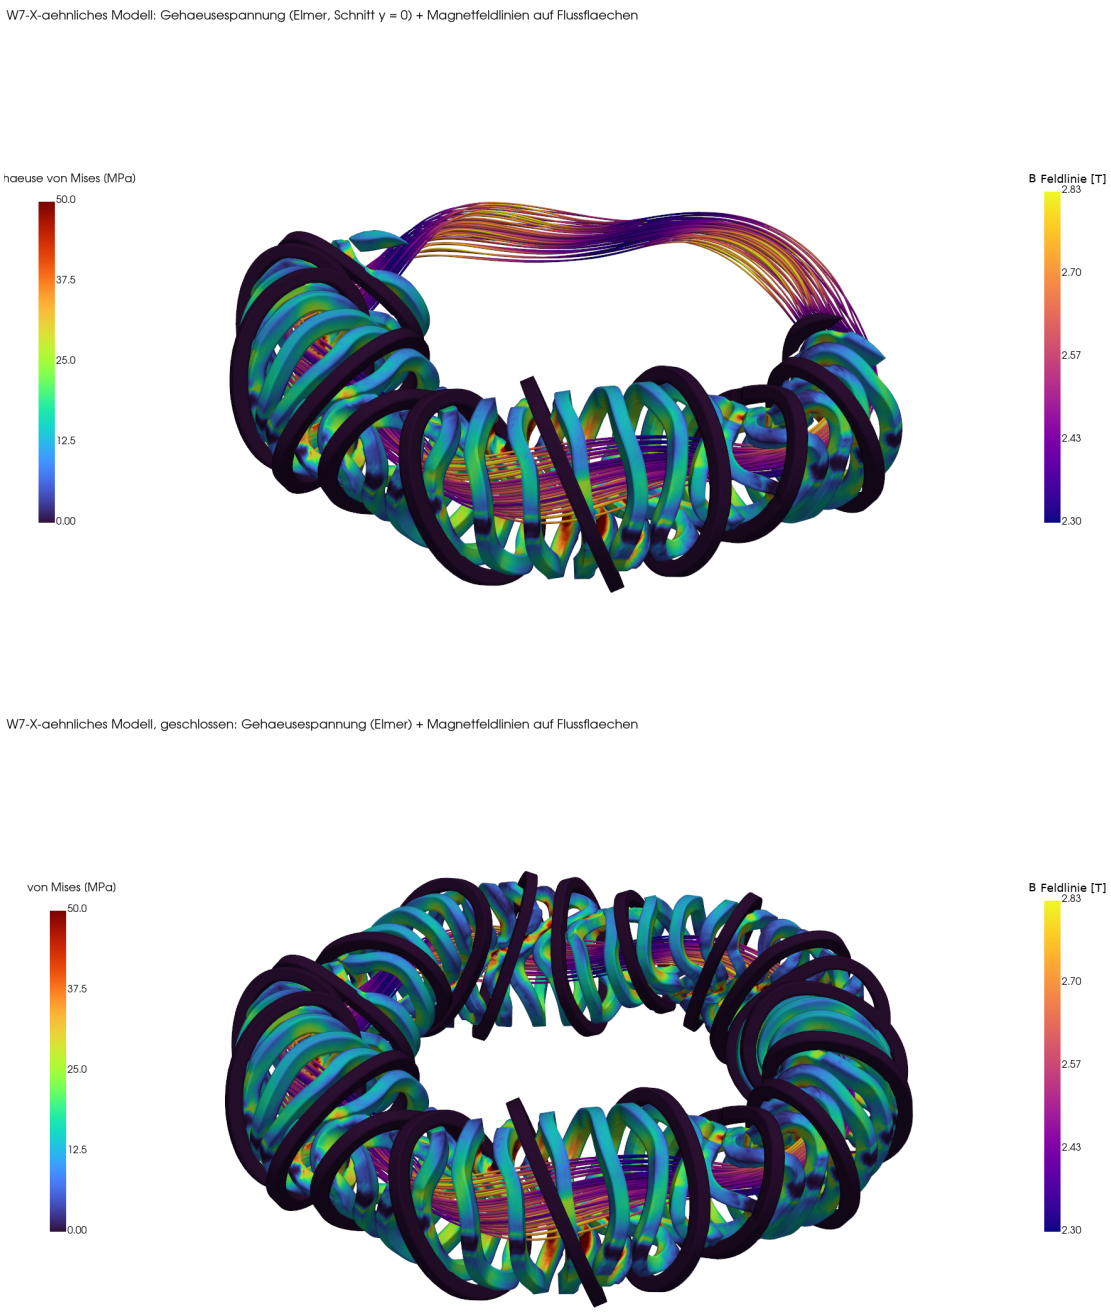


Figure 9: Casing stress with magnetic field lines on flux surfaces, cut open and closed.

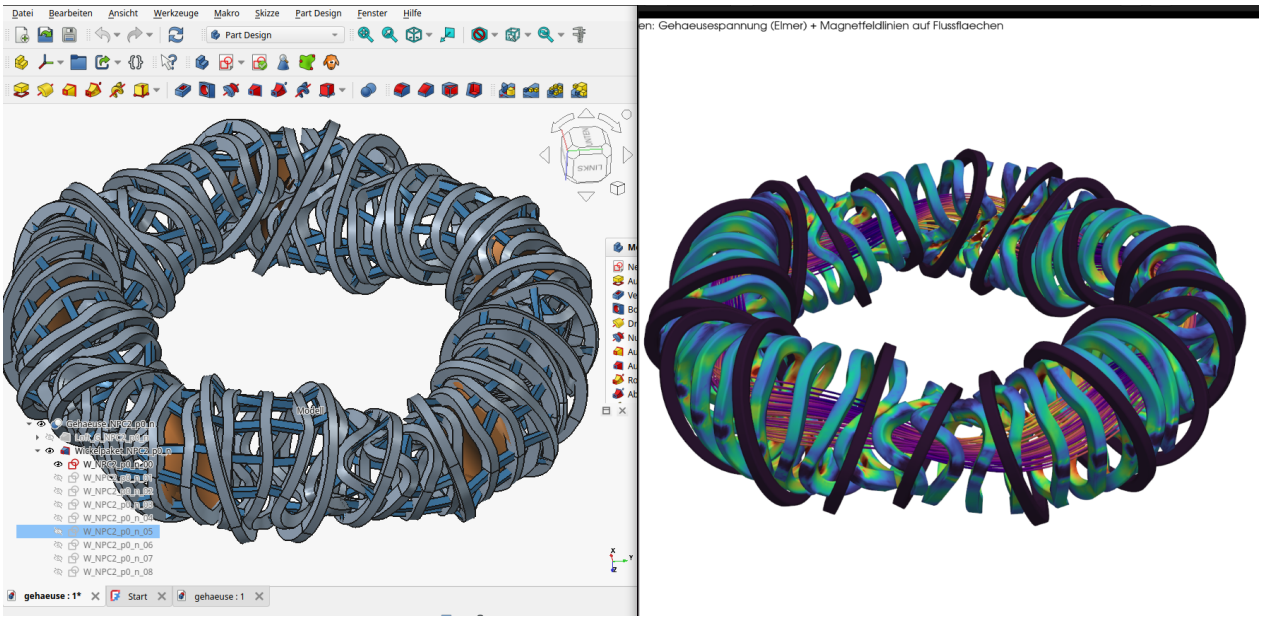


Figure 10: Both branches side by side: parametric CAD (left), stress analysis from the CAD-free pipeline (right).

7 Visualisation: from mesh to Blender

A 3D cut-away model shows each coil's layers: cable (NbTi strands in copper, helium), aluminium jacket, turn insulation, ground insulation, potting (an assumption), steel casing; $12 \times 9 = 108$ conductors per coil. A separate conductor model answers the question "where does the current flow?": almost entirely in the superconducting NbTi; copper only takes over briefly during a quench; jacket and helium carry none.

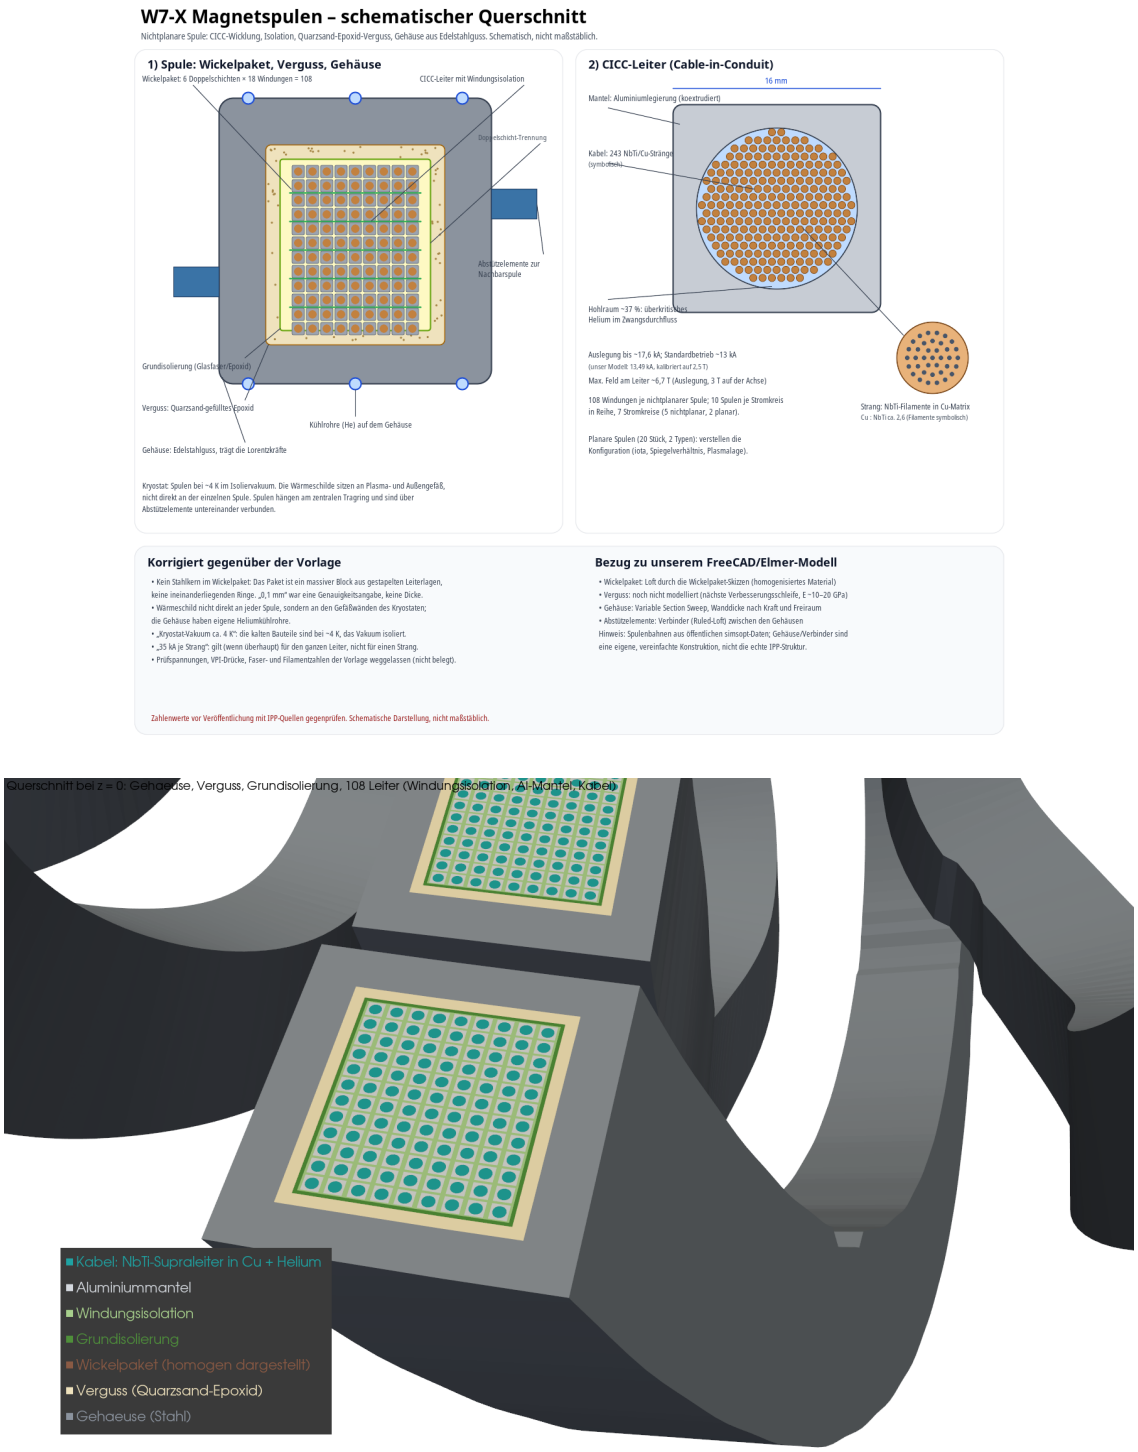


Figure 11: Coil layers (schematic, top) and in the 3D model (bottom).

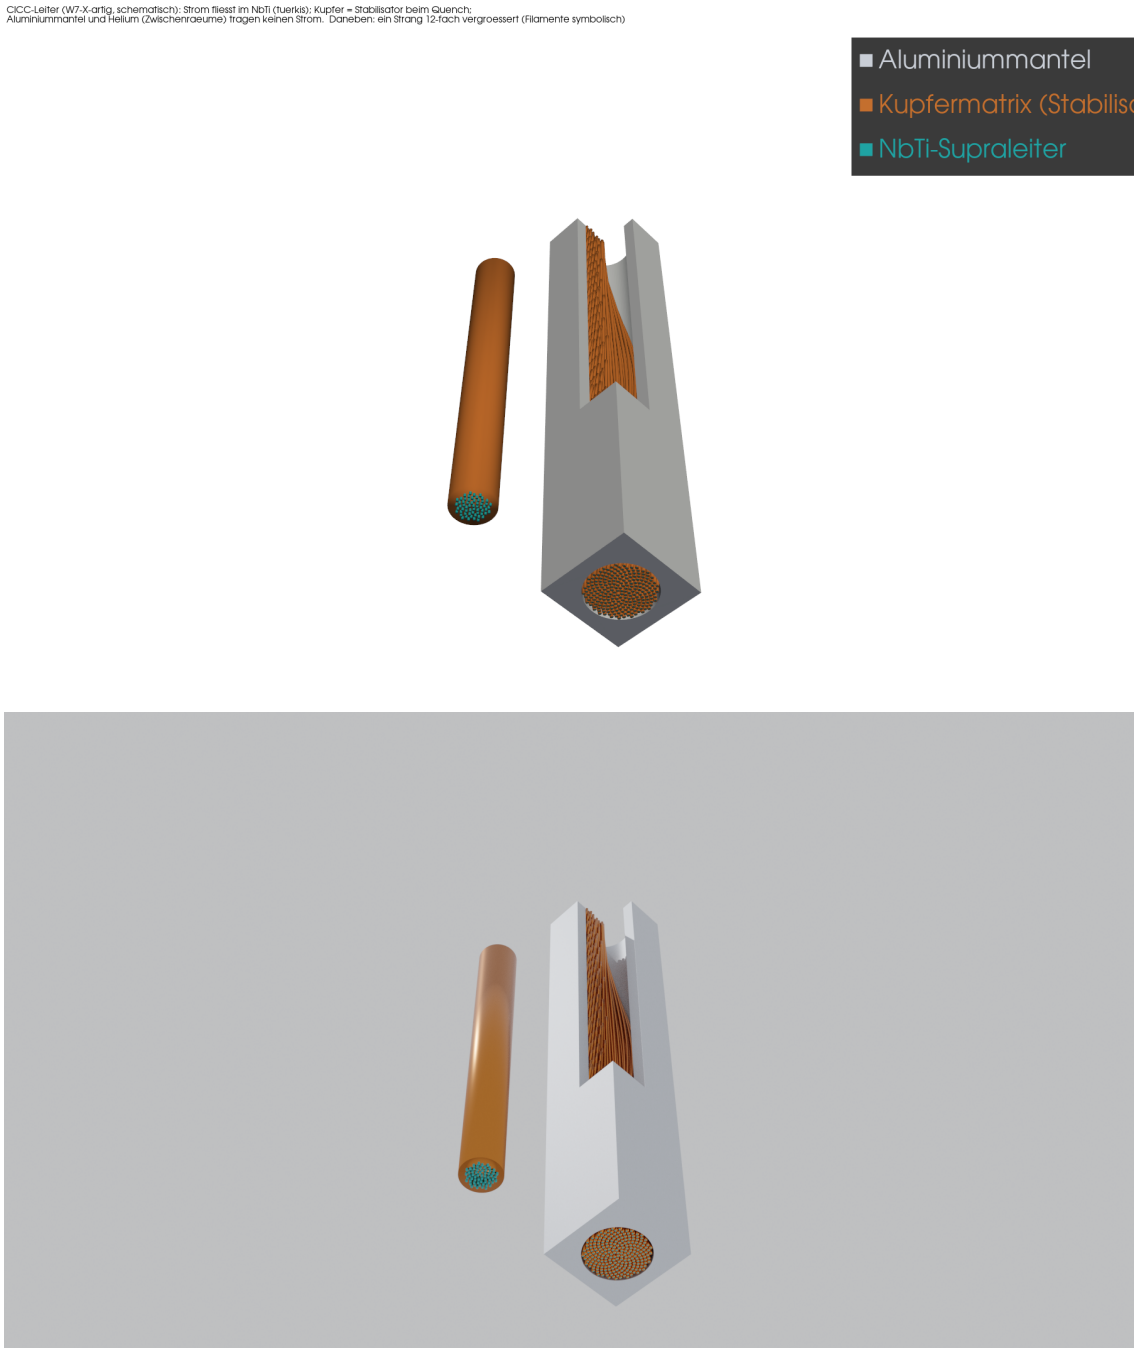


Figure 12: Conductor: opened jacket, twisted strands, one strand magnified with NbTi filaments (count symbolic). PyVista (top), Blender Cycles (bottom).

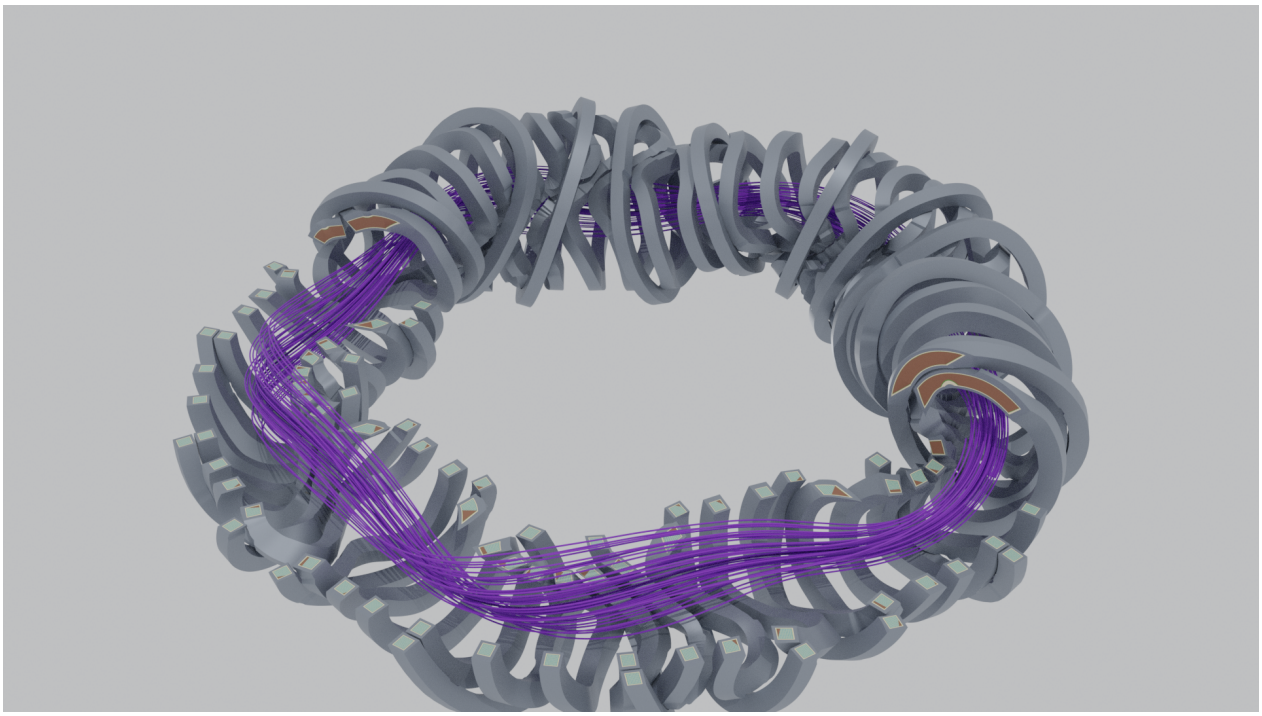


Figure 13: The full coil system with field lines, rendered in Blender.

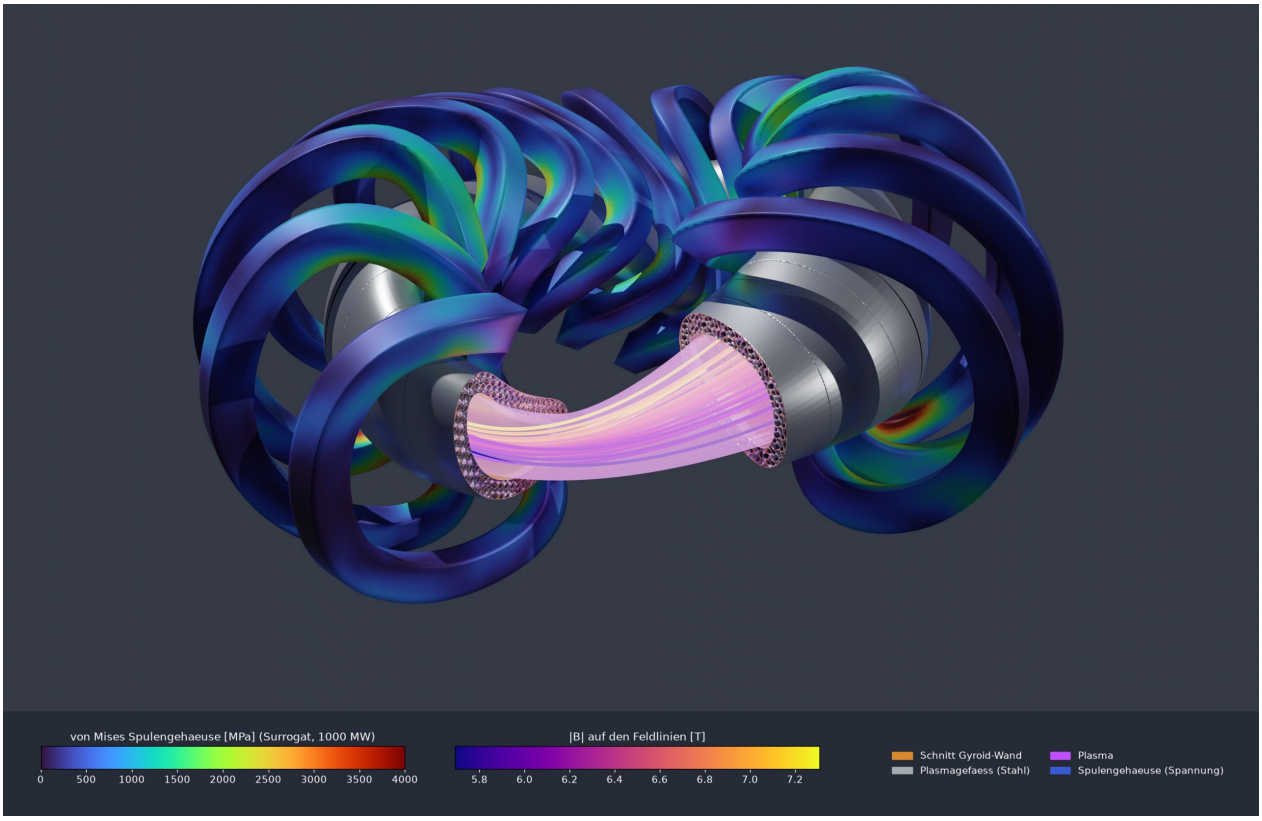


Figure 14: Later step: the 1 GW design with casing stress from the surrogate model (1000 MW), field lines, and a cut-open plasma vessel whose wall has a gyroid cooling core built with PicoGK. A rendering, not a cooling design.

Getting PyVista's glTF export into Blender needed three fixes, each found by measuring instead of guessing: colours arrived as sRGB values in a linear slot (washed out; fixed by linearising), the scene was rotated by 90° about the vertical axis (fixed at the root node), and the distribution's Blender build lacks the

denoiser (render fails; denoising disabled). Rendering runs on the CPU because the GPU is reserved for the local agent.

8 Takeaways for working with coding agents

- **Verification is the job.** The most expensive error looked plausible. Independent checks (invariants, second methods, regression tests from upstream projects) caught every major bug.
- **Silent failure is the default in scientific tools.** Wrong keywords ignored, current directions chosen by the solver, meshers hanging for an hour. Every pipeline stage now asserts on its output.
- **Prefer the boring path.** Direct mesh generation from data beat a CAD kernel with booleans. Keeping a working path alive while exploring the elegant one saved the project.
- **Humans supplied the decisive ideas** (research first, variable section sweep, neighbour subtraction, the parallel branch); agents supplied speed in implementing, testing and documenting them.
- **Files are the shared memory.** A findings log, a handover file and skills let two agents work on one project across sessions and context compactions.
- **References need checking** like any other agent output.

Open items

Potting layer between winding pack and casing; connectors and support ring instead of clamped patches; minimum wall thickness; wedge and casing in one body; wall sizing for the 1 GW design; HTS strain check; tolerance study (± 2 mm).

9 Follow-up: cooling the plasma vessel

Status 27 Oct 2026. Flow computed with cold water; heat only estimated roughly. The numbers show trends; they are not a design.

9.1 Task and layout

The plasma vessel wall of the 1 GW design is 0.9 m thick and has a core of gyroid filaments built with PicoGK [25]. A 10° segment (85°–95°) was studied. The assumed heat flux is 0.5 MW/m^2 on 26.7 m^2 of plasma-facing surface, i.e. 13.3 MW. With a 10 K temperature rise the segment needs 319 l/s of water.

The layout grew step by step from the user's instructions:

- Water flows *through the filaments*, not through holes in a perforated plate.
- The inlet is on the inboard side of the torus, the outlet outboard.
- The core sits between two inclined planes. In front of and behind it, the distribution and collection chambers taper off (Z-manifold). This shape came from a sketch the user drew on a screenshot.
- The gyroid is graded: a 130 mm layer with 100 mm cells on the plasma side, behind it a dense, coarse gyroid as load-bearing structure.

The flow was computed with FluidX3D [17], a lattice Boltzmann solver for GPUs [18, 19, 20, 21, 22, 23, 24]. The lattice has 8 mm spacing and 259 million cells. Each run covers 40 s of physical time with the LES turbulence model and takes about 15 min. FluidX3D has its own licence: free for non-commercial use only, source-available, not an open-source licence. The two modified files (`setup.cpp`, `defines.hpp`) are published as an "altered version" with the original licence in the repository folder `fluidx3d/`.

9.2 History

Variant	Change	Result
a, b	chamber with perforated plate; holes missed the plate (spotted by the user)	unstable, jets
f1	water through the filaments, open chambers	spread 141 %
f2	finer gyroid towards the plasma	spread 40 %
f4, f5	Z-manifold from the sketch, denser outer layer	spread 20 %, 7.9 kPa
f6–f9	thinner front layer, wall thickness by radius	dips at ±52°, cause searched
f10	lattice cells taken directly from PicoGK	dips gone, spread 36 %
f11	front wall 4.9 mm to 5.8 mm by radius	**spread 8 %**, 22 kPa, 288 l/s

Table 1: Variants of the cooling segment. Spread = standard deviation of the flow over 24 circumferential sectors divided by the mean.

9.3 An error in the lattice

In f6 to f9, two bands at ±52° of the circumference received almost no water. It looked like a design flaw. Three hypotheses were tested and rejected in turn: layer thickness, orientation of the layer relative to the gyroid, and inaccurate offset surfaces in PicoGK. Two intermediate findings were our own measurement errors.

The actual cause was the conversion of the water mesh into lattice cells. FluidX3D decides for each cell with a ray test whether it is inside. On nearly horizontal faces of the very fine gyroid mesh this failed: at $z = 1.906$ m a whole lattice layer became steel and blocked the band. It was found only by comparing the lattice with the mesh layer by layer. Since f10, PicoGK writes its voxel occupancy to a file that FluidX3D reads directly; it matches the mesh 100 %. The numbers of f6 to f9 are therefore distorted.

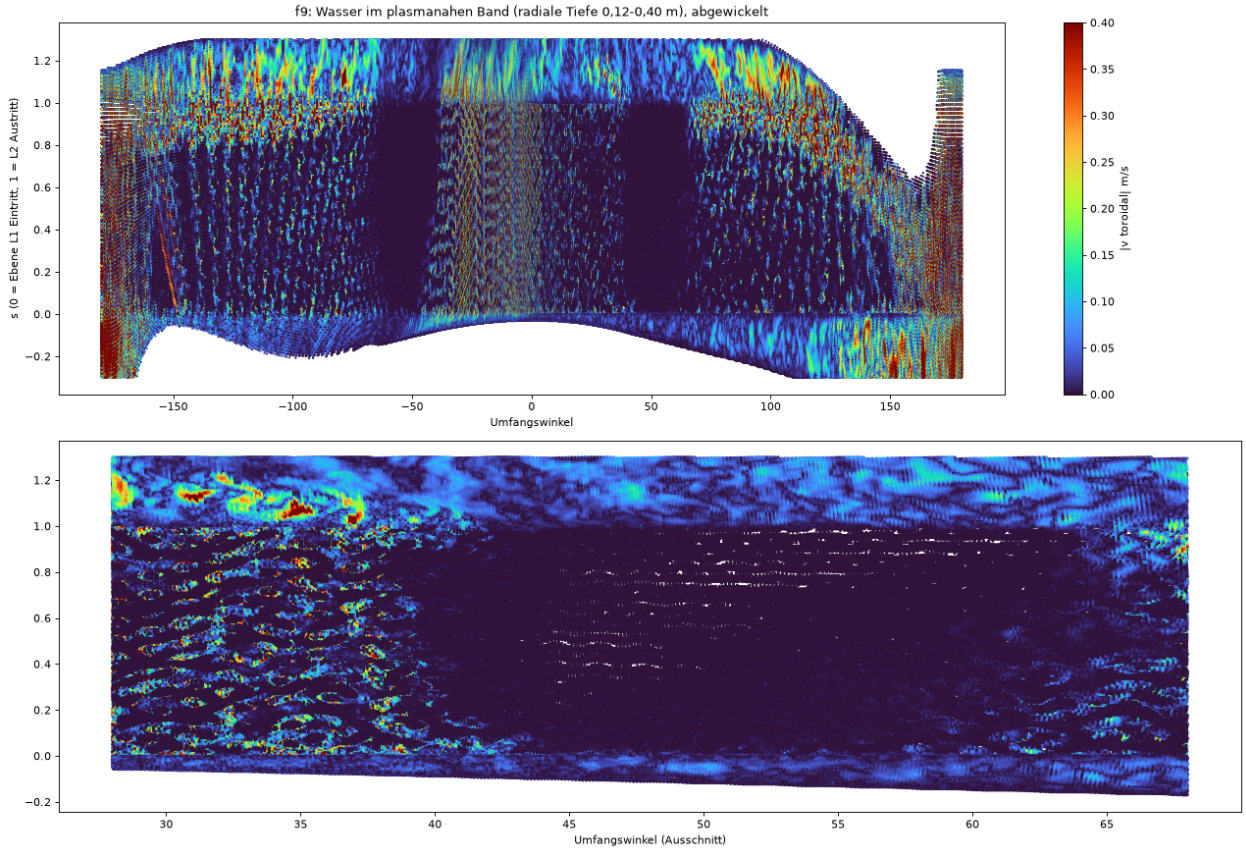


Figure 15: Lattice error (f9): water in the plasma-side layer, unrolled over the circumference (horizontal) and the path from inlet to outlet (vertical), colour = toroidal speed. The dark blocks at about $\pm 50°$ carry no flow because one lattice layer there was wrongly set to steel. Zoom below. Labels in German.

After that, the distribution could be balanced in one step. The wall of the front filaments now varies with radius: 5.8 mm inboard, 4.9 mm top and bottom, 5.2 mm outboard. The spread dropped from 36 % to 8 %.

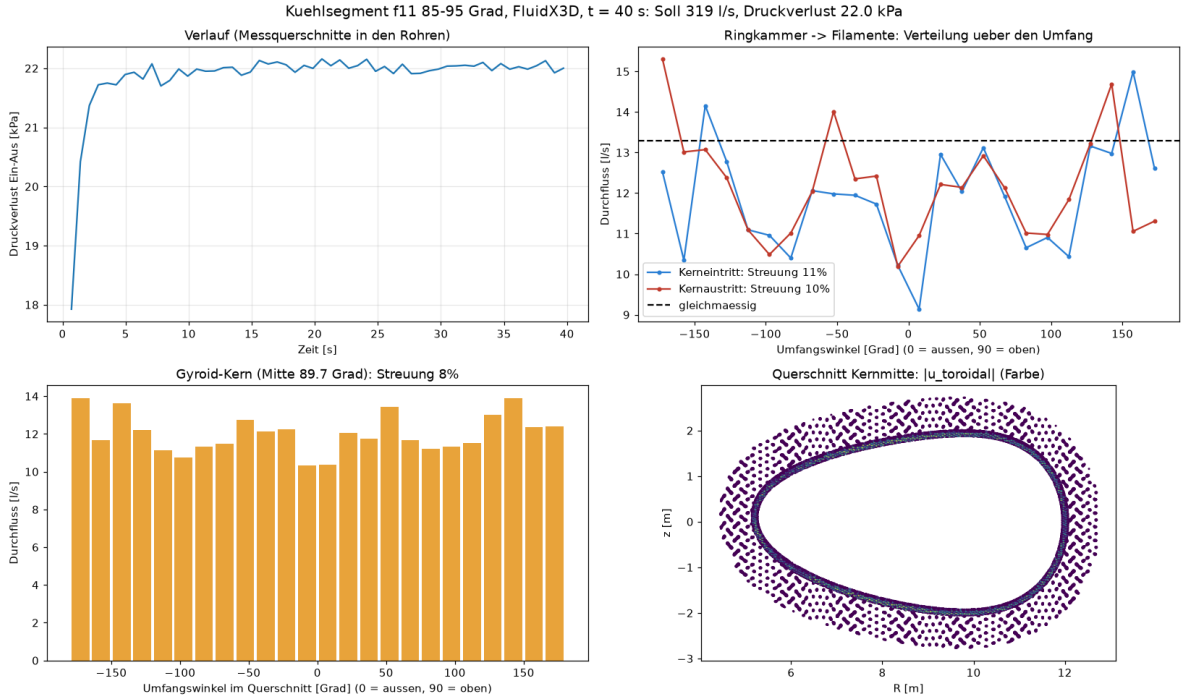


Figure 16: Evaluation of the FluidX3D run f11: pressure drop over time (steady after about 5 s), flow over the circumference at core inlet and outlet, flow through the gyroid core per sector (spread 8 %), and cross-section at mid core. Labels in German.

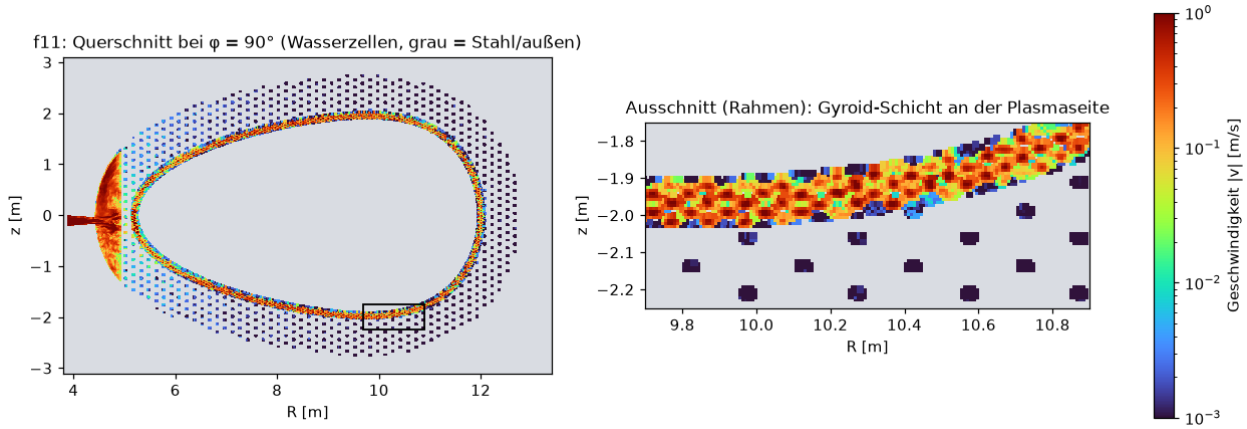


Figure 17: Cross-section through the f11 flow field at $\varphi = 90°$; every 8 mm lattice cell with water is a dot, colour = speed (log scale). Water enters on the left (inboard). Almost all the flow passes through the fine gyroid layer on the plasma side (zoom on the right, 0.1 m/s to 1 m/s). In the coarse dense gyroid behind it the water sits in closed pores; it only carries load. Labels in German.

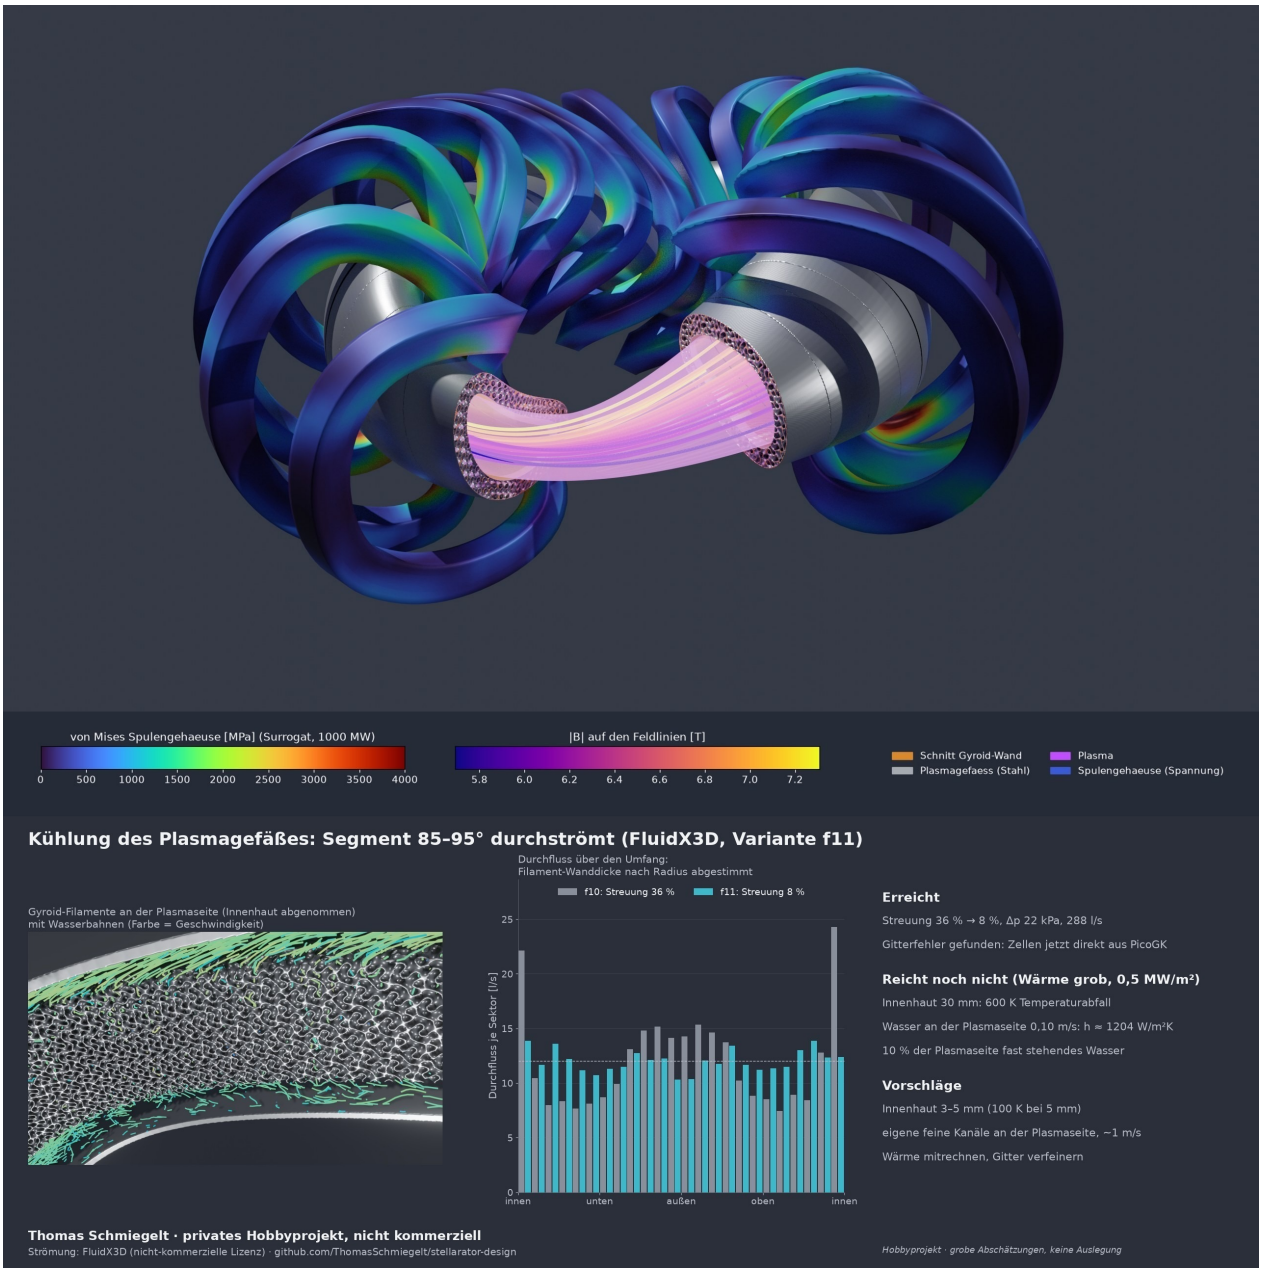


Figure 18: Overview of the 1 GW device extended by the cooling results: gyroid filaments on the plasma side with water paths (left), flow per circumferential sector before and after balancing (middle), findings (right). Labels in German.

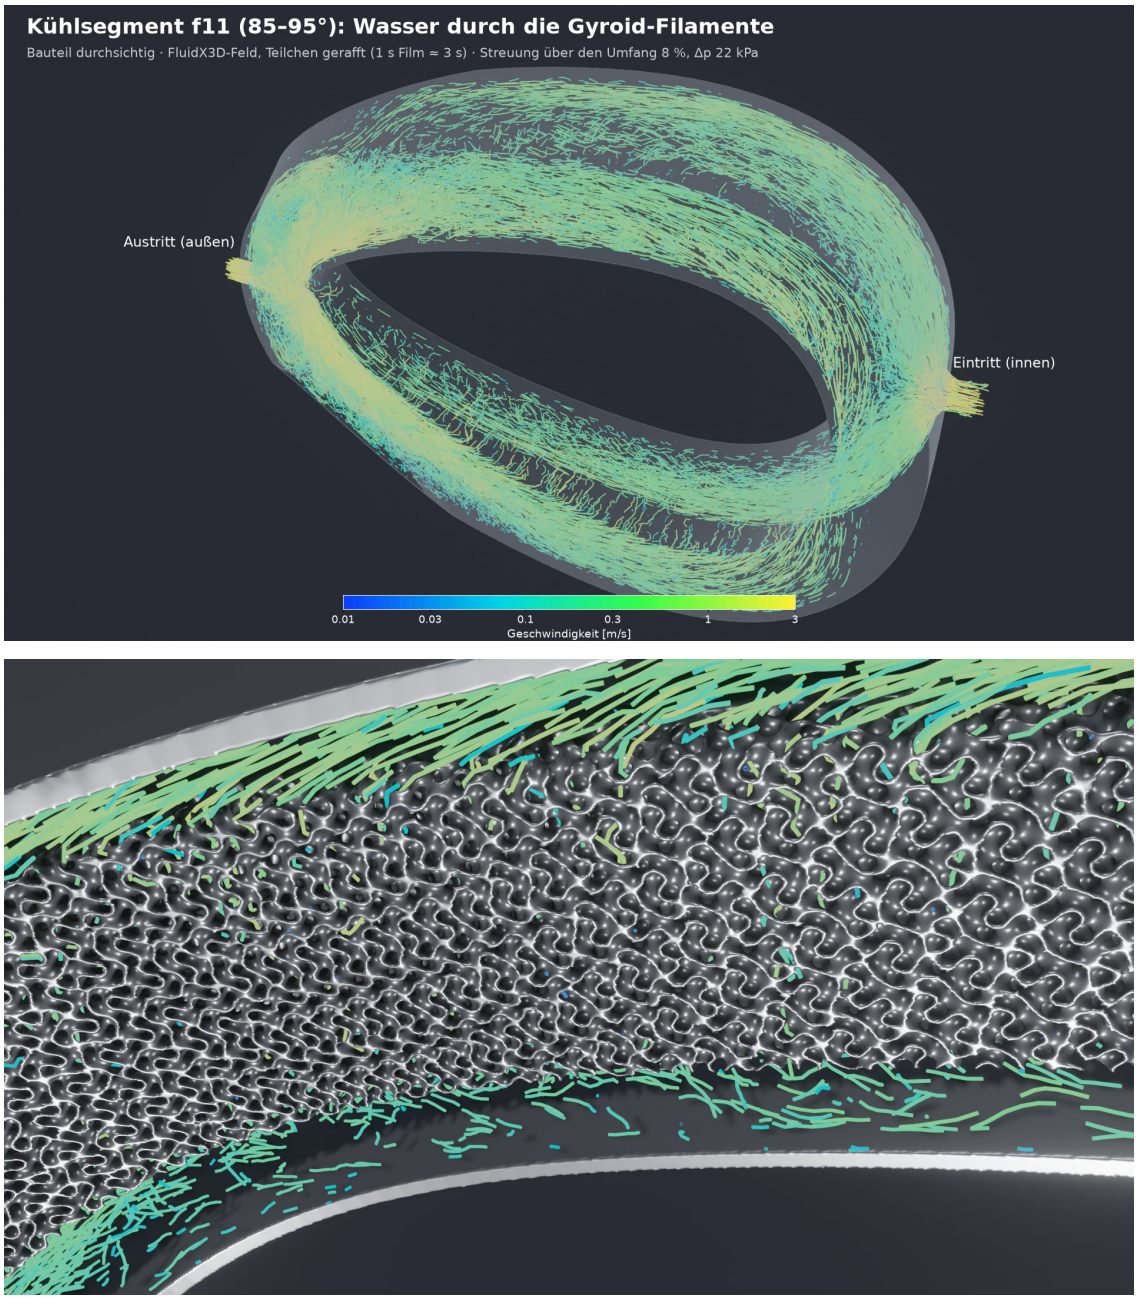


Figure 19: Frames from the flow films (Blender [14]): the segment rendered transparent, inlet inboard right, outlet outboard left (top); the lower half without the inner skin, looking at the plasma-side filaments (bottom). Colour = speed.

9.4 Is it enough? A rough heat estimate

The flow computation contains no heat. A one-dimensional estimate for pressurised water at 300 °C (Dittus–Boelter, steel with $k = 25\,\mathrm{W/(m\,K)}$) shows two bottlenecks.

Quantity at $0.5\,\mathrm{MW/m^2}$	Temperature drop
Conduction through the inner skin, 30 mm (as built)	600 K
Conduction through the inner skin, 5 mm	100 K
Transfer to the water, $v = 0.10\,\mathrm{m/s}$ (f11), without / with fin effect	415 K / 210 K
Transfer to the water, fine 10 mm channels at 1 m/s	48 K

Table 2: Rough heat estimate (`stellarator/kuehlung_waerme.py`).

The distribution is solved, the cooling is not yet. The inner skin is far too thick and the water on the

plasma side is too slow. In addition, on 10 % of the plasma side the water almost stands still (below 1 mm/s); hot spots would form there.

9.5 Suggested improvements

1. **Thinner inner skin:** 3 mm to 5 mm instead of 30 mm, possibly with tungsten armour. The gyroid core behind it carries the load.
2. **Faster flow on the plasma side:** either a thinner front layer (30 mm to 40 mm instead of 130 mm) with a finer cell, or a separate first-wall circuit with parallel channels at about 1 m/s. The gyroid behind then only carries load and removes volumetric heat. The pressure drop rises; at 22 kPa it is still small today.
3. **Remove stagnant water:** the near-zero speeds on the plasma side are probably at the end walls and chamber corners. They must be located in the field; guide vanes or a more open cell there are obvious options.
4. **Balance less sensitively:** one millimetre of wall thickness changes the flow by roughly a factor of 3.7. With typical manufacturing tolerances of ±0.1 mm that is about ±13 %. Defined orifices at the chambers, or a larger cell with thicker walls, are more robust.
5. **Compute the heat:** coupled flow and conduction (FluidX3D has a temperature extension, or Elmer for the steel). Only then can the fin effect of the filaments be judged.
6. **Check the lattice:** the front filament walls are about 10 mm thick and the lattice is 8 mm. A grid study at 4 mm to 6 mm is needed before the numbers are reliable.
7. **Reach the target flow:** 288 l/s instead of 319 l/s. The cause is the inlet boundary condition, not the geometry.

9.6 Collaboration

Here too, the decisive design ideas came from the user: water through the filaments, inlet inboard, inclined chamber boundaries. To agree on shape, a sketch on a screenshot worked best. This led to a small tool for drawing directly on the 3D model and handing the lines to the AI as 3D points [26]. The lattice error was found by the assistant, but only after several wrong guesses. The lesson: investigate an odd result before changing the design. The workflow is described in `docs/kuehlsegment_prozess.md` (German).

10 Follow-up: double filament, clearance and liquid lead

Of the improvements suggested in Section 9, the second one was taken up: the plasma side gets its own cooling circuit with fine channels. The requirements came from the user: the water should not get hotter than 150 °C, with an inlet of about 60 °C. On top of that came the idea of building the wall as a *double filament*: an inner part for cooling and, separate from it, an outer part for stiffness. All of it without colliding with the coils. All numbers in this section are rough estimates.

10.1 Wall layout

From the plasma side outwards: inner skin 5 mm, cooling channels, back wall 5 mm, dry gyroid for stiffness, outer skin 40 mm. The channels run toroidally across the segment from the inlet to the outlet plenum. The gyroid no longer carries water, so it may be more open and no longer has to be tuned to the cooling network.

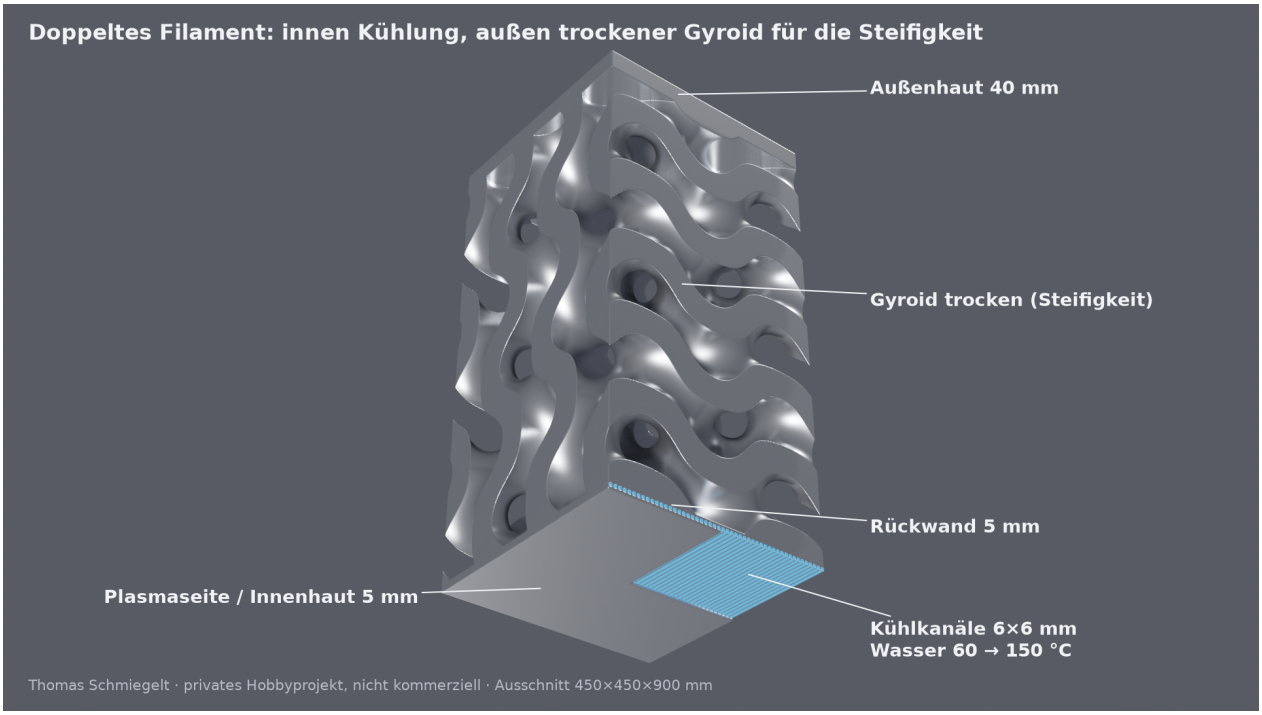


Figure 20: Wall piece with double filament (PicoGK, Blender), cut-out 450 × 450 × 900 mm. A window is cut into the inner skin; the cooling channels (blue) lie below. Labels in German.

10.2 Water at 60 to 150 °C

At 0.5 MW/m^2 the segment (85°–95°) takes up about 12.6 MW on the plasma side. With a temperature rise of 90 K that is only 35 l/s. Spread over all channels in parallel the water would be too slow and would boil at the wall. Therefore three neighbouring channels are connected in series (U-turn at the end face). Method: heat transfer after Gnielinski, friction after Petukhov, temperature in the cross-section with Elmer for the worst channel (`stellarator/erstwand_kanal.py`).

Channels	6 × 6 mm, rib 3 mm, 616 strands with 3 passes each
Velocity	0.96 m/s to 2.15 m/s
Pressure drop	49 kPa
Steel maximum	300 °C
Channel wall maximum	202 °C (boiling from 212 °C at 2 MPa)

Table 3: First-wall channels with water 60 °C → 150 °C at 2 MPa (variant n5).

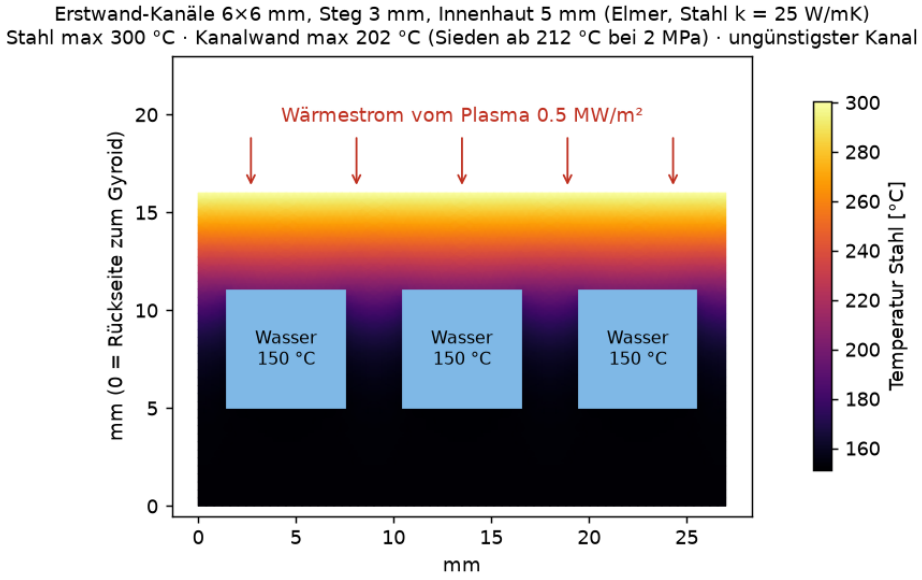


Figure 21: Steel temperature around the channels (Elmer, half pitch mirrored). Labels in German.

So high pressures are not needed at 150 °C; 2 MPa leaves enough margin to boiling. The real drawback lies elsewhere: heat at 150 °C can hardly be turned into electricity.

10.3 Clearance to the coils

For every point of the plasma boundary the smallest distance to the coil casings and connectors of the surrogate model was computed (`stellarator/gefaess_abstand.py`). The minimum is 1.25 m to the casings and 1.70 m to the connectors; in the computed segment it is at least 1.88 m. Including the gap the wall needs 1.05 m; with 5 cm margin it fits everywhere.

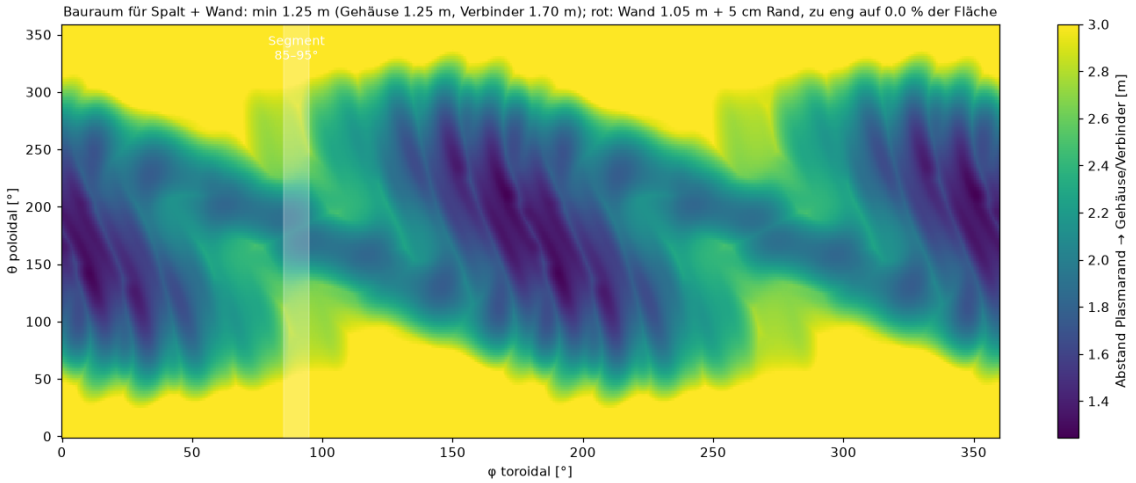


Figure 22: Distance from the plasma boundary to casings and connectors, unrolled over the torus. Red: limit at which the wall would no longer fit; it is reached nowhere. Labels in German.

10.4 Liquid lead instead of water?

As a next step the user proposed a *gyroid within a gyroid*: cooled near the plasma with liquid lead, which passes its heat on to the water. A gyroid wall does split space into two separate, intertwined networks and would make a compact heat exchanger. Lead needs no pressure, and lead-lithium would also breed tritium. Three points speak against putting the exchanger into the wall, however:

1. Lead melts at 327 °C, lead-lithium at about 235 °C. Next to water at 60 °C to 150 °C it would freeze.

2. A leak between hot lead and water right at the plasma is dangerous and cannot be repaired.
3. In the magnetic field, induced currents brake the metal as soon as it flows across the field.

For the third point the pressure drop was roughly estimated (`stellarator/blei_kanal.py`): $\mathrm{d}p/\mathrm{d}x = \sigma\, v\, B_\perp^2\, k_p$ with $k_p = c/(1+c)$ for conducting steel walls (c wall conductance ratio) and $k_p \approx 1/\mathrm{Ha}$ for an insulating layer. The field of a stellarator is mainly toroidal, so toroidal channels run almost along the field; the pitch of the field lines was estimated at about 6°.

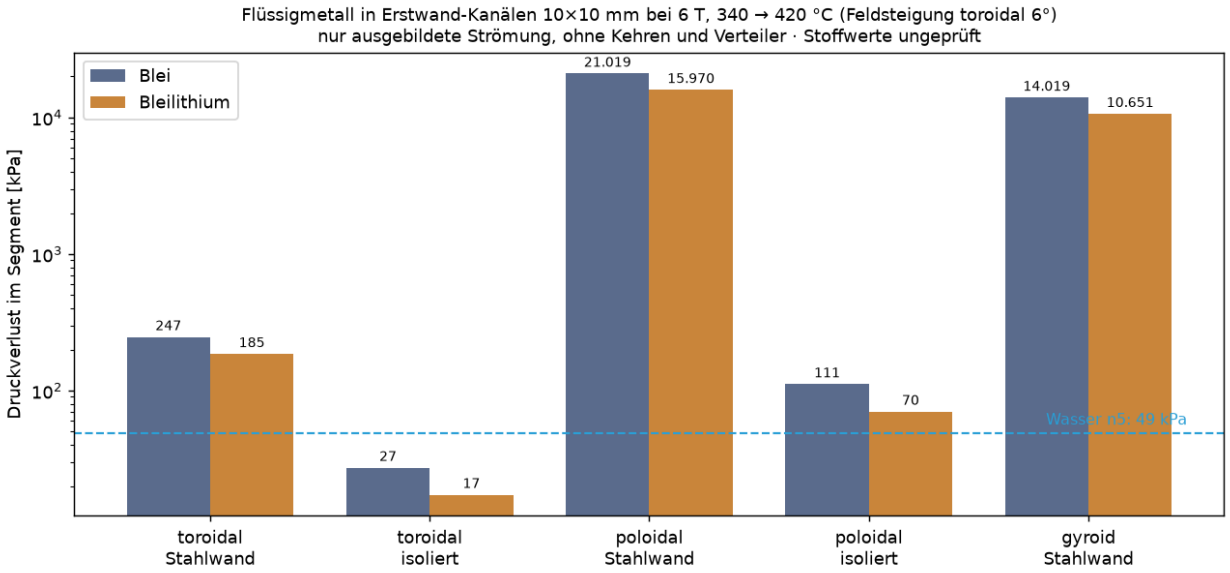


Figure 23: Pressure drop in the segment for lead and lead-lithium per channel direction at 6 T, logarithmic. Water (n5) for comparison. Labels in German.

Channel direction	Pressure drop	Pumping power / heat	
toroidal, steel wall	247 kPa	0.2 %	suitable
toroidal, insulated	27 kPa	0.02 %	suitable
poloidal, steel wall	21 MPa	17 %	unsuitable
gyroid, steel wall	14 MPa	11 %	unsuitable

Table 4: Lead in 10×10 mm channels, 340 °C → 420 °C, 0.87 m/s, 103 l/s per segment (variant b5).

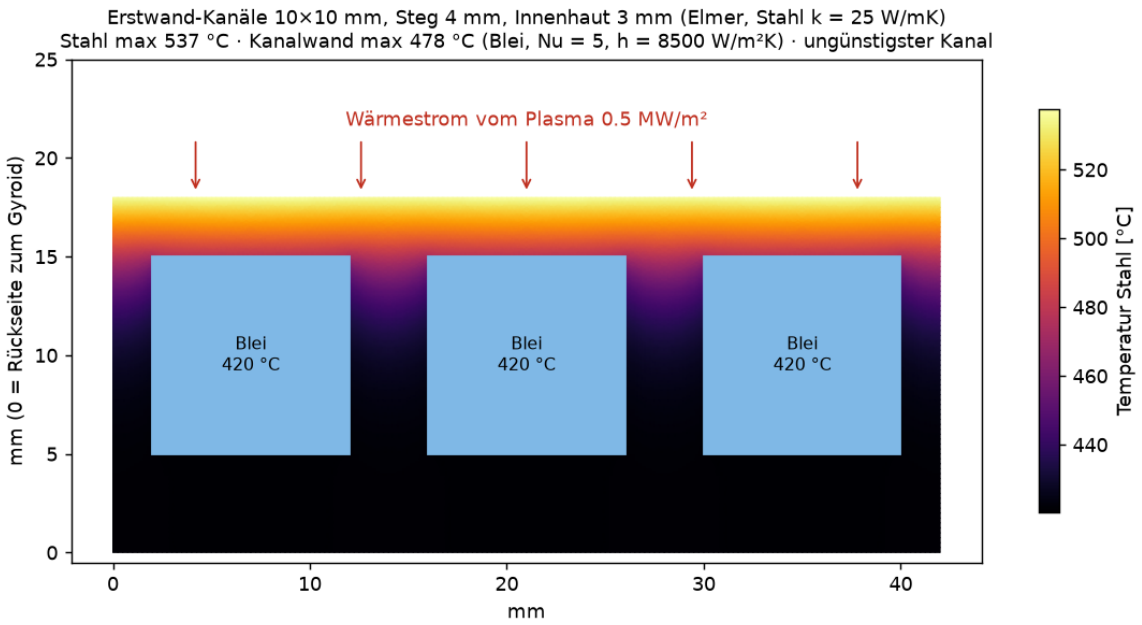


Figure 24: Steel temperature with lead (b5, inner skin 3 mm): at most 537 °C. With a 500 °C outlet and 20 mm channels it would be 711 °C. Labels in German.

Result: **the gyroid is not suitable for liquid metal**, because it keeps turning the flow across the field. Toroidal channels, on the other hand, work well. A sensible layout: lead or lead-lithium in toroidal channels on the plasma side, the dry gyroid behind it for stiffness, and a steam generator outside the coils that passes the heat to the water. At about 420 °C the heat can also be converted to electricity much better.

Not included are the additional losses in bends, manifolds and changes of cross-section (often the largest share), corrosion of steel in lead, volumetric heat from neutrons, and a more accurate heat transfer (conservatively Nu = 5 here). The material properties have not been checked against sources.

10.5 Collaboration

The ideas in this section came from the user: the temperature limit, the double filament, the check against the coils and the gyroid within a gyroid with lead. The assistant calculated them, named the limits and proposed variants. Along the way a small browser tool was built for looking at flow sections: one specifies section planes and hands them to the AI with "send and next section" or "send and finish" (`stellarator/schnitt_werkzeug.py`).

11 Coolant, strength and heat balance up to the turbine

This section follows the cooling from a single segment to the grid. Three questions started it: what does the thinner inner skin gain? Which variant is best when the figures of merit contradict each other? And which pressure and temperature are realistic? The user settled the last one: *reuse an existing turbine from a nuclear power plant and keep everything else as simple as possible.*

11.1 Variant f12: 5 mm inner skin

f12 is f11 with a 5 mm instead of 30 mm inner skin and a 500 mm outlet pipe (rule: outlet always larger than inlet). Flow with FluidX3D as in Section 9.

	f11	f12
Spread around the circumference (gyroid core)	8 %	11 %
Pressure drop	22.0 kPa	19.4 kPa
Flow rate	288 l/s	289 l/s
Velocity on the plasma side	0.10 m/s	0.10 m/s
Temperature drop across the inner skin at 0.5 MW/m^2	600 K	100 K

Table 5: f11 and f12 compared.

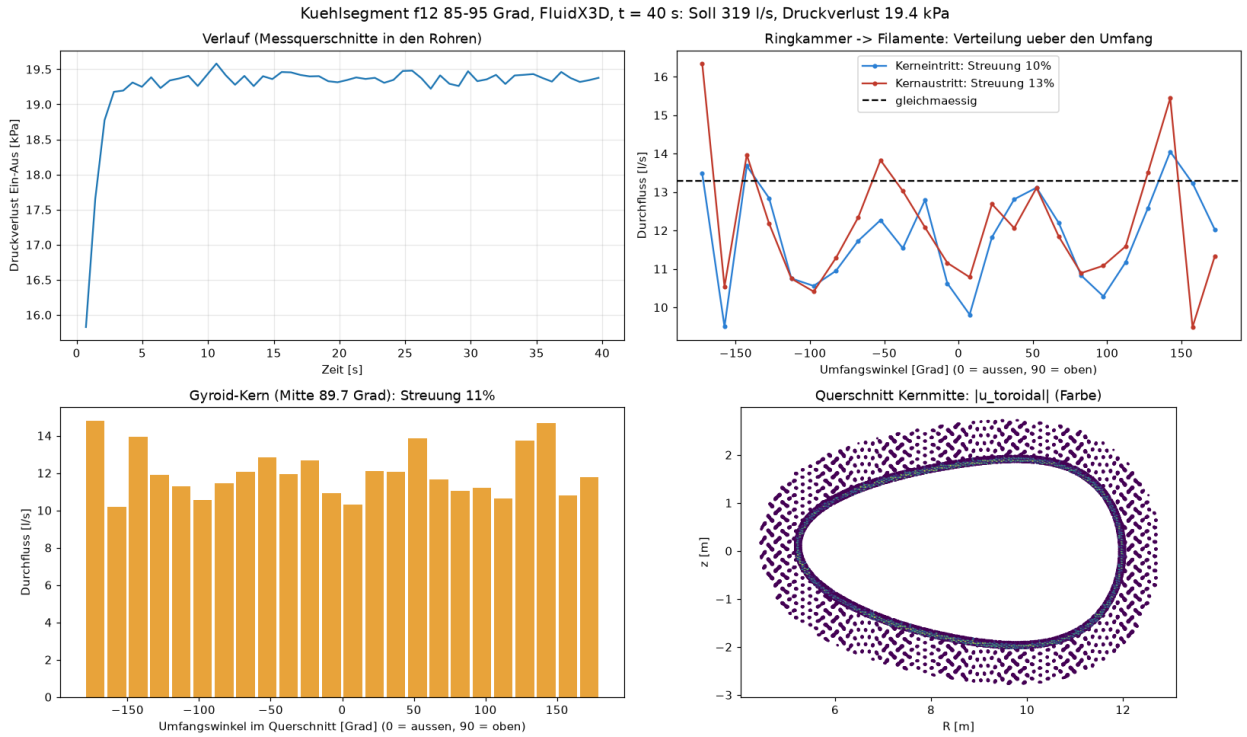


Figure 25: Evaluation of f12: pressure drop over time, distribution around the circumference and flow through the gyroid core.

The wall is no longer the bottleneck. What remains is the transfer into the slow water on the plasma side, about 410 K.

11.2 Ranking by pairwise comparison

Each variant yields several figures of merit that often conflict: f12 beats f11 on temperature and pressure drop but loses on spread. The user suggested deciding how to proceed by pairwise comparison so that resources are used well. Implemented in `stellarator/varianten_vergleich.py`:

1. **Pareto front:** a variant drops out if another one is worse in no criterion and better in at least one. No weights are needed.

2. **Weights by AHP** (Analytic Hierarchy Process): the user compares the criteria in pairs on a scale of 1 to 9. From the comparison matrix A with $a_{ji} = 1/a_{ij}$ the weight vector w is the eigenvector of the largest eigenvalue $\lambda_{\max}$. Consistency is checked by $\mathrm{CR} = \frac{(\lambda_{\max}-n)/(n-1)}{\mathrm{RI}_n}$ with the random index $\mathrm{RI}_4 = 0.90$; below 0.1 the judgements count as consistent.
3. **Rank:** weighted sum of the figures of merit normalised to 0 to 1.

The user's judgements: temperature is five times as important as spread, pressure drop and stagnant water each; pressure drop is three times as important as spread and five times as important as stagnant water. This gives weights of 0.61 for temperature, 0.23 for pressure drop, 0.08 for spread and 0.08 for stagnant water, with CR = 0.10, right at the limit.

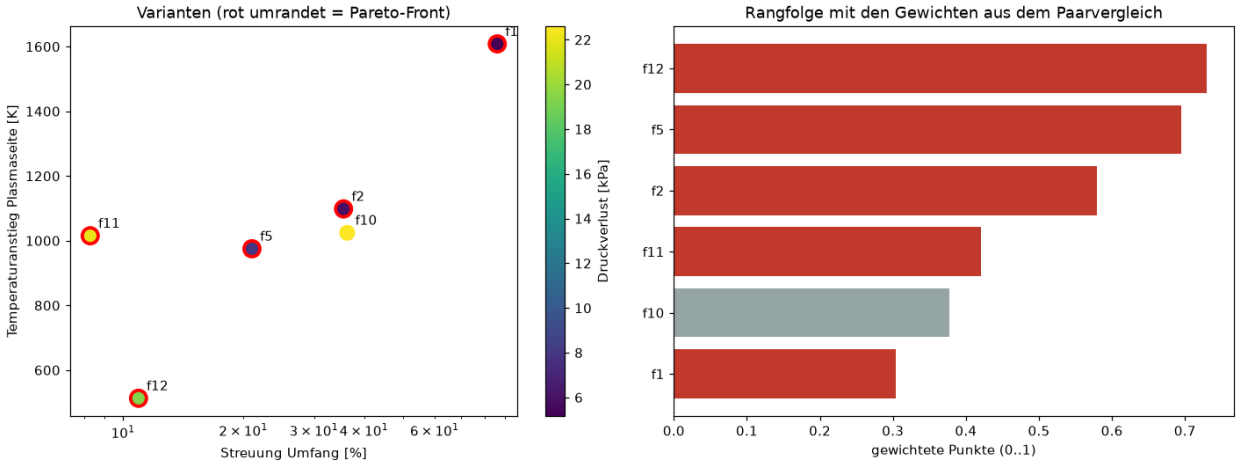


Figure 26: Left: variants by temperature rise and spread, colour pressure drop, red outline the Pareto front. Right: ranking with the user's weights.

Ranking: f12 (0.73) ahead of f5 (0.69), f2 (0.58), f11 (0.42), f10 (0.38) and f1 (0.30). Notably, f11 had been optimised for an even distribution, which the user in the end weighted low. The pairwise comparison thus showed that effort had gone into a secondary goal.

11.3 Does the thin skin hold the pressure?

A thin skin is good thermally but has to carry the coolant pressure. `stellarator/haut_festigkeit.py` computes with Elmer a steel plate of thickness t over 2×2 gyroid cells ($a = 100\,\mathrm{mm}$). It is clamped where the 10 mm gyroid walls touch the skin. In the plane of the skin these are the points with

$$\frac{|f|}{|\nabla f|}\,\frac{a}{2\pi} < \frac{w}{2}, \qquad f = \sin X\cos Y + \sin Y\cos Z + \sin Z\cos X, \quad X = 2\pi x/a.$$

In between, the coolant pushes towards the plasma. As a check, a strip with straight supports at a spacing of $b = a/2 - w = 40\,\mathrm{mm}$ was computed. With both edges clamped, $\sigma = p\,b^2/(2t^2)$: at $t = 5\,\mathrm{mm}$ Elmer gives 34.3 MPa per MPa of pressure, the formula 32.0 MPa.

Skin thickness	Stress per MPa of pressure	Allowable pressure at 150 MPa
5 mm	39 MPa	3.8 MPa
10 mm	14 MPa	10.5 MPa
15 mm	9 MPa	16 MPa

Table 6: Inner skin over the gyroid (Elmer, linear elastic, no thermal stresses).

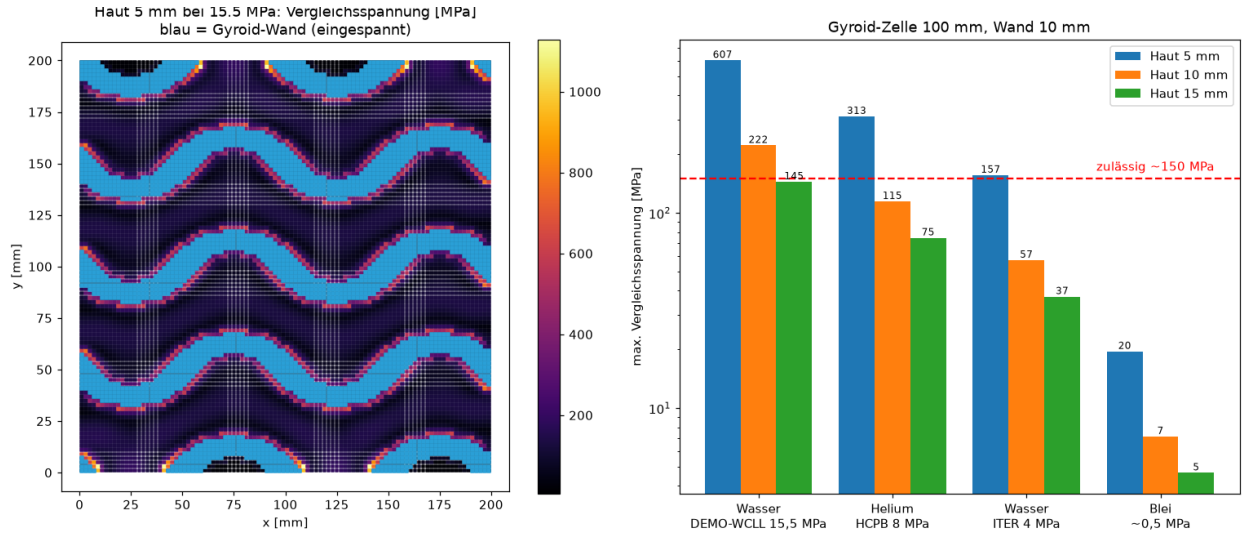


Figure 27: Left: von Mises stress in the 5 mm skin at 15.5 MPa, gyroid walls in blue. Right: peak stress per coolant and skin thickness.

For pressurised water as in a nuclear plant (≈ 16 MPa) a water-filled gyroid is therefore unsuitable. High-pressure water belongs in small channels: a 6 mm channel under a 5 mm skin reaches only about 11 MPa by the same formula. That is the double filament of Section 10.

11.4 Sourced reference values

Coolant data of real plants were researched with the local agent PI. A first attempt with another local agent (Hermes) produced invented values and links. For PI a skill `web-recherche` was therefore written: it actually fetches pages and records each value only with a verbatim quote and its source. A checker confirms that each quote appears in the fetched text. Because the local model has little context, PI researches each plant in a separate call; the user suggested this after a long run had produced nothing. All 29 quotes were confirmed. PI still misread two rows (inlet and outlet swapped at the steam generator); they were corrected on review.

Plant	Coolant	In / out	Pressure	Source
EPR	water	295 / 320 °C	15.5 MPa	Wikipedia
VVER-1000	water	289 / 319 °C	16.0 MPa	Wikipedia
BREST-OD-300	lead	340 / 540 °C	–	Wikipedia
ALFRED	lead	390 / 520 °C	unpressurised	Wikipedia
HTR-PM	helium	250 / 750 °C	7.0 MPa	OSTI report
ITER blanket	water	70 °C / –	4 MPa	iter.org
EU-DEMO HCPB	helium	300 / 500 °C	8 MPa	KIT

Table 7: Sourced coolant data; full quotes in `docs/recherche/pi/ergebnis.md`.

11.5 Decision: pressurised water as in a nuclear plant

With an existing turbine from a pressurised water reactor the steam conditions are fixed, and with them the primary circuit: pressurised water at about 289 °C → 319 °C and 16 MPa as in the VVER-1000. Steam generators, main coolant pumps and pressuriser are off-the-shelf. Lead (freezing at 327 °C, braking in the magnetic field) and helium (8 MPa, blowers) are dropped. The first-wall channels (6×6 mm, three passes) were recomputed for these conditions: steel at most 441 °C, channel wall 340 °C (boiling at 345 °C), pressure drop 0.24 MPa. The steel thus sits in the working range of EUROFER (about 300 °C to 550 °C); with water at 60 °C to 150 °C it was below it. The margin to boiling, 5 K, is tight.

11.6 Derivation of the heat balance

The chain runs backwards from the grid to the plasma, with the same assumptions as the power balance of stage 0 (`invers/stufe0_bilanz.py`), here split by component (`stellarator/waermebilanz.py`).

1. Thermal power from grid power. P_{grid} goes to the grid; the turbine converts the heat P_{th} into electricity with efficiency η. Pumps ($f_p P_{\text{th}}$) and the remaining house load P_{house} (coil refrigeration, tritium plant, buildings) are subtracted:

$$P_{\text{grid}} = \eta P_{\text{th}} - f_p P_{\text{th}} - P_{\text{house}} \quad \Rightarrow \quad P_{\text{th}} = \frac{P_{\text{grid}} + P_{\text{house}}}{\eta - f_p} = \frac{1000 + 60}{0.35 - 0.04}\,\text{MW} = 3419\,\text{MW}.$$

2. Fusion power. One D-T fusion releases 17.59 MeV; the alpha particle carries 3.52 MeV, the neutron 14.07 MeV. The alpha fraction is $f_\alpha = 3.52/17.59 = 0.200$. Nuclear reactions in the blanket multiply the neutron energy by $M = 1.15$:

$$P_{\text{th}} = P_{\text{fus}}\left[(1 - f_\alpha)M + f_\alpha\right] = 1.120\,P_{\text{fus}} \quad \Rightarrow \quad P_{\text{fus}} = 3053\,\text{MW}.$$

3. Split by component. The neutrons pass through the first wall and deposit their energy in the blanket. The alpha power stays in the plasma and leaves it as radiation (onto the whole first wall) or as particle flow along the field lines (onto the divertor). Half each way is assumed ($f_{\text{div}} = 0.5$):

$$P_{\text{blanket}} = (1 - f_\alpha)\,M\,P_{\text{fus}} = 2808\,\text{MW},$$
$$P_{\text{first wall}} = (1 - f_{\text{div}})\,f_\alpha\,P_{\text{fus}} = 305\,\text{MW}, \qquad q_{\text{first wall}} = P_{\text{first wall}}/S = 305\,\text{MW}/1230\,\text{m}^2 = 0.25\,\text{MW/m}^2,$$
$$P_{\text{divertor}} = f_{\text{div}}\,f_\alpha\,P_{\text{fus}} = 305\,\text{MW}, \qquad A_{\text{divertor}} \geq P_{\text{divertor}}/q_{\max} = 305\,\text{MW}/10\,\text{MW/m}^2 = 31\,\text{m}^2.$$

S is the plasma surface of the HTS design (from stage 0: neutron wall load 1.99 MW/m^2). The first-wall channels are designed for 0.5 MW/m^2, a factor of 2 margin for local peaks. 10 MW/m^2 is the design load of the water-cooled divertor elements of Wendelstein 7-X (Wikipedia).

4. Cooling water. With $\Delta T = 30\,\text{K}$ and $c_p = 5380\,\text{J/(kg K)}$, $\rho = 732\,\text{kg/m}^3$ at about 300 °C:

$$\dot{m} = \frac{P}{c_p\,\Delta T}: \qquad \dot{m}_{\text{blanket}} = 17\,400\,\text{kg/s}\ (23.8\,\text{m}^3\text{/s}), \qquad \dot{m}_{\text{first wall}} = 1890\,\text{kg/s}\ (2.6\,\text{m}^3\text{/s}).$$

The pumping power for the first-wall channels alone is small: $\dot{V}\,\Delta p/\eta_{\text{pump}} = 26.4\,\text{m}^3\text{/s} \cdot 0.24\,\text{MPa}/0.8 \approx 8\,\text{MW}$, even if all the water passed through such channels. The assumed $f_p P_{\text{th}} = 137\,\text{MW}$ leaves room for steam generators, piping and blanket.

5. Turbine. $P_{\text{gross}} = \eta\,P_{\text{th}} = 1197\,\text{MW}$, waste heat at the condenser $(1 - \eta)P_{\text{th}} = 2223\,\text{MW}$. Comparison with sourced plants:

- **Isar 2 (Konvoi):** 3950 MW thermal, 1485 MW gross, 1410 MW net, i.e. $\eta_{\text{gross}} = 0.376$; shut down since 15 April 2023. The fusion plant would need 87 % of its heat and 81 % of its turbine output: **one turbine of this class is enough**.
- **VVER-1000:** 3000 MW thermal, net efficiency 31.7 %, about 1 GW electric. The fusion plant delivers 114 % of this heat; a single turbine of this size would be too small.

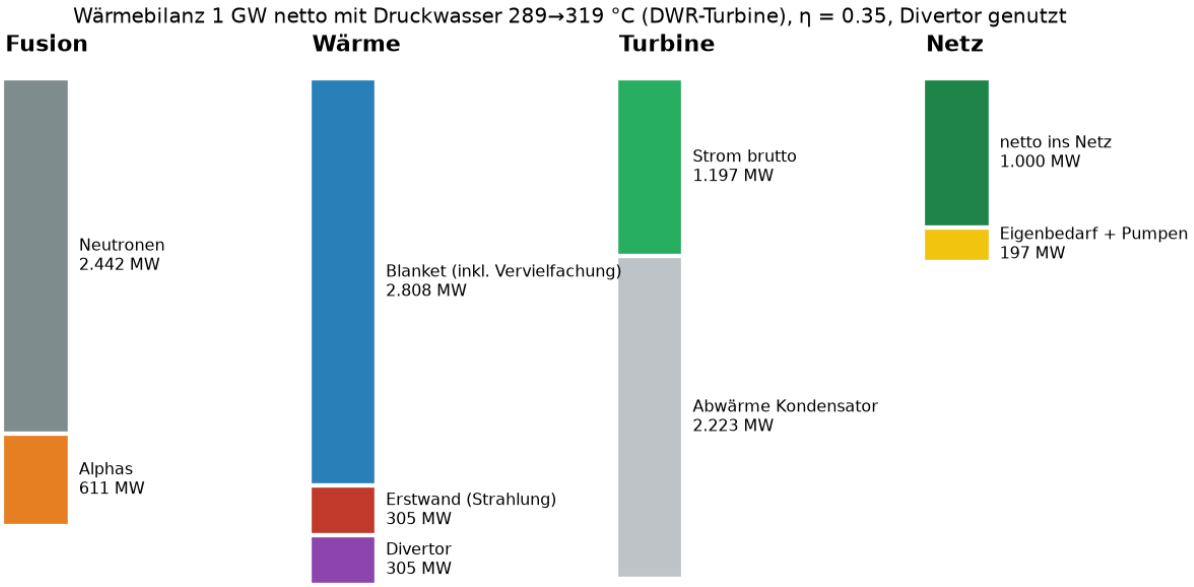


Figure 28: Heat balance for 1 GW net with pressurised water and a turbine from a pressurised water reactor (labels in German).

6. Sensitivity. Efficiency has the largest effect on the required fusion power:

Case	P_{fus}	P_{th}	P_{gross}
$\eta = 0.35$, divertor heat used	3053 MW	3419 MW	1197 MW
$\eta = 0.376$ (as Isar 2)	2817 MW	3155 MW	1186 MW
$\eta = 0.317$ (VVER-1000, net)	3417 MW	3827 MW	1213 MW
$\eta = 0.35$, divertor heat lost	3353 MW	3419 MW	1197 MW

Table 8: Fusion power required for 1 GW net.

The divertor is cooled at low temperature (copper alloy). Its heat can preheat the feedwater. If it is lost, the required fusion power rises by 10 %.

Limits. f_{div} and M are assumptions. Part of the neutron heat is already deposited in the first wall and is counted here with the blanket. The blanket itself is open: it must remove 2.8 GW at pressurised-water conditions and breed tritium at the same time. The EU-DEMO WCLL concept (water as coolant, lead-lithium as breeder) aims at exactly this; sourced values are still missing here because the publishers' sites block automated access. Whether a turbine from a decommissioned plant is actually available was not checked; the reference plants only show the matching size class.

11.7 Do the pipes fit through the gaps between the coils?

All the cooling water has to go from the vessel through the coils to the steam generators and back. On the inboard side the winding packs of the HTS design leave only 0.58 m (stage 3). This question could rule out the pressurised-water solution, so it was checked before anything else (`stellarator/rohr_durchfuehrung.py`).

Demand. From the heat balance: $\dot{V} = (\dot{m}_{\mathrm{blanket}} + \dot{m}_{\mathrm{first\ wall}})/\rho = 19\,290\,\mathrm{kg/s} \,/\, 732\,\mathrm{kg/m^3} = 26.4\,\mathrm{m^3/s}$ in each direction, so twice that pipe area for supply and return. With velocity v and inner diameter d_i one needs

$$n = \frac{2\,\dot{V}}{v\,\pi d_i^2/4} \quad \text{pipes, e.g. } n = \frac{2 \cdot 26.4}{10 \cdot \pi \cdot 0.8^2/4} = 11 \text{ at } d_i = 0.8\,\mathrm{m},\ v = 10\,\mathrm{m/s}.$$

The pipe wall follows from the boiler formula $s = p\,d_i/(2\sigma)$: at 16 MPa and 150 MPa, $s = 0.053\,d_i$, so the outer diameter is $D = d_i(1 + p/\sigma) = 1.11\,d_i$.

Space. From each point of the vessel's outer surface (plasma boundary plus 1.05 m) a straight pipe runs 4 m outward along the surface normal. The smallest distance $a_{\min}$ of this path to winding pack, potting, casing and connectors of the surrogate model ($R_0 = 10.4$ m) limits the pipe diameter:

$$D_{\max} = 2\,(a_{\min} - a_{\text{shield}}), \qquad a_{\text{shield}} = 0.15\,\text{m}.$$

a_{shield} is room for a gap and a thermal shield: the pipe is at 300 °C, the superconducting coils are cryogenic. Pipes must be at least $D + 0.2$ m apart; counting was greedy, freest spots first.

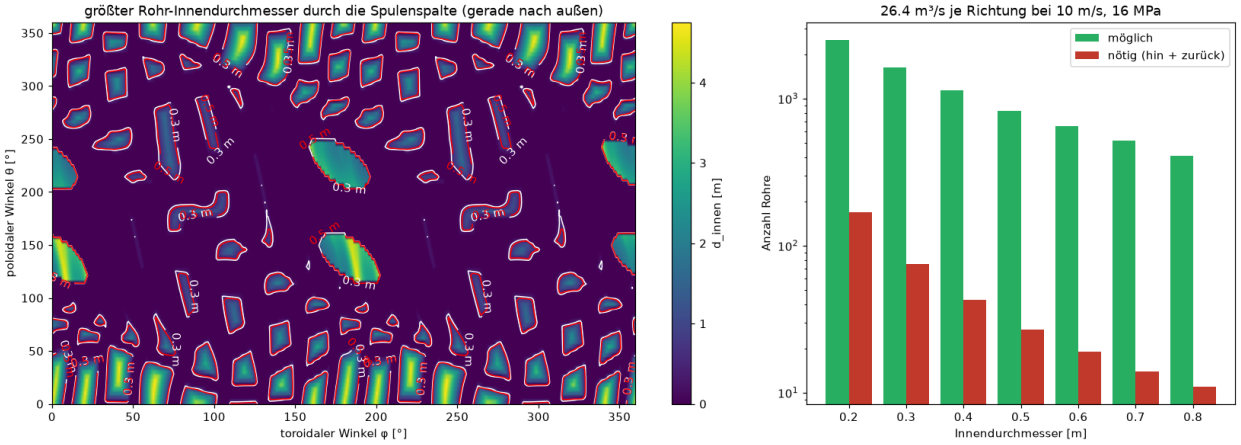


Figure 29: Left: largest inner diameter of a straight pipe over the unrolled vessel surface ($\theta = 0$ outboard, 180° inboard); the windows between coils and connectors are clearly visible. Right: possible and required pipes per size (labels in German).

d_i	required ($v = 10$ m/s)	required ($v = 5$ m/s, shield 0.3 m)	possible (default / cautious)
0.3 m	75	150	1634 / 1386
0.5 m	27	54	824 / 732
0.8 m	11	22	410 / 340

Table 9: Pipes through the gaps between the coils.

Result. **The pipes fit with plenty of room**, by more than a factor of 10 even in the cautious case. On 19 % of the vessel surface a straight path for 0.8 m pipes is free. The windows lie mainly outboard (θ near 0) and at top and bottom. Inboard (θ around 180°) almost everything is covered by coils: the inboard blanket has to be fed through headers from above and below. The initial concern about 0.58 m applied only to the narrowest inboard gap.

Limits. The support structure between the coils is still open (forces about 1000 times those of W7-X). It will occupy part of the windows, as will access for heating, diagnostics, vacuum pumps and maintenance. Only straight pipes along the normal were considered; with bends more room would be found.

11.8 Collaboration

From the user came: the thinner wall with a larger outlet (f12), pairwise comparison as a tool for decisions and optimisation loops together with the judgements, the point that the walls of the pressure space must withstand the pressure, the point about the small context of the local model, and the decision for an existing turbine. The assistant did the calculations, built the tools, checked the research and named the consequences.

Code and data

Software: ElmerFEM, gmsh, FreeCAD, simsopt, DESC, PyVista, ParaView, Blender, FluidX3D 3.8 [17], PicoGK [25]. `https://github.com/ThomasSchmiegelt/stellarator-design` (MIT licence): pipeline code, findings log, figures.

References

[1] H.-S. Bosch, R. C. Wolf, T. Andreeva et al.: Technical challenges in the construction of the steady-state stellarator Wendelstein 7-X. *Nucl. Fusion* 53 (2013) 126001. doi:10.1088/0029-5515/53/12/126001

[2] Th. Rummel, K. Rieße, H. Ehmler: Manufacture and test of the non-planar coils for WENDELSTEIN 7-X. *Fusion Eng. Des.* 75–79 (2005) 117–121. doi:10.1016/j.fusengdes.2005.06.118

[3] R. K. Maix, V. Bagnato, M. Fricke et al.: Design, production and QA test results of the NbTi CIC conductors for the W7-X magnet system. *J. Phys.: Conf. Ser.* 43 (2006) 753–758. doi:10.1088/1742-6596/43/1/184

[4] C. Bachmann, M. Siccinio, A. Ciula et al.: Re-design of EU DEMO with a low aspect ratio. *Fusion Eng. Des.* 204 (2024) 114518.

[5] P. F. Gil, V. Brack, T. Schuler, P. Huslage, E. V. Stenson: Manufacturing tolerances of non-planar coils for an optimized tabletop stellarator. arXiv:2507.22516 (2025).

[6] D. Schmeling, M. Russo, B. T. Gebreamlak et al.: Experimental results from early non-planar NI-HTS magnet prototypes for the Columbia Stellarator eXperiment (CSX). *IEEE Trans. Plasma Sci.* (2026), doi:10.1109/TPS.2026.3681798; arXiv:2510.07042.

[7] M. Landreman, B. Medasani, F. Wechsung, A. Giuliani, R. Jorge, C. Zhu: SIMSOPT: A flexible framework for stellarator optimization. *J. Open Source Softw.* 6(65) (2021) 3525. doi:10.21105/joss.03525. W7-X coil data: `simsopt/configs/W7-X.dat`, coil set `w7x_v001`.

[8] D. W. Dudt, E. Kolemen: DESC: A stellarator equilibrium solver. *Phys. Plasmas* 27 (2020) 102513.

[9] C. Geuzaine, J.-F. Remacle: Gmsh: A 3-D finite element mesh generator with built-in pre- and post-processing facilities. *Int. J. Numer. Meth. Eng.* 79 (2009) 1309–1331.

[10] ElmerFEM, version 9.0. CSC – IT Center for Science, Finland. `https://www.elmerfem.org`

[11] FreeCAD, version 1.1. `https://www.freecad.org`

[12] C. B. Sullivan, A. Kaszynski: PyVista: 3D plotting and mesh analysis through a streamlined interface for the Visualization Toolkit (VTK). *J. Open Source Softw.* 4(37) (2019) 1450.

[13] J. Ahrens, B. Geveci, C. Law: ParaView: An end-user tool for large data visualization. In: *Visualization Handbook*, Elsevier (2005).

[14] Blender, version 4.0. Blender Foundation. `https://www.blender.org`

[15] Nous Research: Hermes Agent. `https://github.com/NousResearch/hermes-agent` (MIT licence).

[16] Anthropic: Claude Code. `https://www.anthropic.com/claude-code`

[17] M. Lehmann: FluidX3D, version 3.8. `https://github.com/ProjectPhysX/FluidX3D`. Licence: free for non-commercial use only, source-available, not open source. Altered version of `setup.cpp` and `defines.hpp` in the repository folder `fluidx3d/`.

[18] M. Lehmann: Computational study of microplastic transport at the water-air interface with a memory-optimized lattice Boltzmann method. PhD thesis, University of Bayreuth (2023). doi:10.15495/EPub_UBT_00006977

[19] M. Lehmann: Esoteric Pull and Esoteric Push: Two simple in-place streaming schemes for the lattice Boltzmann method on GPUs. *Computation* 10 (2022) 92. doi:10.3390/computation10060092

[20] M. Lehmann, M. Krause, G. Amati, M. Sega, J. Harting, S. Gekle: Accuracy and performance of the lattice Boltzmann method with 64-bit, 32-bit, and customized 16-bit number formats. *Phys. Rev. E* 106 (2022) 015308.

[21] M. Lehmann: Combined scientific CFD simulation and interactive raytracing with OpenCL. *IWOCL'22: International Workshop on OpenCL* (2022) 3, 1–2.

[22] M. Lehmann, L. M. Oehlschlägel, F. Häusl, A. Held, S. Gekle: Ejection of marine microplastics by raindrops: a computational and experimental study. *Microplastics and Nanoplastics* 1 (2021) 18. doi:10.1186/s43591-021-00018-8

[23] M. Lehmann: High performance free surface LBM on GPUs. Master's thesis, University of Bayreuth (2019). doi:10.15495/EPub_UBT_00005400

[24] M. Lehmann, S. Gekle: Analytic solution to the piecewise linear interface construction problem and its application in curvature calculation for volume-of-fluid simulation codes. *Computation* 10 (2022) 21. doi:10.3390/computation10020021

[25] LEAP 71: PicoGK. https://github.com/LEAP71/PicoGK (Apache 2.0 licence).

[26] sketch-to-prompt: sketches on 3D models as input for AI assistants. https://github.com/ThomasSchmiegelt/sketch-to-prompt (MIT licence).