Two AI Agents, One Stellarator

Building an engineering simulation pipeline from an open dataset – including the dead ends

Thomas Schmiegelt | hobby project `stellarator-design`

25 October 2026, updated 27 October 2026

Hobby project, driven by curiosity. Not affiliated with any research institute or company, not peer reviewed. All numbers are estimates from simplified models and may contain errors. The W7-X model is *inspired by* Wendelstein 7-X: coil paths come from public data [7]; casings, connectors, potting, supports and wall thicknesses are my own simplified assumptions and do not describe the real device. Code, computations and the draft text were produced with the help of AI agents; all figures are my own computations and renderings.

Introduction and motivation

I am not a fusion researcher. This is a private project, driven by curiosity: how does a stellarator work, and how far can one person get with open software and AI assistance on such a demanding part? Stellarator coils are a good example because design, manufacturing and strength meet here. Physics dictates the path, current and field create the forces, and the structure has to fit in between. Later the plasma vessel was added, with a cooling core of gyroid filaments.

A stellarator confines a fusion plasma with twisted 3D superconducting coils that carry large magnetic forces. In four days, a human and two AI agents built an end-to-end pipeline from a public coil dataset to a structural analysis of the coil casings: coil path → magnetic field → Lorentz force → casing geometry → stress. This report is mostly about *how* it was built: the division of labour between a local agent (Hermes) and a coding agent (Claude Code), a first model that was wrong while looking right, a fully parametric CAD model that broke the mesher, and a second, CAD-free code path that finally worked. For a W7-X-like system of 70 coils the result is 99 % of casing stresses below 43 MPa – under strongly simplified supports. A follow-up (Section 9) computes the flow through the gyroid cooling core of the 1 GW plasma vessel with FluidX3D and suggests improvements.

Contents

1 Setup **3**
- 1.1 Agents and responsibilities . . . 3
- 1.2 Tool chain . . . 3

2 Research phase **3**

3 First results – and why they were wrong **4**
- 3.1 Validating the solver . . . 4
- 3.2 The surrogate model (report v0.1) . . . 4

4 Restart: mesh directly from the data **6**

5 The parametric CAD branch **10**

6 The parallel branch: no CAD kernel **12**

7 Visualisation: from mesh to Blender **16**

8 Takeaways for working with coding agents **20**

9 Follow-up: cooling the plasma vessel **21**
- 9.1 Task and layout . . . 21
- 9.2 History . . . 21
- 9.3 An error in the lattice . . . 21
- 9.4 Is it enough? A rough heat estimate . . . 25
- 9.5 Suggested improvements . . . 26
- 9.6 Collaboration . . . 26

10 Follow-up: double filament, clearance and liquid lead **27**
- 10.1 Wall layout . . . 27
- 10.2 Water at 60 to 150 °C . . . 27
- 10.3 Clearance to the coils . . . 28
- 10.4 Liquid lead instead of water? . . . 28
- 10.5 Collaboration . . . 30

11 Coolant, strength and heat balance up to the turbine **31**
- 11.1 Variant f12: 5 mm inner skin . . . 31
- 11.2 Ranking by pairwise comparison . . . 31
- 11.3 Does the thin skin hold the pressure? . . . 32
- 11.4 Sourced reference values . . . 33
- 11.5 Decision: pressurised water as in a nuclear plant . . . 33
- 11.6 Derivation of the heat balance . . . 34
- 11.7 From plasma to grid: power flow and influences . . . 35
- 11.8 The components of the chain . . . 36
- 11.9 Do the pipes fit through the gaps between the coils? . . . 38
- 11.10 Blanket: comparison with the EU-DEMO WCLL concept . . . 39
- 11.11 Graded structure: flat contact at the skin, branching outward . . . 39
- 11.12 How does the heat spread sideways and outward? . . . 41
- 11.13 Divertor preheater: which design yields the most electricity? . . . 44
- 11.14 Exchanger modules built and simulated with FluidX3D . . . 47
- 11.15 Surrogate model from the flow simulations and the question of grading . . . 49
- 11.16 Collaboration . . . 52

1 Setup

Stellarator in five sentences. A fusion plasma (over 100 million degrees) is held by a magnetic field inside a ring-shaped vacuum vessel. To keep particles from drifting out, field lines must wind helically around the ring. A tokamak gets this twist from a large current in the plasma itself; a stellarator gets it entirely from external coils, which therefore have to be twisted in 3D. This makes steady-state operation easier but turns each coil into a demanding 3D part with tight tolerances. The reference device is Wendelstein 7-X (IPP Greifswald) with 50 non-planar and 20 planar superconducting coils [1].

1.1 Agents and responsibilities

Human	Goal and direction, reviews results *as images*, decides, supplies the key design ideas when stuck.
Hermes	Hermes Agent [15], run locally with open-weight models: literature research, project bookkeeping, writing reusable *skills*, long-running jobs overnight.
Claude Code	Coding agent [16]: simulation pipeline, debugging, post-processing, figures.

The agents never talk directly. They coordinate through files in the repository: an append-only findings log (`BEFUNDE.md`; every entry has observation, measurement, location, status), a handover file for decisions and running jobs (`ERINNERUNGEN.md`), and skills – short Markdown playbooks an agent loads before a task. Long simulations are started detached (`setsid nohup`) so they survive a session ending or a context compaction. The very first human intervention was a stop: Hermes started building tooling right away and was told to do the research first.

1.2 Tool chain

Physics data	simsopt [7] (coil Fourier coefficients), DESC [8] (plasma shape optimisation)
Geometry	Python + gmsh/OpenCASCADE [9] (direct meshing), FreeCAD 1.1 headless [11] (parametric CAD)
Solver	ElmerFEM 9.0 [10]: magnetostatics and linear elasticity in one coupled run
Own code	Biot-Savart (numba), FFT field solver, field-line tracer, collision checks
Visualisation	PyVista/VTK [12], ParaView [13], Blender 4.0 [14] via glTF
Hardware	one workstation with one GPU (reserved for the local agent)

2 Research phase

Hermes collected the basics: W7-X has a major radius of 5.5 m, five field periods and 2.5 T on axis [1]. Each coil is wound from a cable-in-conduit conductor: 243 NbTi strands of 0.57 mm, each with 144 filaments in copper, inside a co-extruded 16 mm × 16 mm aluminium jacket with 37 % void for helium cooling [2, 3]. From a set of 33 PDFs supplied by the user, four were relevant; the most useful one sizes coil structures for EU-DEMO with a 660 MPa stress limit and wedged inner legs [4]. A rate-limited arXiv search added work on HTS strain sensitivity [6] and on manufacturing tolerances (W7-X: 2 mm) [5].

Lesson – verify citations. The bibliography an agent produced during research contained entries without authors and with wrong venues or years. For this report every reference was checked individually against full text or publisher metadata (Crossref, arXiv API). Treat agent-written references as unverified until checked.

3 First results – and why they were wrong

3.1 Validating the solver

Before simulating coils, ElmerFEM was tested on a case with an exact answer: an infinite straight wire, $B = \mu_0 I / 2\pi r$. The first runs returned zero everywhere – silently. Root causes: a mesh format the converter misread, and wrong keywords in the solver input file, which Elmer ignores without an error. Starting from Elmer's own regression test instead of hand-written input fixed it: 0.09 % error inside the wire, 2 % outside (Fig. 4).

3.2 The surrogate model (report v0.1)

Hermes believed the real coil paths were unavailable (the URLs it tried returned 404) and generated a parametric surrogate: ellipses and bands, current calibrated to 2.5 T. Compared with W7-X the ampere-turns per coil matched to 0.90×. A report was written.

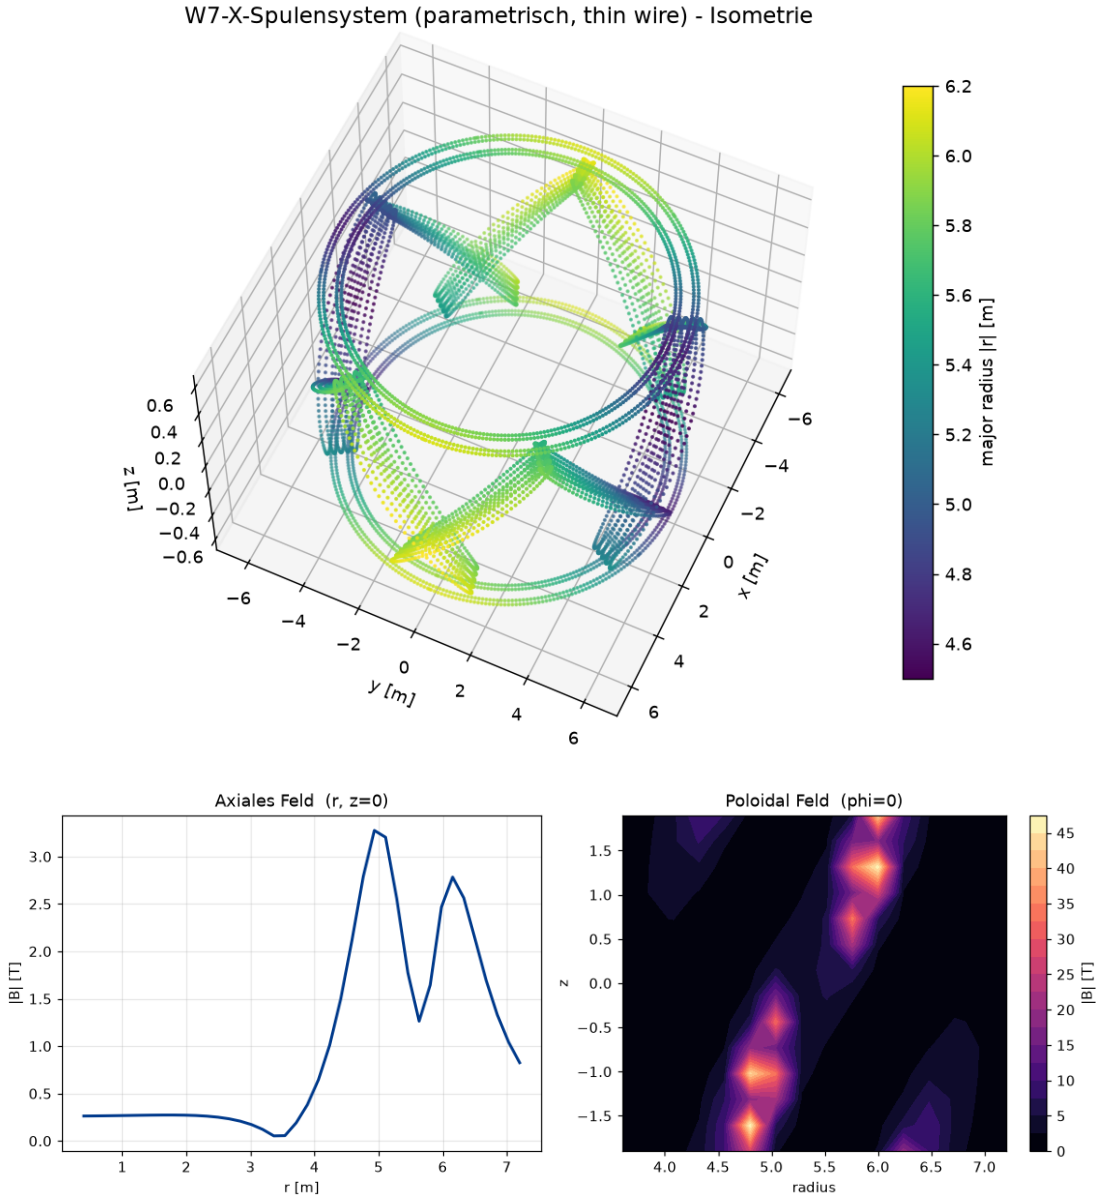

Figure 1: Report v0.1: the surrogate coil set and its field. Plausible numbers – but not a stellarator.

A topological check by the coding agent – does each coil actually encircle the plasma axis? – failed for 20 of 50 coils; the "planar" coils were horizontal bands above and below the plasma. The good match was a coincidence. And the data had been available all along: simsopt ships the Fourier coefficients of the seven

W7-X base coils in its package data [7]. Two more bugs surfaced later and changed published numbers: a sign error in the axis formula (calibration current 12.18 kA instead of 13.49 kA per turn), and Elmer choosing the current direction itself, so 25 mirrored coils ran backwards (1.38 T instead of 2.50 T on axis). The latter is now caught by an automated check.

Lesson – plausible is not correct. Agents are good at producing numbers that look right. Every geometry gets an independent invariant check (winding number around the axis, clearances), and every field or force is computed by a second, independent method before it goes into a report.

4 Restart: mesh directly from the data

With the real coil paths (7 base coils, 70 by symmetry) the expected physics appeared immediately: bean-shaped flux surfaces and the characteristic island chain (Fig. 2). A custom FFT voxel solver matched exact Biot-Savart to 1.3 % but destroyed the flux surfaces during field-line tracing, because a Cartesian grid breaks the device's symmetry – so field lines are always traced with Biot-Savart.

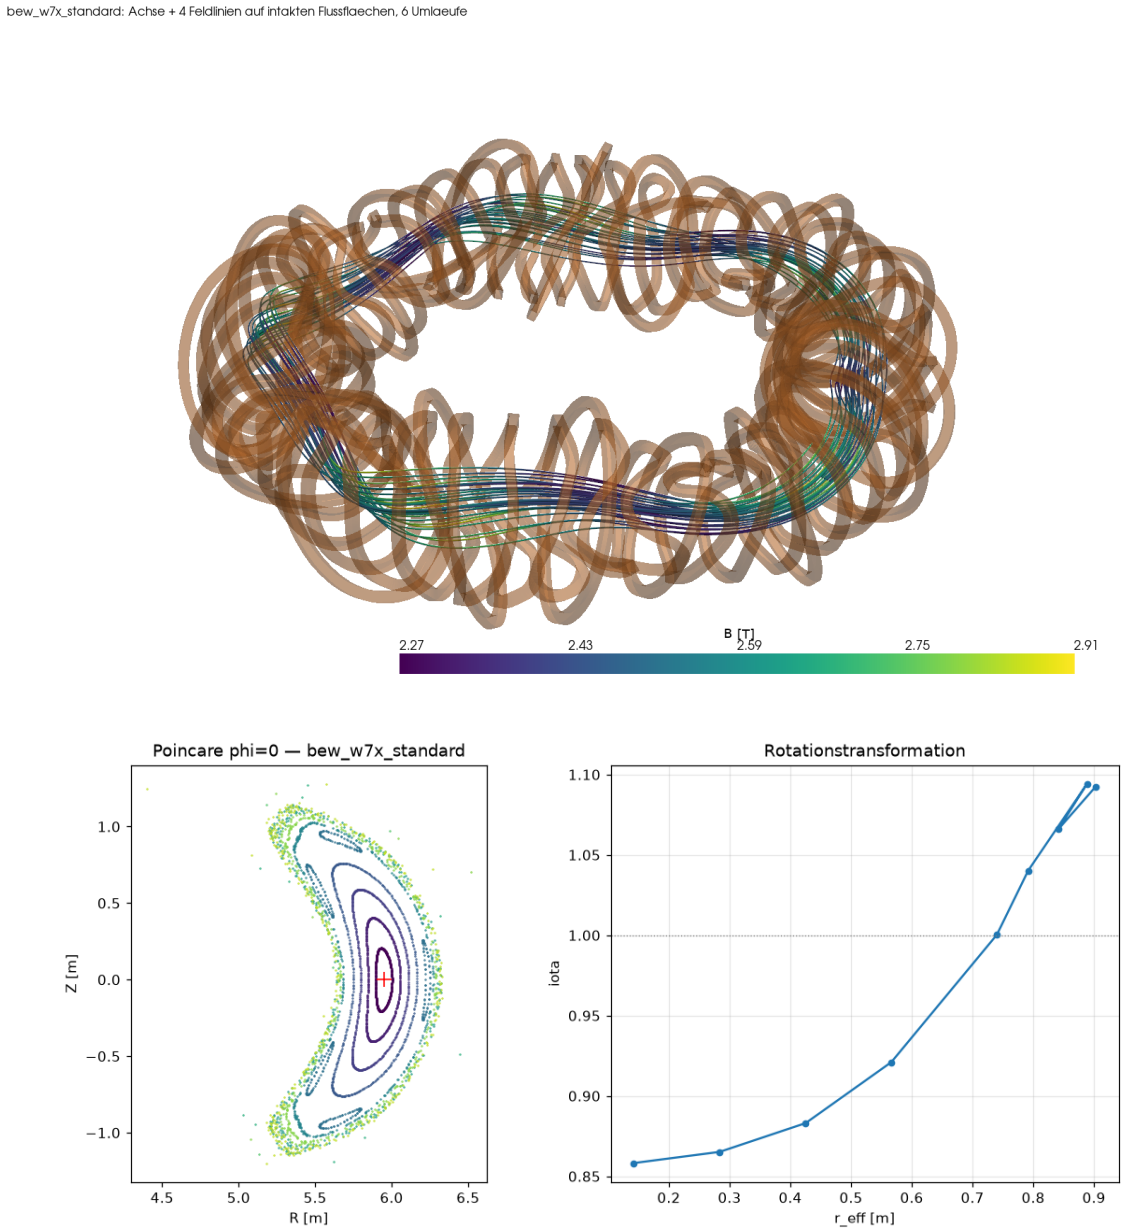

Figure 2: Field lines on flux surfaces (top), Poincaré section and rotational transform (bottom).

The plan had been to model coils in CAD. Generating the mesh directly was faster and more robust: a rectangle at every path point, two ruled half-surfaces per coil, meshed by gmsh. Sweeps and smooth lofts failed with self-intersecting surfaces. A clearance check on centre lines replaced boolean intersection tests (1.5 s instead of more than 45 min). Elmer with all 70 coils (3.3 M tetrahedra, 680 s, 10 GB) matched Biot-Savart on the axis to a median of 0.24 %.

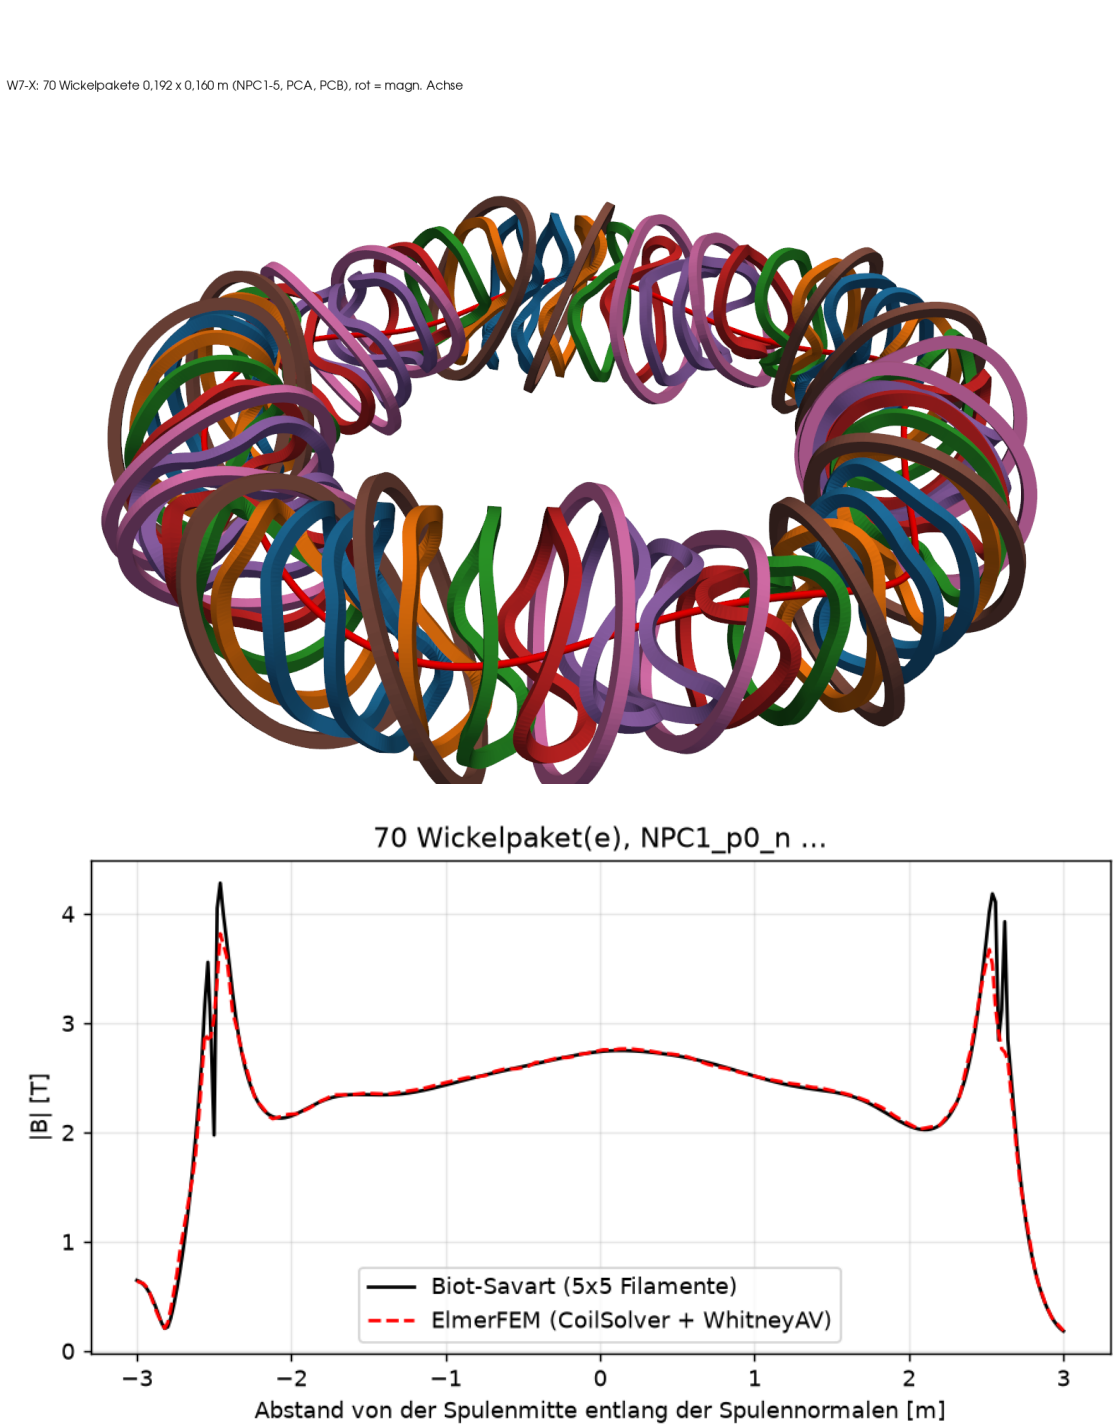

Figure 3: 70 winding packs meshed straight from the path data (top); FEM vs. Biot-Savart along a coil normal (bottom).

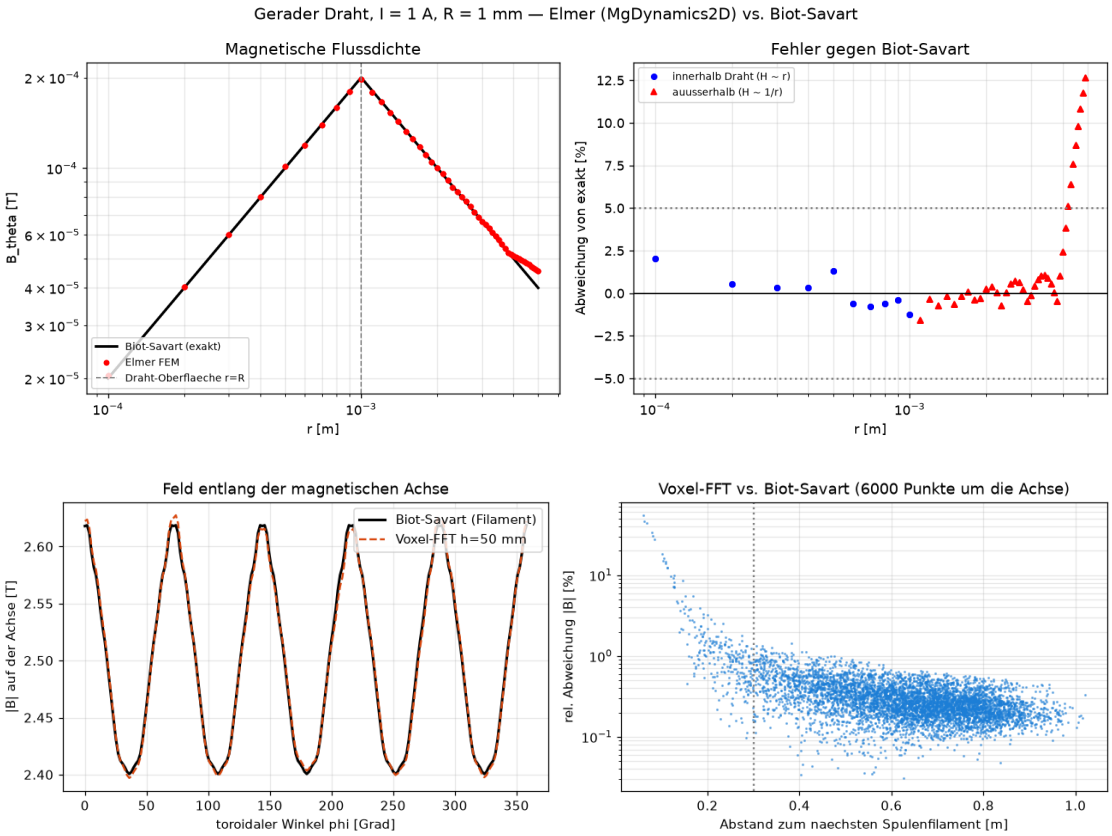

Figure 4: Validation: FEM on a straight wire (top), FFT voxel solver vs. Biot-Savart (bottom).

A side study ran the pipeline backwards from a target of 1 GW net electric power: plasma size and field → own plasma shape (DESC) → coils (simsopt). With high-temperature superconductors the device ends up at about 10 m major radius with forces of roughly 1 GN per coil – about a thousand times W7-X. From here on, the support structure is the real problem.

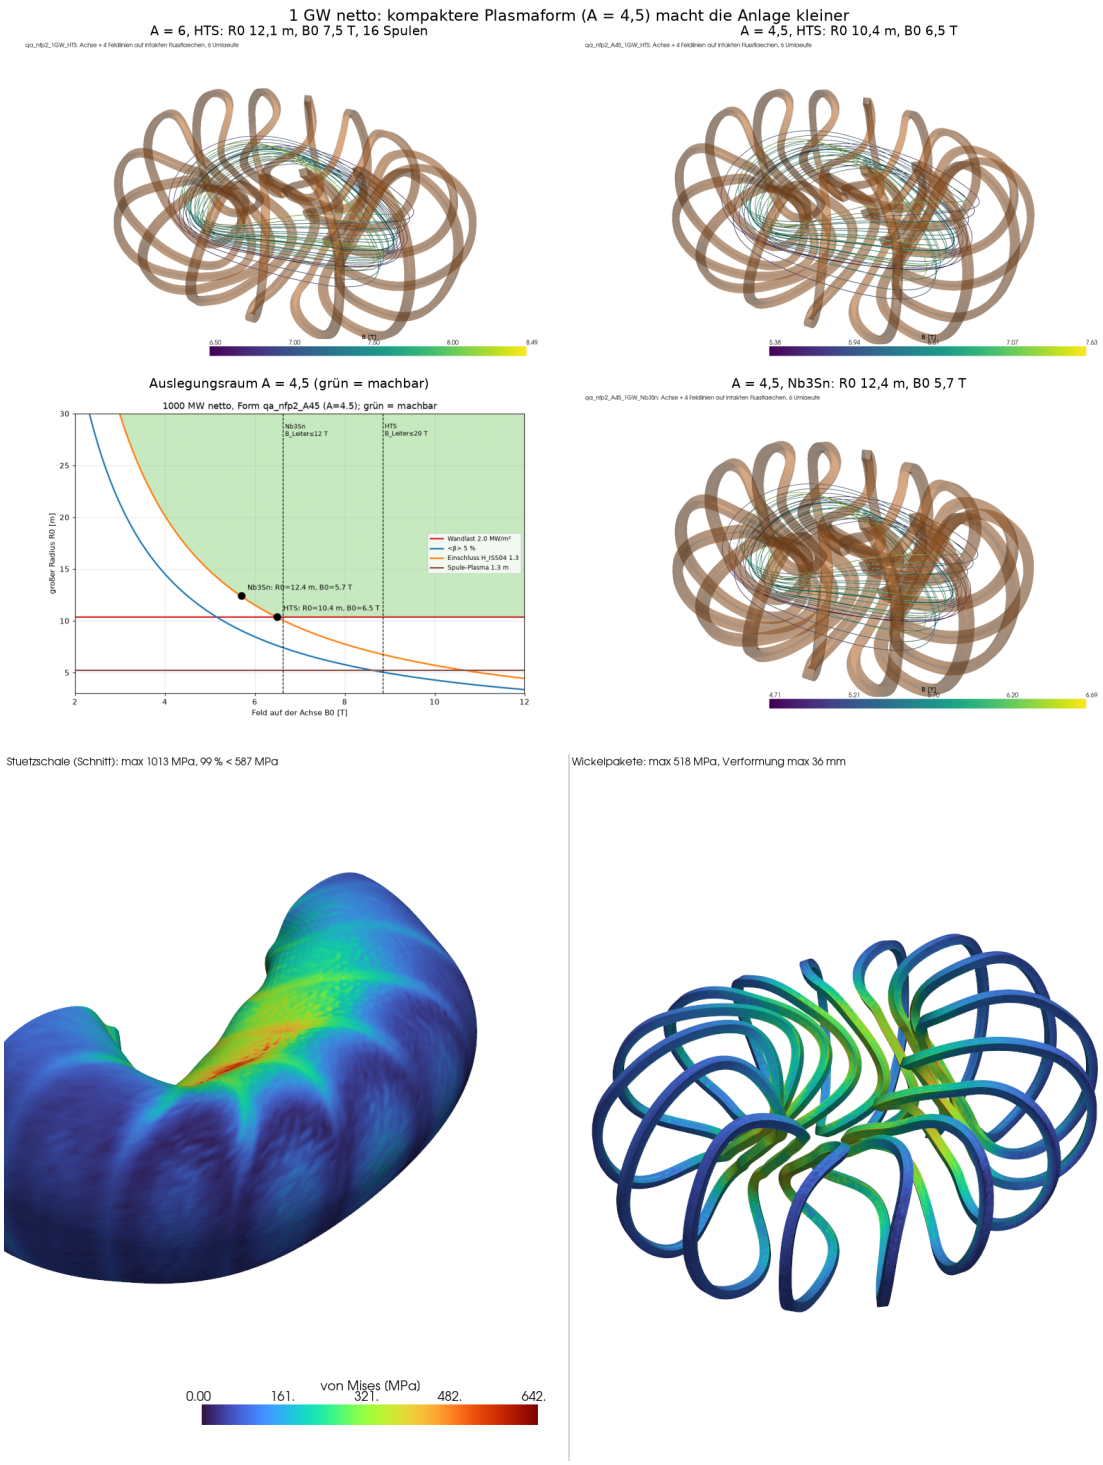

Figure 5: 1 GW sizing study (top) and a first support shell with coupled field → force → stress analysis (bottom).

5 The parametric CAD branch

A continuous support shell could not be constructed cleanly. The human proposed a technique from commercial CAD (Creo's *Variable Section Sweep*): per coil a closed guide curve as the casing's inner contour, sketches normal to it, a loft through them; all walls grow outwards. The agent implemented it in headless FreeCAD as a fully parametric model: spreadsheet parameters, sketch dimensions as expressions, symmetric coils as clones, connectors between neighbours.

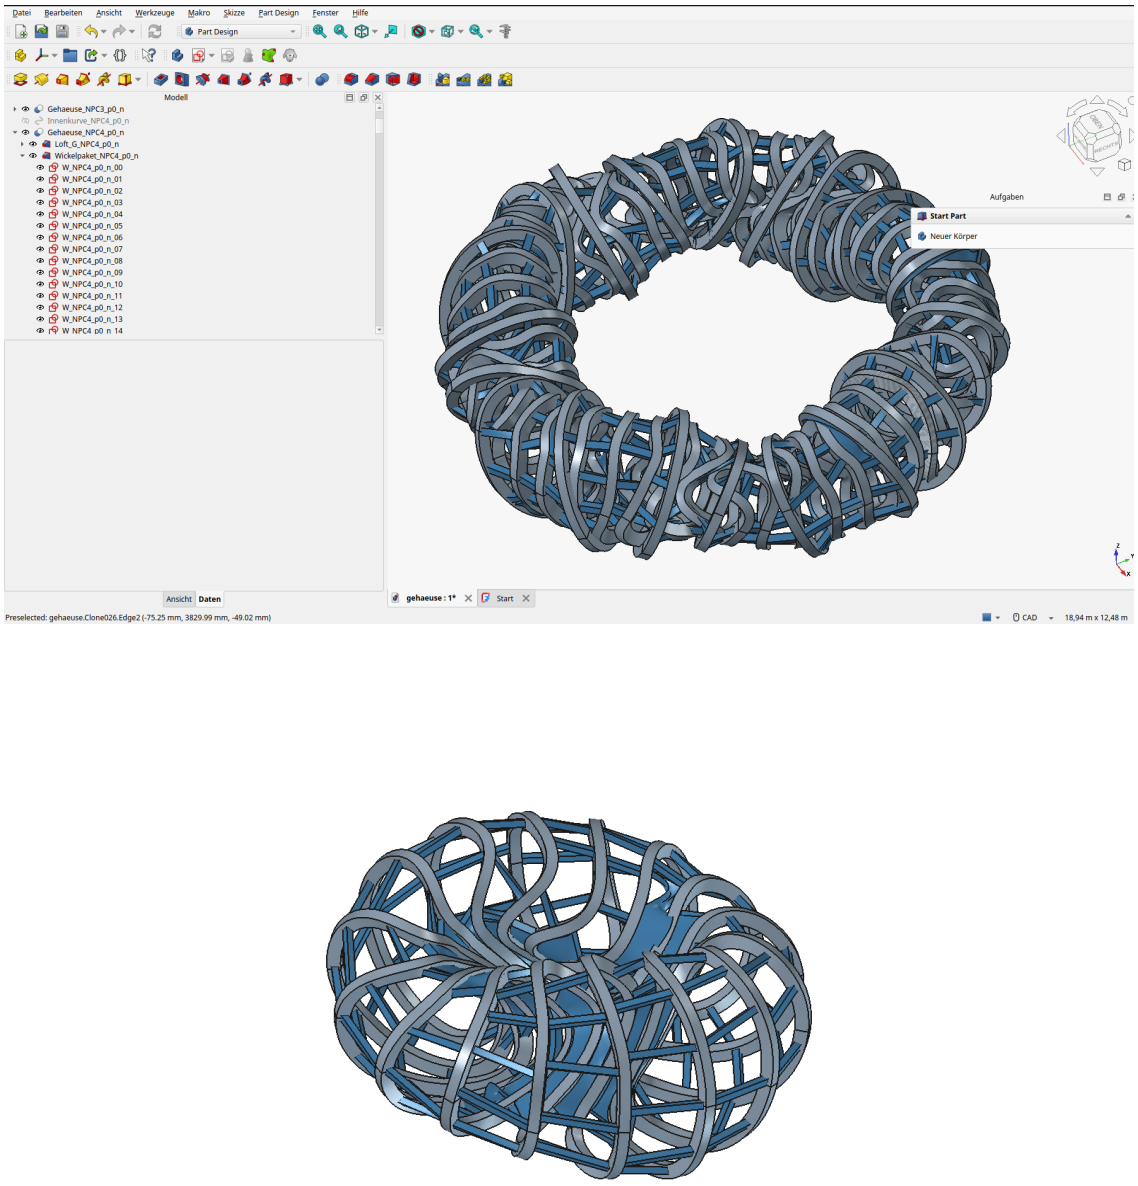

Figure 6: Parametric FreeCAD model: W7-X-like casings with connectors (top), compact 1 GW design (bottom).

What went wrong:

1. **Estimated, not designed, walls.** The first coupled run worked but gave stresses up to eight times the allowable value.
2. **Wedges break the mesher.** Inner wedges (after [4]) produced a valid CAD solid, but gmsh failed with "overlapping facets": several boolean cuts on the same winding pack created coincident faces. Six

attempts later (tolerances, fusing in FreeCAD, deliberate overlap; the last one hung gmsh for over an hour) the wedge was disabled.

3. **Collisions.** A loft bulges between sketch stations and corners of tilted neighbours overlap. An incremental assembly check (add coils one by one, count collisions, render a video) found about 100.
4. **Fix in CAD (human idea): subtract the neighbours.** Each casing minus its neighbours' envelopes plus a gap: 0 collisions in all 70 steps. Meshing the whole assembly stayed fragile.

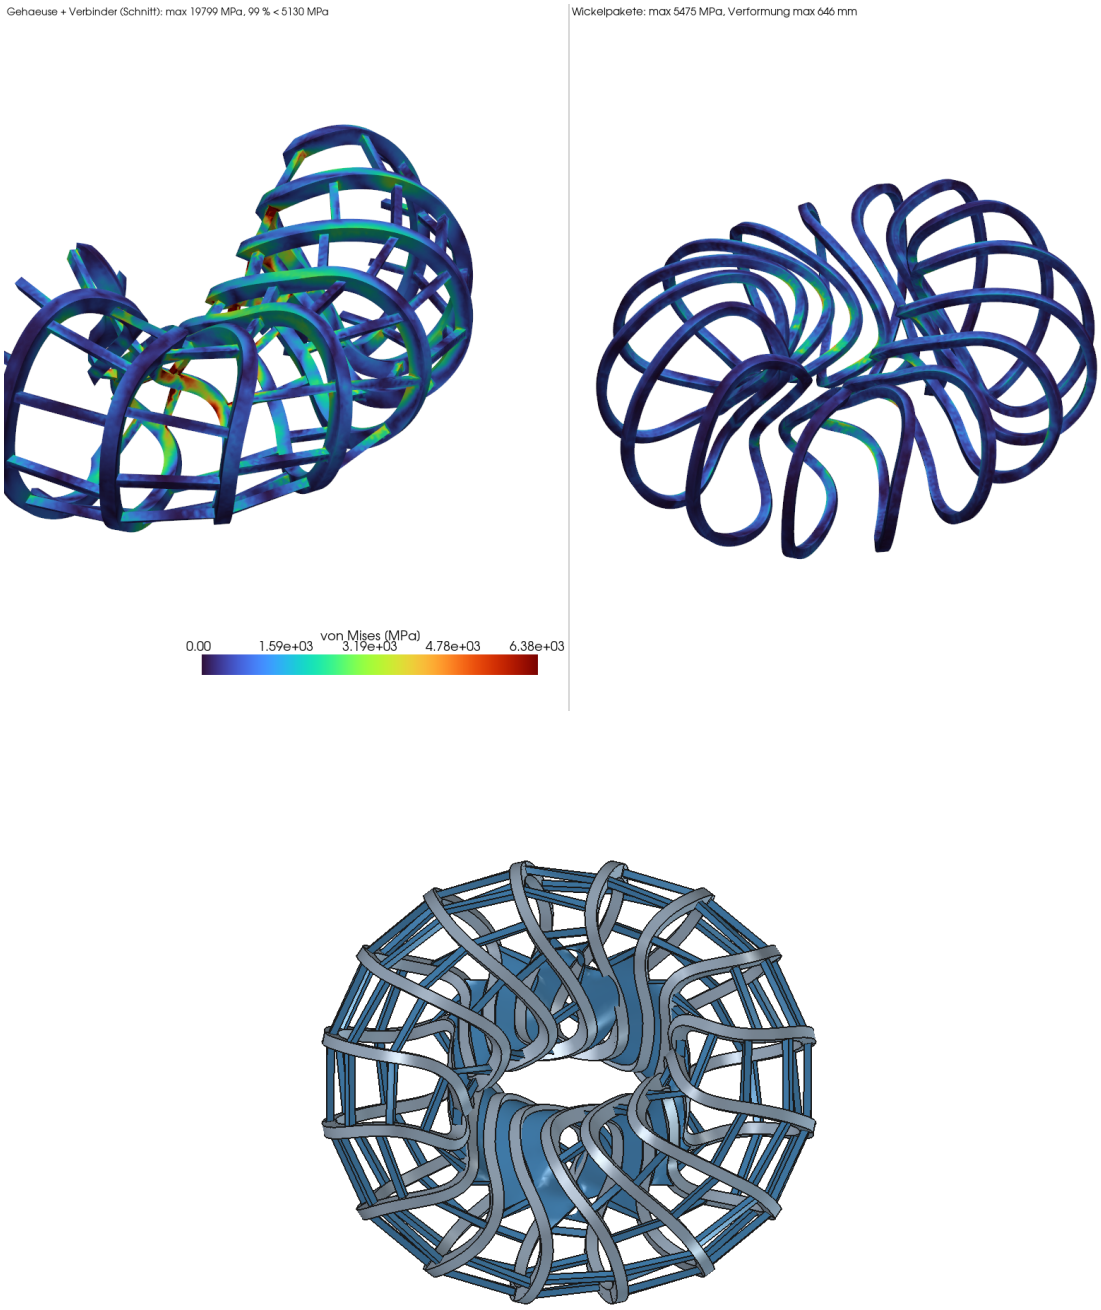

Figure 7: First stress run of the CAD casings (top); casings with inner wedges – valid in CAD, not meshable (bottom).

Lesson – valid is not meshable. The more boolean operations meet on the same faces, the more sliver faces appear, and the mesher fails far away from where the problem was created. Keep the analysis model separate from the design model.

6 The parallel branch: no CAD kernel

The human's suggestion: *"keep this morning's working approach for the analysis and run the CAD path in parallel as an add-on."* Each casing became two rectangular tubes on the exact coil path (outer contour and winding-pack cavity) as separate surface meshes – no boolean operations at all. Wall thicknesses follow from load and available space; a disentangling step tests surface points of foreign casings (including intermediate rings between stations) against the local cross-section and shrinks the cheapest wall: 4980 → 40 → 0 contact points.

Model	steel casings ($E = 200\,\mathrm{GPa}$), homogenised winding packs ($E = 100\,\mathrm{GPa}$), linear elastic
Load	Lorentz force $\vec{J} \times \vec{B}$ from the magnetostatic solution in the same run
Supports	one clamped patch per casing (simplification; no connectors, no support ring)
Mesh / run time	3.02 M tetrahedra in 78 s; solve 55 min, about 14 GB RAM
Force per coil	FEM 2.42 MN vs. Biot-Savart 2.39 MN (+1 %)
Casing von Mises	mean 11 MPa, 99 % below 43 MPa, max 110 MPa; deformation below 1 mm

Because of the simplified supports and walls as thin as 3 mm in places this is *not* a strength proof – but it is a reproducible end-to-end pipeline.

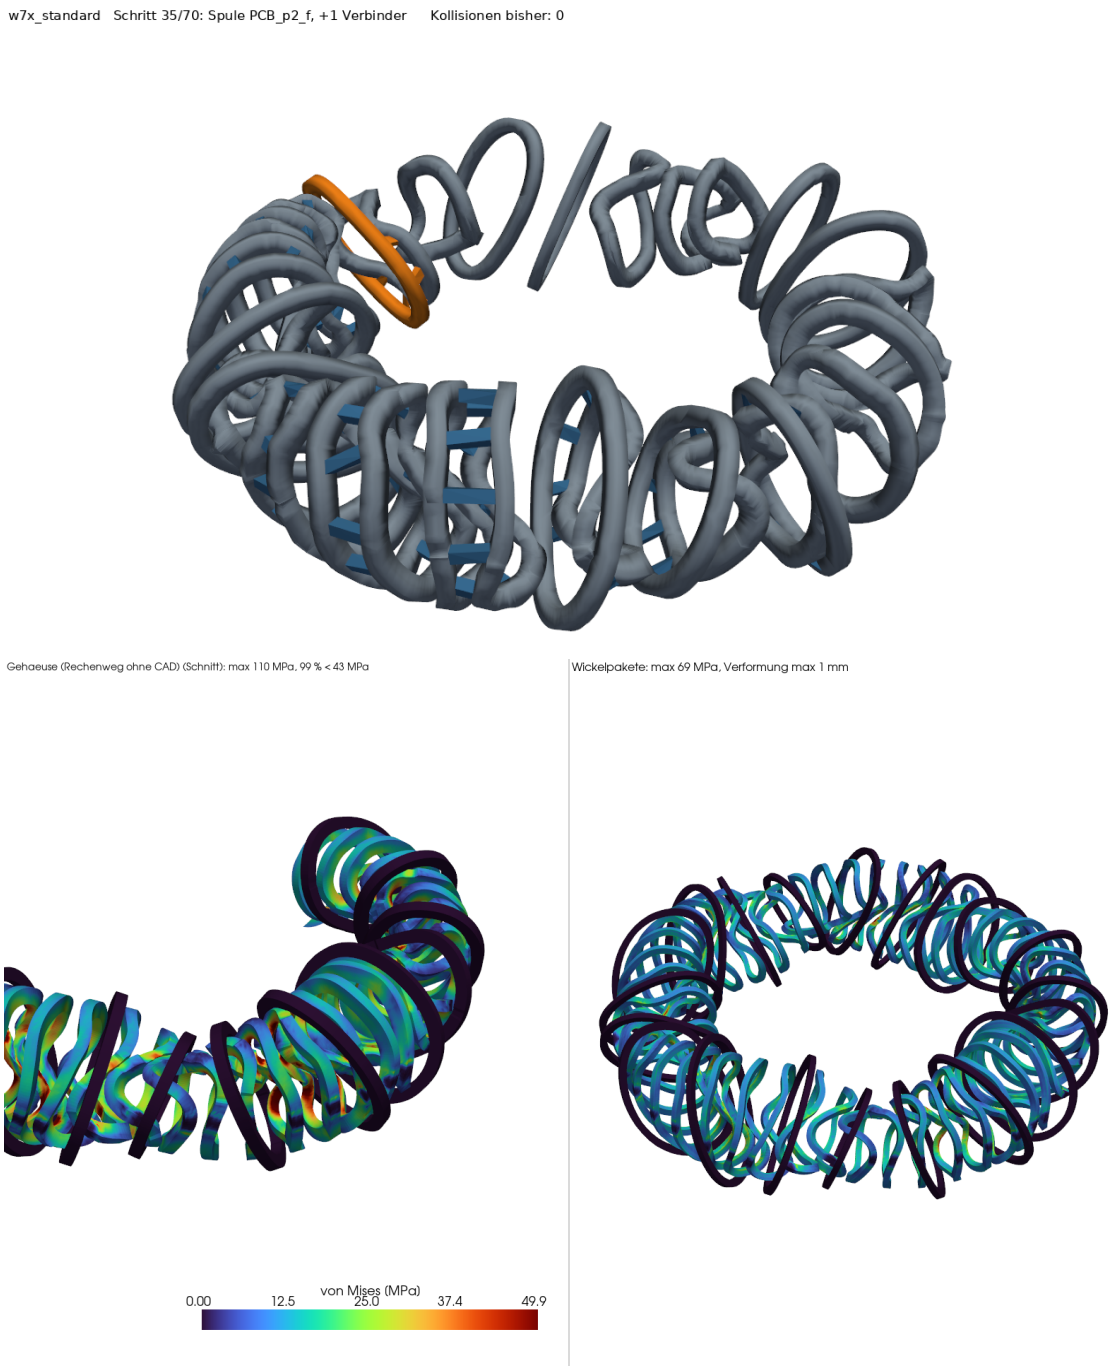

Figure 8: Incremental collision check in the CAD branch, step 35/70 (top); von Mises stress from the CAD-free branch (bottom).

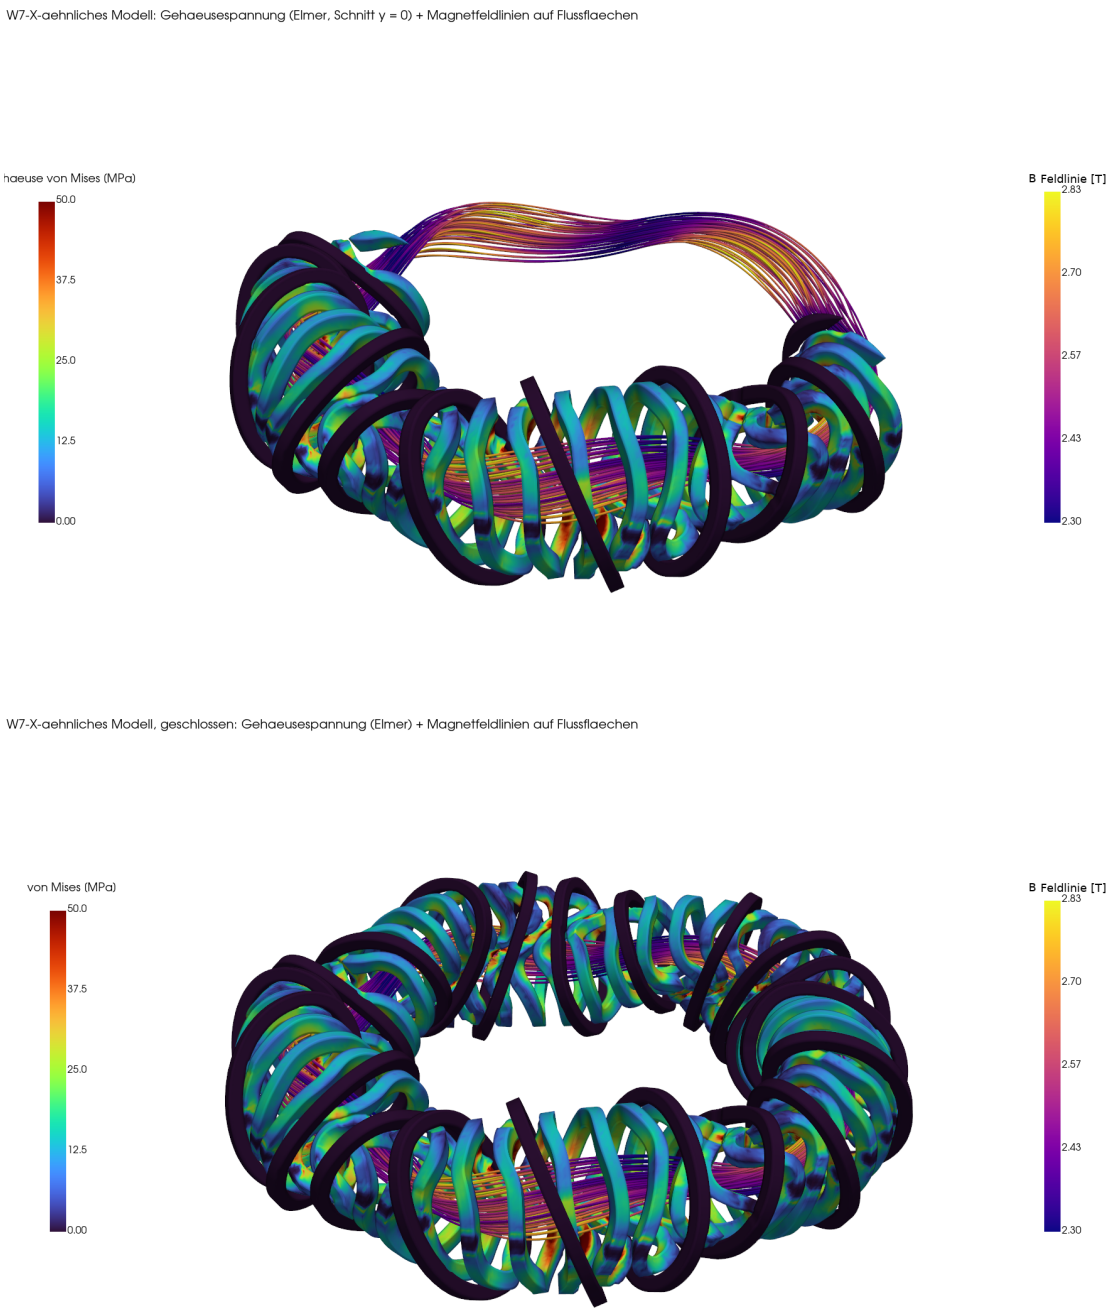

Figure 9: Casing stress with magnetic field lines on flux surfaces, cut open and closed.

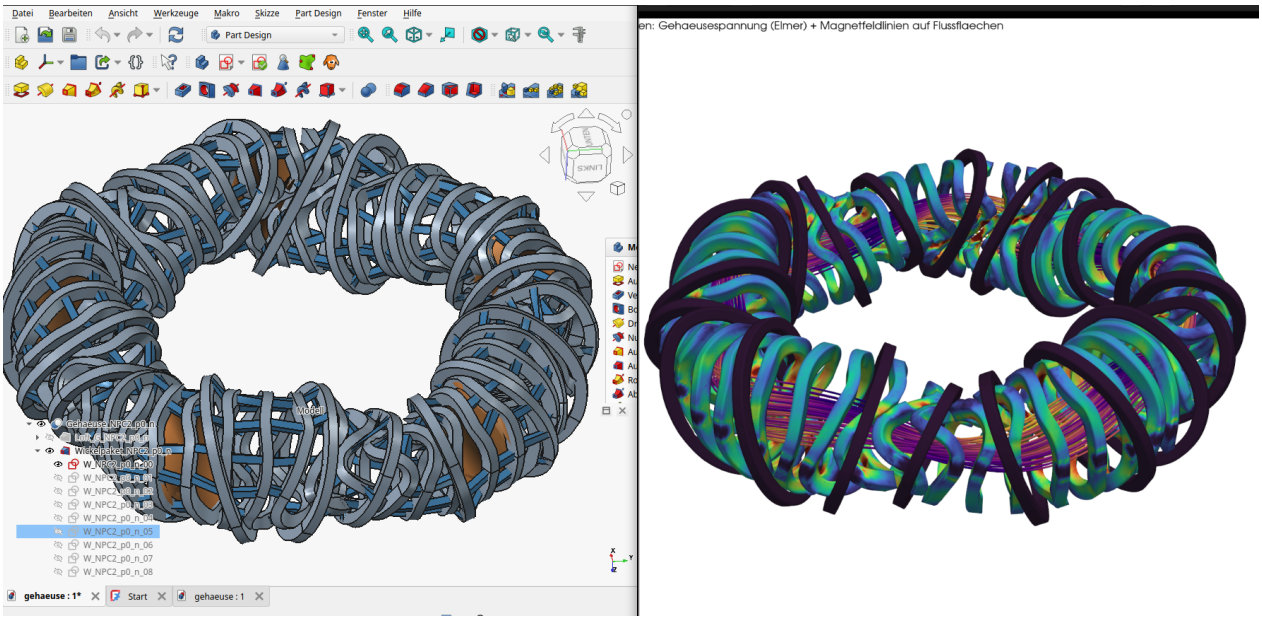

Figure 10: Both branches side by side: parametric CAD (left), stress analysis from the CAD-free pipeline (right).

7 Visualisation: from mesh to Blender

A 3D cut-away model shows each coil's layers: cable (NbTi strands in copper, helium), aluminium jacket, turn insulation, ground insulation, potting (an assumption), steel casing; $12 \times 9 = 108$ conductors per coil. A separate conductor model answers the question "where does the current flow?": almost entirely in the superconducting NbTi; copper only takes over briefly during a quench; jacket and helium carry none.

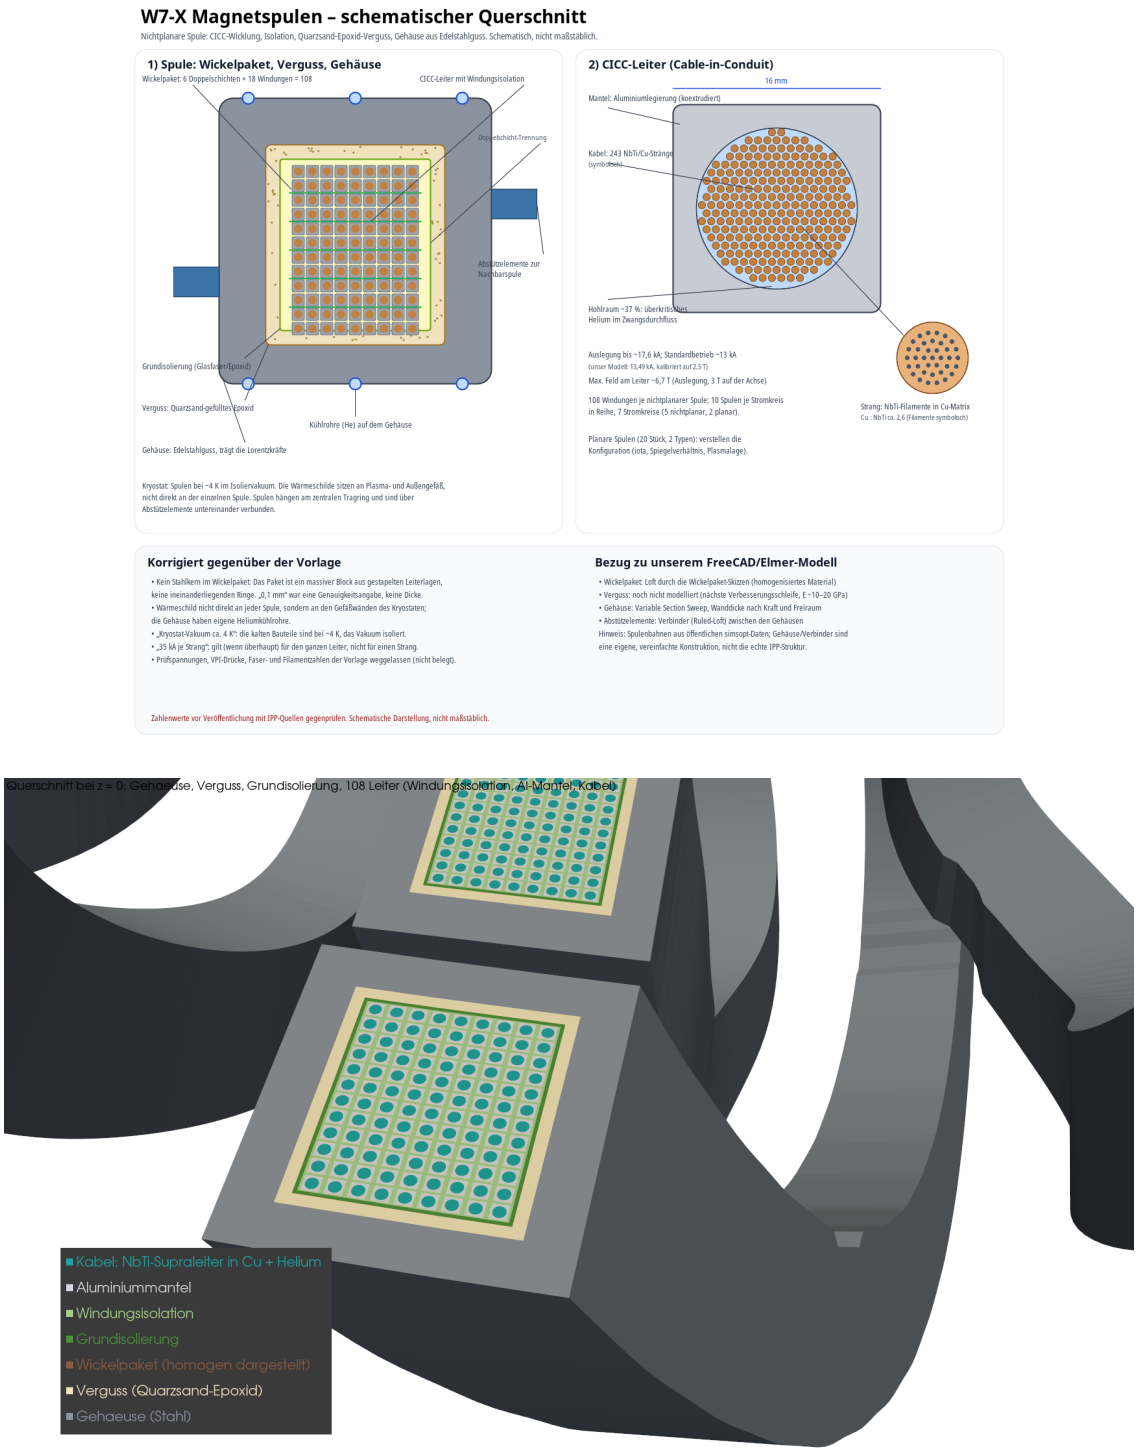

Figure 11: Coil layers (schematic, top) and in the 3D model (bottom).

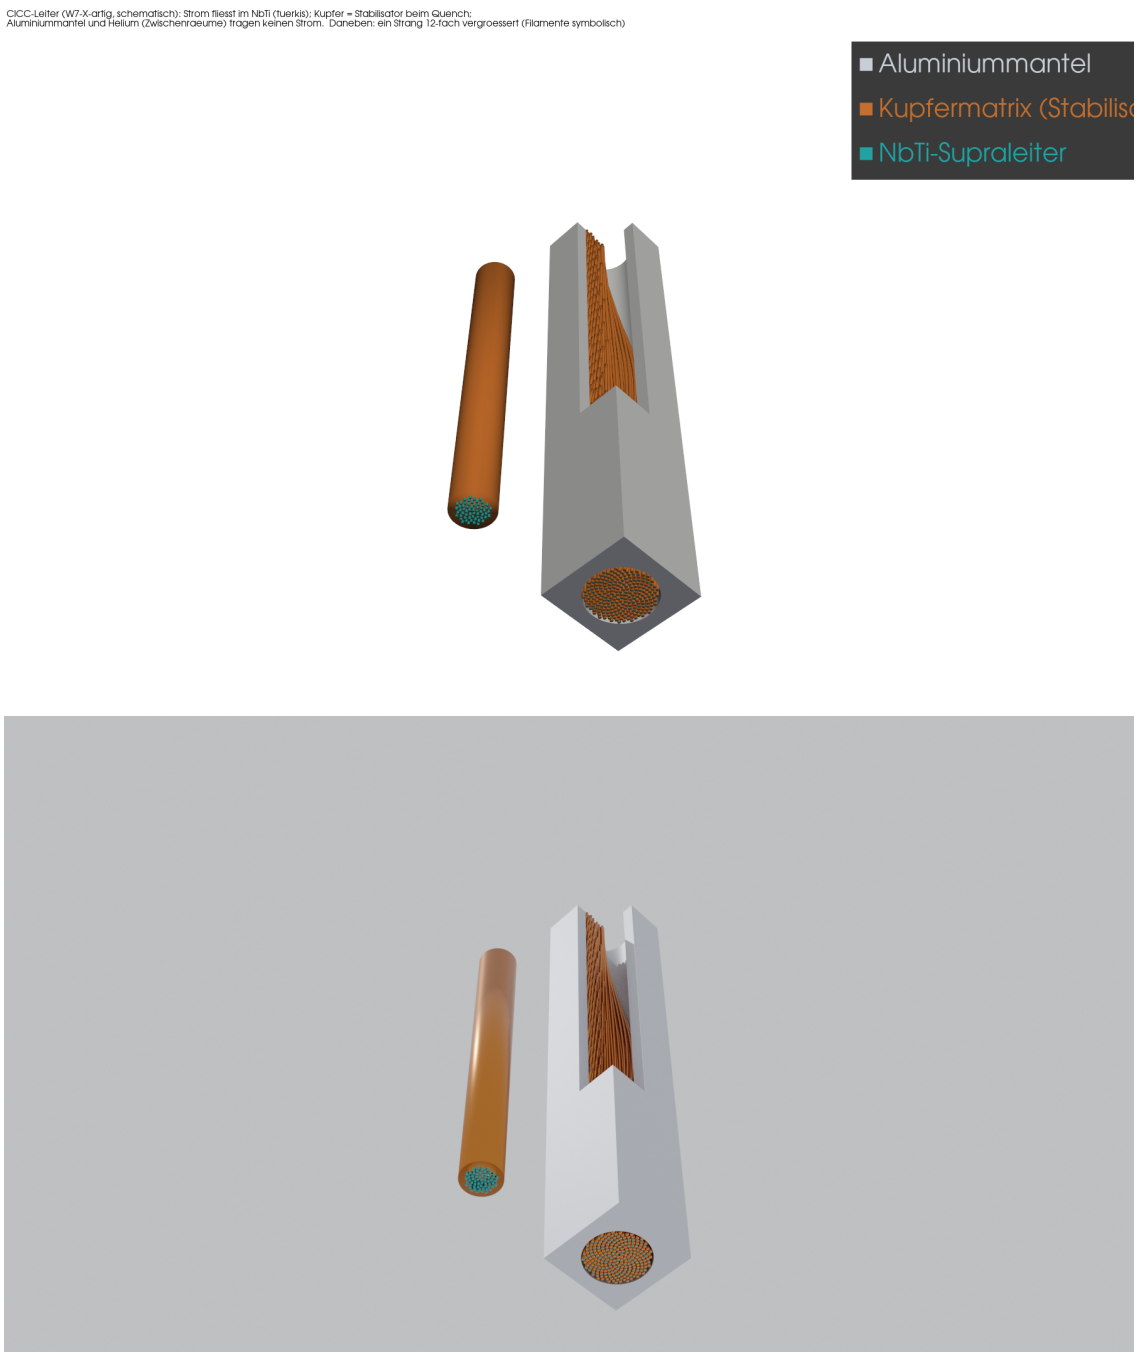

Figure 12: Conductor: opened jacket, twisted strands, one strand magnified with NbTi filaments (count symbolic). PyVista (top), Blender Cycles (bottom).

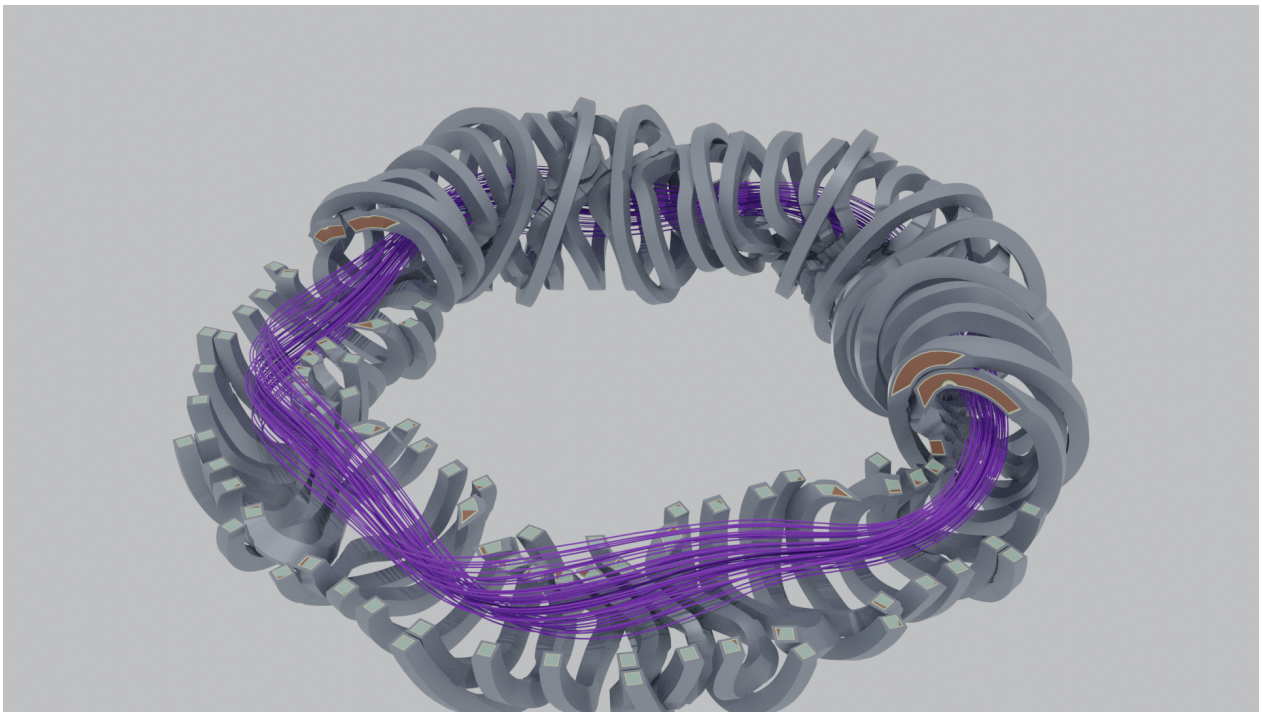

Figure 13: The full coil system with field lines, rendered in Blender.

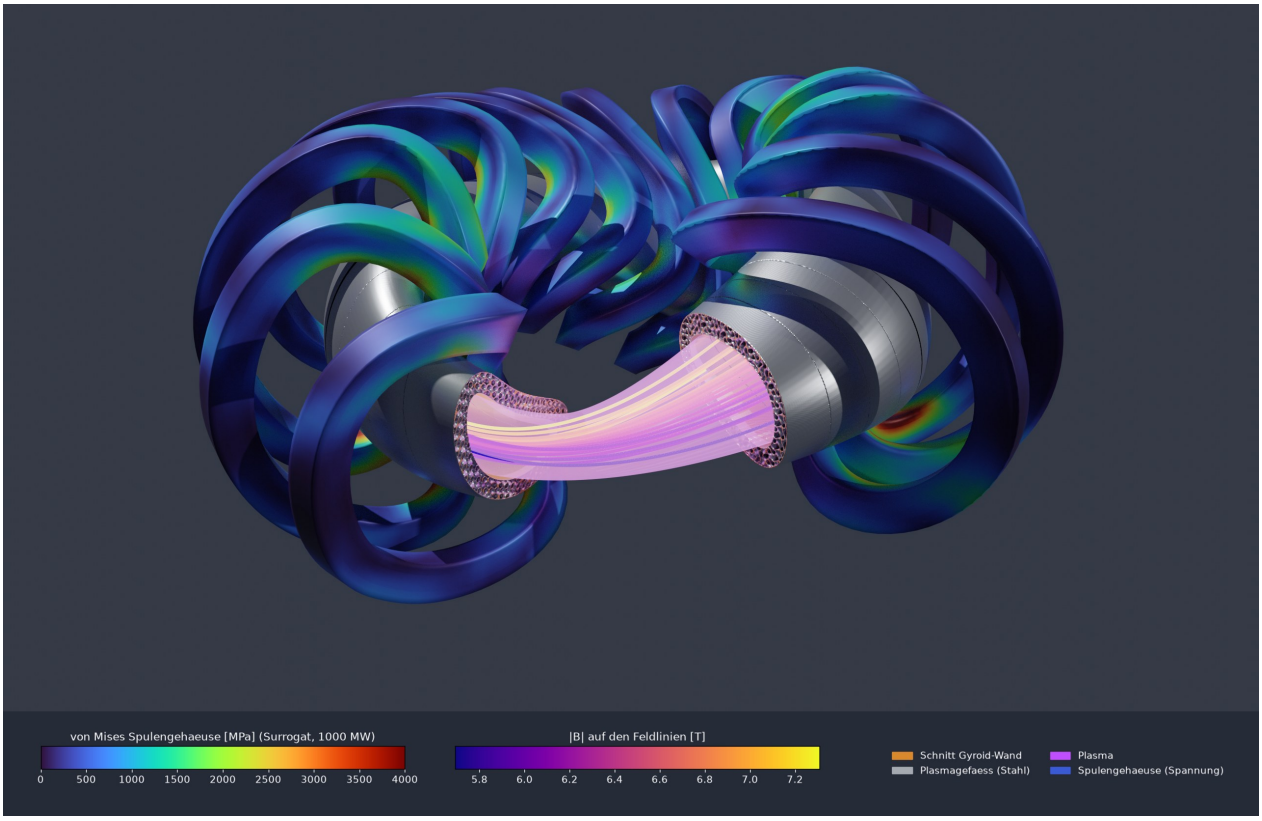

Figure 14: Later step: the 1 GW design with casing stress from the surrogate model (1000 MW), field lines, and a cut-open plasma vessel whose wall has a gyroid cooling core built with PicoGK. A rendering, not a cooling design.

Getting PyVista's glTF export into Blender needed three fixes, each found by measuring instead of guessing: colours arrived as sRGB values in a linear slot (washed out; fixed by linearising), the scene was rotated by 90° about the vertical axis (fixed at the root node), and the distribution's Blender build lacks the

denoiser (render fails; denoising disabled). Rendering runs on the CPU because the GPU is reserved for the local agent.

8 Takeaways for working with coding agents

- **Verification is the job.** The most expensive error looked plausible. Independent checks (invariants, second methods, regression tests from upstream projects) caught every major bug.
- **Silent failure is the default in scientific tools.** Wrong keywords ignored, current directions chosen by the solver, meshers hanging for an hour. Every pipeline stage now asserts on its output.
- **Prefer the boring path.** Direct mesh generation from data beat a CAD kernel with booleans. Keeping a working path alive while exploring the elegant one saved the project.
- **Humans supplied the decisive ideas** (research first, variable section sweep, neighbour subtraction, the parallel branch); agents supplied speed in implementing, testing and documenting them.
- **Files are the shared memory.** A findings log, a handover file and skills let two agents work on one project across sessions and context compactions.
- **References need checking** like any other agent output.

Open items

Potting layer between winding pack and casing; connectors and support ring instead of clamped patches; minimum wall thickness; wedge and casing in one body; wall sizing for the 1 GW design; HTS strain check; tolerance study (± 2 mm).

9 Follow-up: cooling the plasma vessel

Status 27 Oct 2026. Flow computed with cold water; heat only estimated roughly. The numbers show trends; they are not a design.

9.1 Task and layout

The plasma vessel wall of the 1 GW design is 0.9 m thick and has a core of gyroid filaments built with PicoGK [25]. A 10° segment (85°–95°) was studied. The assumed heat flux is 0.5 MW/m^2 on 26.7 m^2 of plasma-facing surface, i.e. 13.3 MW. With a 10 K temperature rise the segment needs 319 l/s of water.

The layout grew step by step from the user's instructions:

- Water flows *through the filaments*, not through holes in a perforated plate.
- The inlet is on the inboard side of the torus, the outlet outboard.
- The core sits between two inclined planes. In front of and behind it, the distribution and collection chambers taper off (Z-manifold). This shape came from a sketch the user drew on a screenshot.
- The gyroid is graded: a 130 mm layer with 100 mm cells on the plasma side, behind it a dense, coarse gyroid as load-bearing structure.

The flow was computed with FluidX3D [17], a lattice Boltzmann solver for GPUs [18, 19, 20, 21, 22, 23, 24]. The lattice has 8 mm spacing and 259 million cells. Each run covers 40 s of physical time with the LES turbulence model and takes about 15 min. FluidX3D has its own licence: free for non-commercial use only, source-available, not an open-source licence. The two modified files (`setup.cpp`, `defines.hpp`) are published as an "altered version" with the original licence in the repository folder `fluidx3d/`.

9.2 History

Variant	Change	Result
a, b	chamber with perforated plate; holes missed the plate (spotted by the user)	unstable, jets
f1	water through the filaments, open chambers	spread 141 %
f2	finer gyroid towards the plasma	spread 40 %
f4, f5	Z-manifold from the sketch, denser outer layer	spread 20 %, 7.9 kPa
f6–f9	thinner front layer, wall thickness by radius	dips at ±52°, cause searched
f10	lattice cells taken directly from PicoGK	dips gone, spread 36 %
f11	front wall 4.9 mm to 5.8 mm by radius	**spread 8 %**, 22 kPa, 288 l/s

Table 1: Variants of the cooling segment. Spread = standard deviation of the flow over 24 circumferential sectors divided by the mean.

9.3 An error in the lattice

In f6 to f9, two bands at ±52° of the circumference received almost no water. It looked like a design flaw. Three hypotheses were tested and rejected in turn: layer thickness, orientation of the layer relative to the gyroid, and inaccurate offset surfaces in PicoGK. Two intermediate findings were our own measurement errors.

The actual cause was the conversion of the water mesh into lattice cells. FluidX3D decides for each cell with a ray test whether it is inside. On nearly horizontal faces of the very fine gyroid mesh this failed: at $z = 1.906$ m a whole lattice layer became steel and blocked the band. It was found only by comparing the lattice with the mesh layer by layer. Since f10, PicoGK writes its voxel occupancy to a file that FluidX3D reads directly; it matches the mesh 100 %. The numbers of f6 to f9 are therefore distorted.

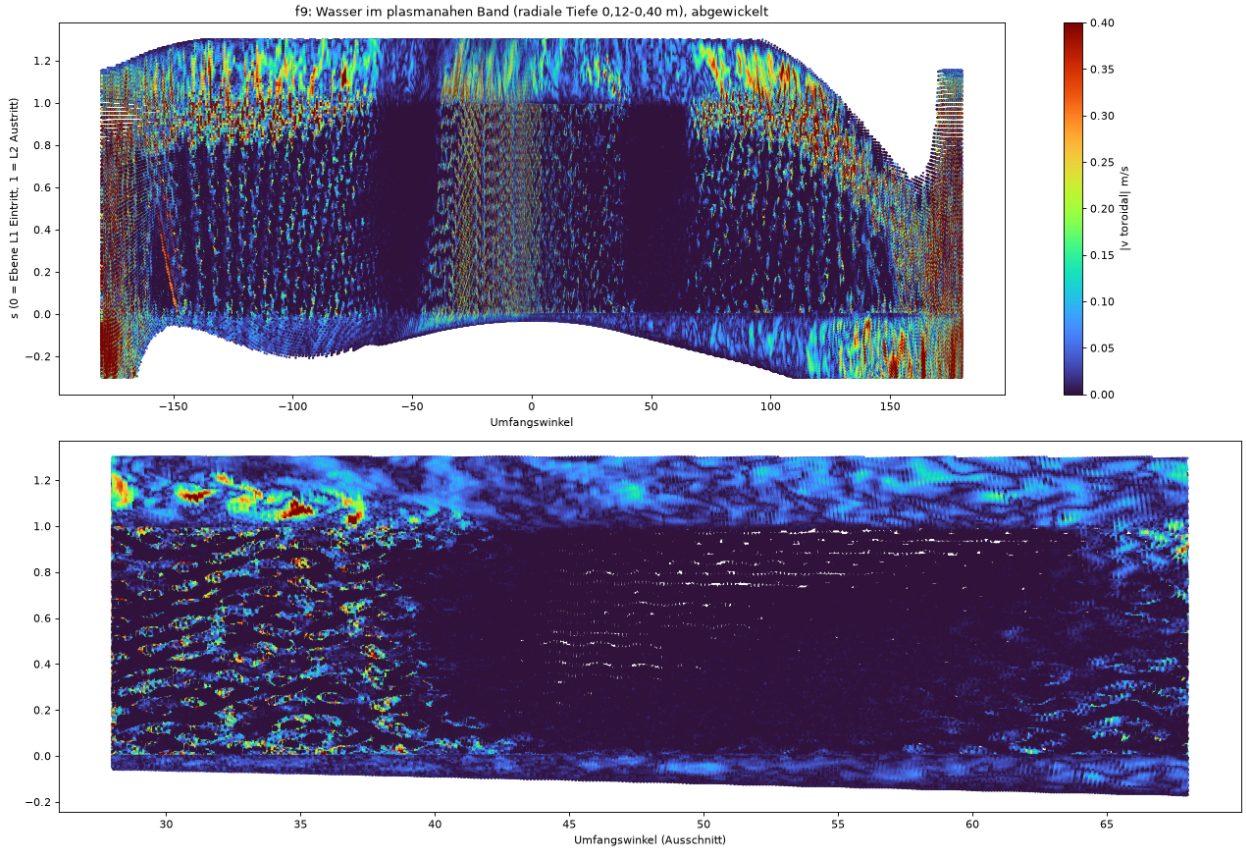

Figure 15: Lattice error (f9): water in the plasma-side layer, unrolled over the circumference (horizontal) and the path from inlet to outlet (vertical), colour = toroidal speed. The dark blocks at about $\pm 50^\circ$ carry no flow because one lattice layer there was wrongly set to steel. Zoom below. Labels in German.

After that, the distribution could be balanced in one step. The wall of the front filaments now varies with radius: 5.8 mm inboard, 4.9 mm top and bottom, 5.2 mm outboard. The spread dropped from 36 % to 8 %.

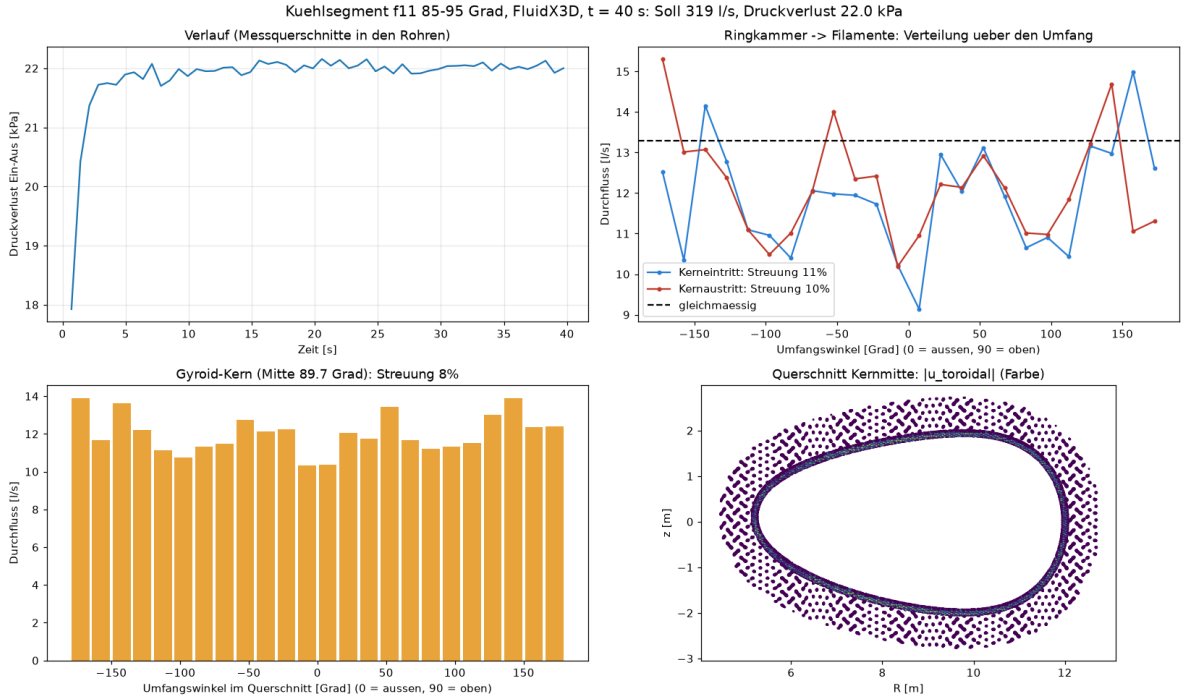

Figure 16: Evaluation of the FluidX3D run f11: pressure drop over time (steady after about 5 s), flow over the circumference at core inlet and outlet, flow through the gyroid core per sector (spread 8 %), and cross-section at mid core. Labels in German.

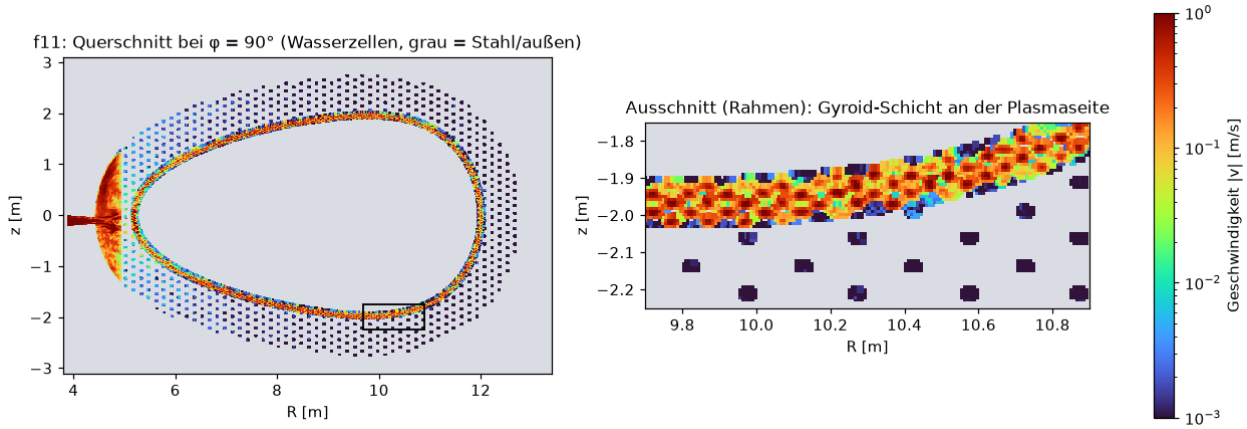

Figure 17: Cross-section through the f11 flow field at $\varphi = 90°$; every 8 mm lattice cell with water is a dot, colour = speed (log scale). Water enters on the left (inboard). Almost all the flow passes through the fine gyroid layer on the plasma side (zoom on the right, 0.1 m/s to 1 m/s). In the coarse dense gyroid behind it the water sits in closed pores; it only carries load. Labels in German.

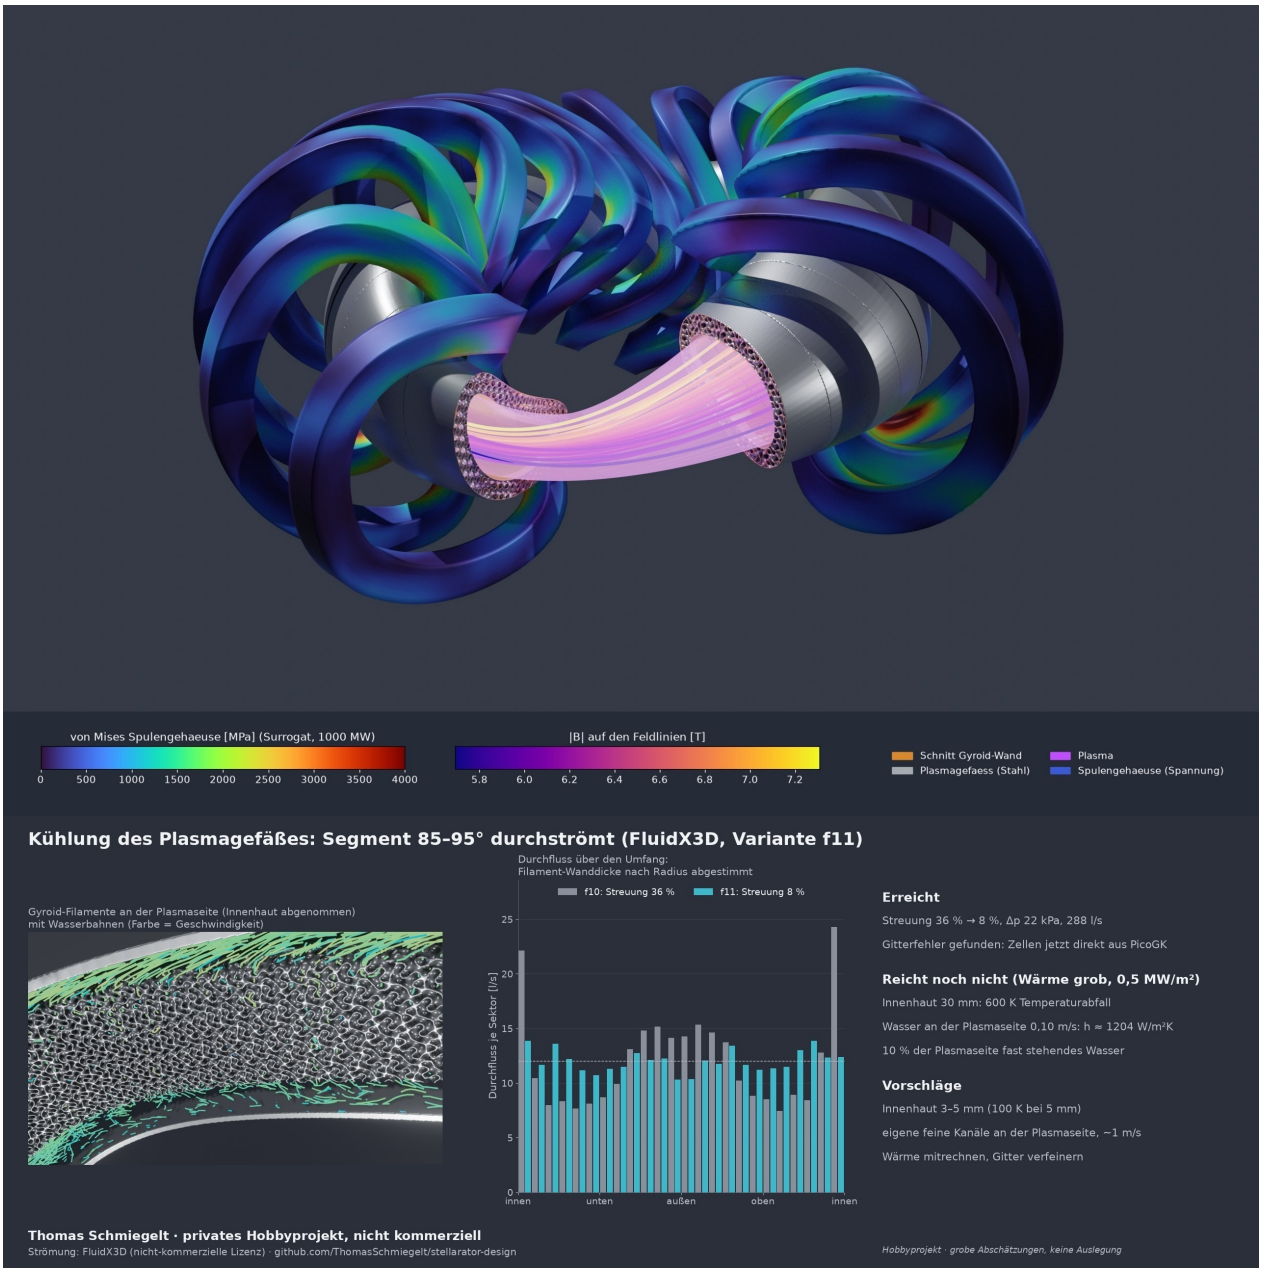

Figure 18: Overview of the 1 GW device extended by the cooling results: gyroid filaments on the plasma side with water paths (left), flow per circumferential sector before and after balancing (middle), findings (right). Labels in German.

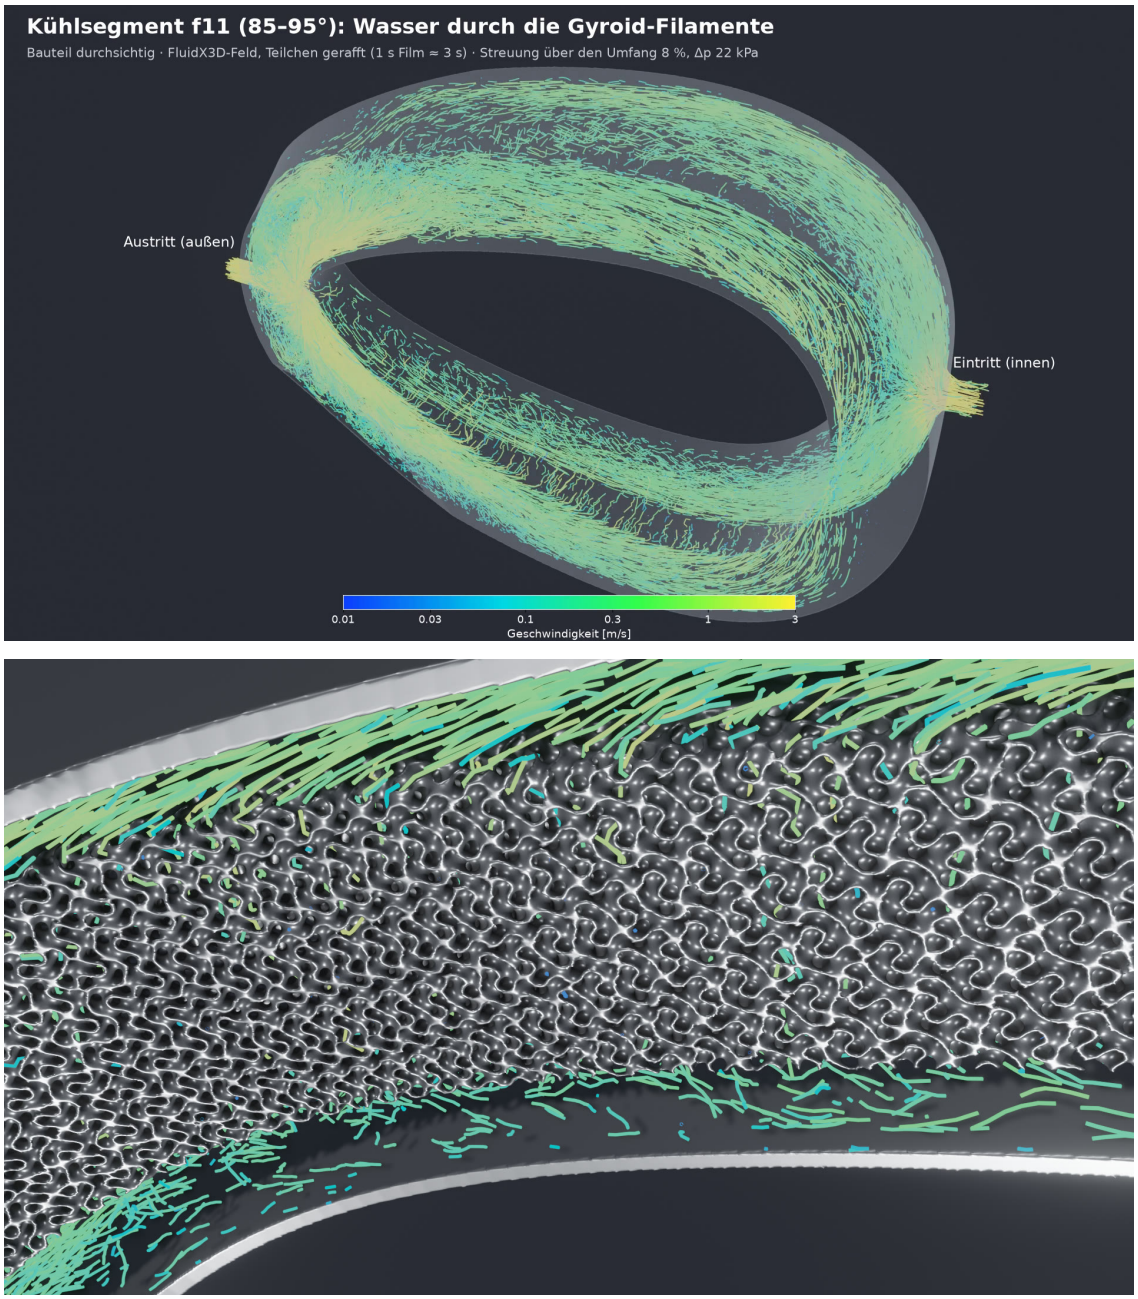

Figure 19: Frames from the flow films (Blender [14]): the segment rendered transparent, inlet inboard right, outlet outboard left (top); the lower half without the inner skin, looking at the plasma-side filaments (bottom). Colour = speed.

9.4 Is it enough? A rough heat estimate

The flow computation contains no heat. A one-dimensional estimate for pressurised water at 300 °C (Dittus–Boelter, steel with $k = 25\,\mathrm{W/(m\,K)}$) shows two bottlenecks.

Quantity at $0.5\,\mathrm{MW/m^2}$	Temperature drop
Conduction through the inner skin, 30 mm (as built)	600 K
Conduction through the inner skin, 5 mm	100 K
Transfer to the water, $v = 0.10\,\mathrm{m/s}$ (f11), without / with fin effect	415 K / 210 K
Transfer to the water, fine 10 mm channels at 1 m/s	48 K

Table 2: Rough heat estimate (`stellarator/kuehlung_waerme.py`).

The distribution is solved, the cooling is not yet. The inner skin is far too thick and the water on the

plasma side is too slow. In addition, on 10 % of the plasma side the water almost stands still (below 1 mm/s); hot spots would form there.

9.5 Suggested improvements

1. **Thinner inner skin:** 3 mm to 5 mm instead of 30 mm, possibly with tungsten armour. The gyroid core behind it carries the load.
2. **Faster flow on the plasma side:** either a thinner front layer (30 mm to 40 mm instead of 130 mm) with a finer cell, or a separate first-wall circuit with parallel channels at about 1 m/s. The gyroid behind then only carries load and removes volumetric heat. The pressure drop rises; at 22 kPa it is still small today.
3. **Remove stagnant water:** the near-zero speeds on the plasma side are probably at the end walls and chamber corners. They must be located in the field; guide vanes or a more open cell there are obvious options.
4. **Balance less sensitively:** one millimetre of wall thickness changes the flow by roughly a factor of 3.7. With typical manufacturing tolerances of ±0.1 mm that is about ±13 %. Defined orifices at the chambers, or a larger cell with thicker walls, are more robust.
5. **Compute the heat:** coupled flow and conduction (FluidX3D has a temperature extension, or Elmer for the steel). Only then can the fin effect of the filaments be judged.
6. **Check the lattice:** the front filament walls are about 10 mm thick and the lattice is 8 mm. A grid study at 4 mm to 6 mm is needed before the numbers are reliable.
7. **Reach the target flow:** 288 l/s instead of 319 l/s. The cause is the inlet boundary condition, not the geometry.

9.6 Collaboration

Here too, the decisive design ideas came from the user: water through the filaments, inlet inboard, inclined chamber boundaries. To agree on shape, a sketch on a screenshot worked best. This led to a small tool for drawing directly on the 3D model and handing the lines to the AI as 3D points [26]. The lattice error was found by the assistant, but only after several wrong guesses. The lesson: investigate an odd result before changing the design. The workflow is described in `docs/kuehlsegment_prozess.md` (German).

10 Follow-up: double filament, clearance and liquid lead

Of the improvements suggested in Section 9, the second one was taken up: the plasma side gets its own cooling circuit with fine channels. The requirements came from the user: the water should not get hotter than 150 °C, with an inlet of about 60 °C. On top of that came the idea of building the wall as a *double filament*: an inner part for cooling and, separate from it, an outer part for stiffness. All of it without colliding with the coils. All numbers in this section are rough estimates.

10.1 Wall layout

From the plasma side outwards: inner skin 5 mm, cooling channels, back wall 5 mm, dry gyroid for stiffness, outer skin 40 mm. The channels run toroidally across the segment from the inlet to the outlet plenum. The gyroid no longer carries water, so it may be more open and no longer has to be tuned to the cooling network.

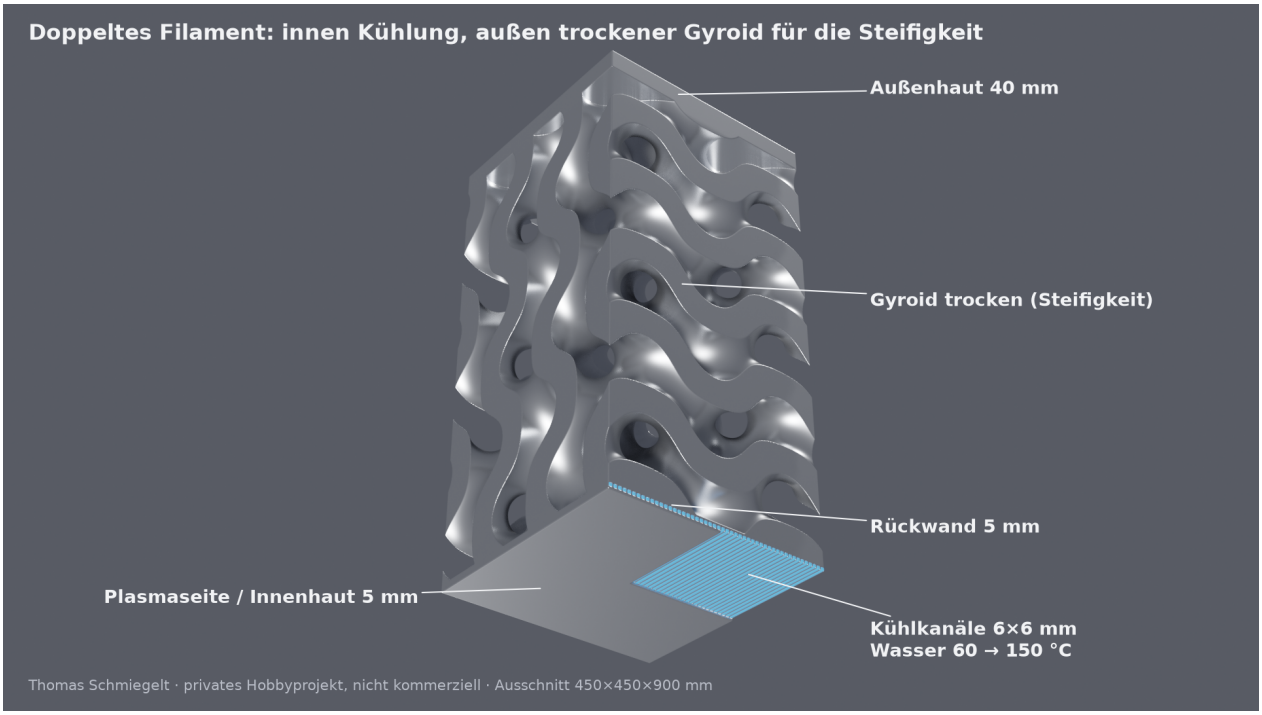

Figure 20: Wall piece with double filament (PicoGK, Blender), cut-out 450 × 450 × 900 mm. A window is cut into the inner skin; the cooling channels (blue) lie below. Labels in German.

10.2 Water at 60 to 150 °C

At 0.5 MW/m^2 the segment (85°–95°) takes up about 12.6 MW on the plasma side. With a temperature rise of 90 K that is only 35 l/s. Spread over all channels in parallel the water would be too slow and would boil at the wall. Therefore three neighbouring channels are connected in series (U-turn at the end face). Method: heat transfer after Gnielinski, friction after Petukhov, temperature in the cross-section with Elmer for the worst channel (`stellarator/erstwand_kanal.py`).

Channels	6 × 6 mm, rib 3 mm, 616 strands with 3 passes each
Velocity	0.96 m/s to 2.15 m/s
Pressure drop	49 kPa
Steel maximum	300 °C
Channel wall maximum	202 °C (boiling from 212 °C at 2 MPa)

Table 3: First-wall channels with water 60 °C → 150 °C at 2 MPa (variant n5).

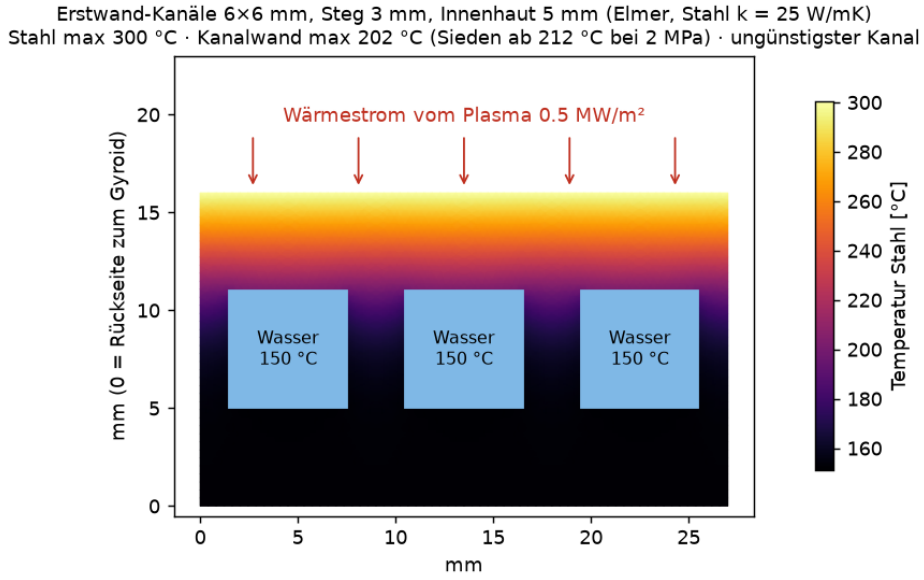

Figure 21: Steel temperature around the channels (Elmer, half pitch mirrored). Labels in German.

So high pressures are not needed at 150 °C; 2 MPa leaves enough margin to boiling. The real drawback lies elsewhere: heat at 150 °C can hardly be turned into electricity.

10.3 Clearance to the coils

For every point of the plasma boundary the smallest distance to the coil casings and connectors of the surrogate model was computed (`stellarator/gefaess_abstand.py`). The minimum is 1.25 m to the casings and 1.70 m to the connectors; in the computed segment it is at least 1.88 m. Including the gap the wall needs 1.05 m; with 5 cm margin it fits everywhere.

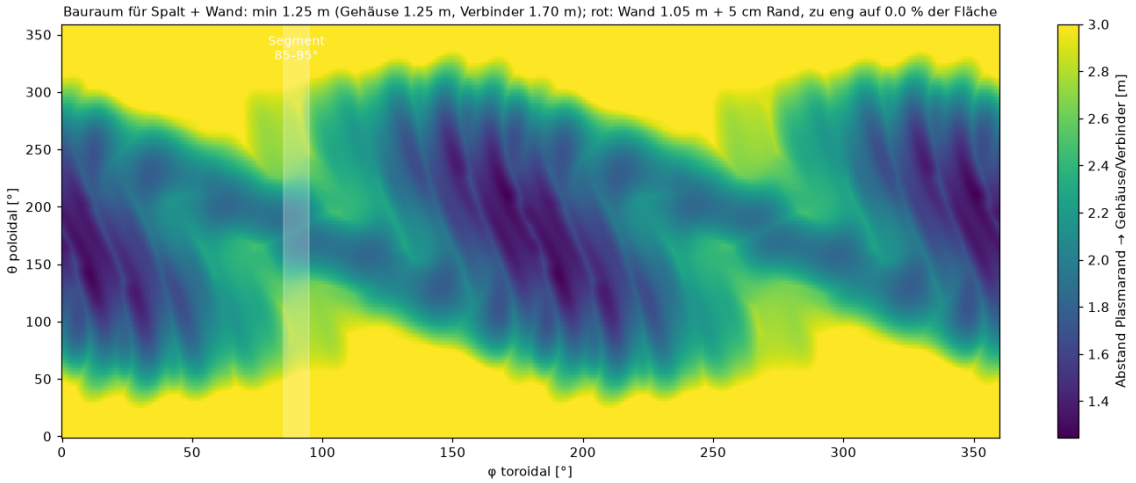

Figure 22: Distance from the plasma boundary to casings and connectors, unrolled over the torus. Red: limit at which the wall would no longer fit; it is reached nowhere. Labels in German.

10.4 Liquid lead instead of water?

As a next step the user proposed a *gyroid within a gyroid*: cooled near the plasma with liquid lead, which passes its heat on to the water. A gyroid wall does split space into two separate, intertwined networks and would make a compact heat exchanger. Lead needs no pressure, and lead-lithium would also breed tritium. Three points speak against putting the exchanger into the wall, however:

1. Lead melts at 327 °C, lead-lithium at about 235 °C. Next to water at 60 °C to 150 °C it would freeze.

2. A leak between hot lead and water right at the plasma is dangerous and cannot be repaired.
3. In the magnetic field, induced currents brake the metal as soon as it flows across the field.

For the third point the pressure drop was roughly estimated (`stellarator/blei_kanal.py`): $\mathrm{d}p/\mathrm{d}x = \sigma\, v\, B_\perp^2\, k_p$ with $k_p = c/(1+c)$ for conducting steel walls (c wall conductance ratio) and $k_p \approx 1/\mathrm{Ha}$ for an insulating layer. The field of a stellarator is mainly toroidal, so toroidal channels run almost along the field; the pitch of the field lines was estimated at about 6°.

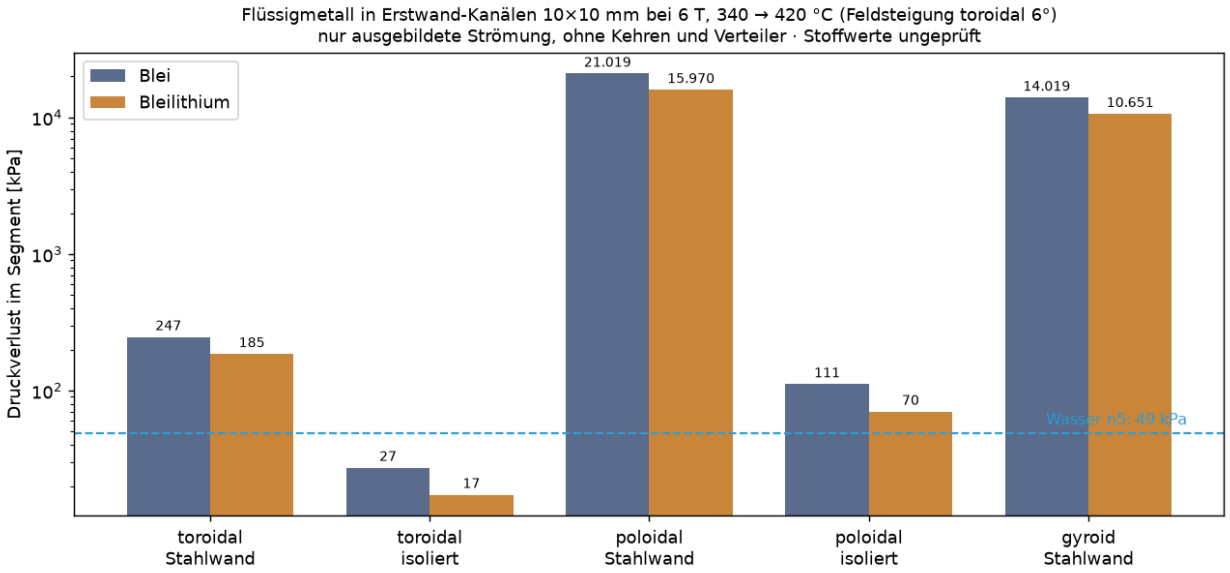

Figure 23: Pressure drop in the segment for lead and lead-lithium per channel direction at 6 T, logarithmic. Water (n5) for comparison. Labels in German.

Channel direction	Pressure drop	Pumping power / heat	
toroidal, steel wall	247 kPa	0.2 %	suitable
toroidal, insulated	27 kPa	0.02 %	suitable
poloidal, steel wall	21 MPa	17 %	unsuitable
gyroid, steel wall	14 MPa	11 %	unsuitable

Table 4: Lead in 10×10 mm channels, 340 °C → 420 °C, 0.87 m/s, 103 l/s per segment (variant b5).

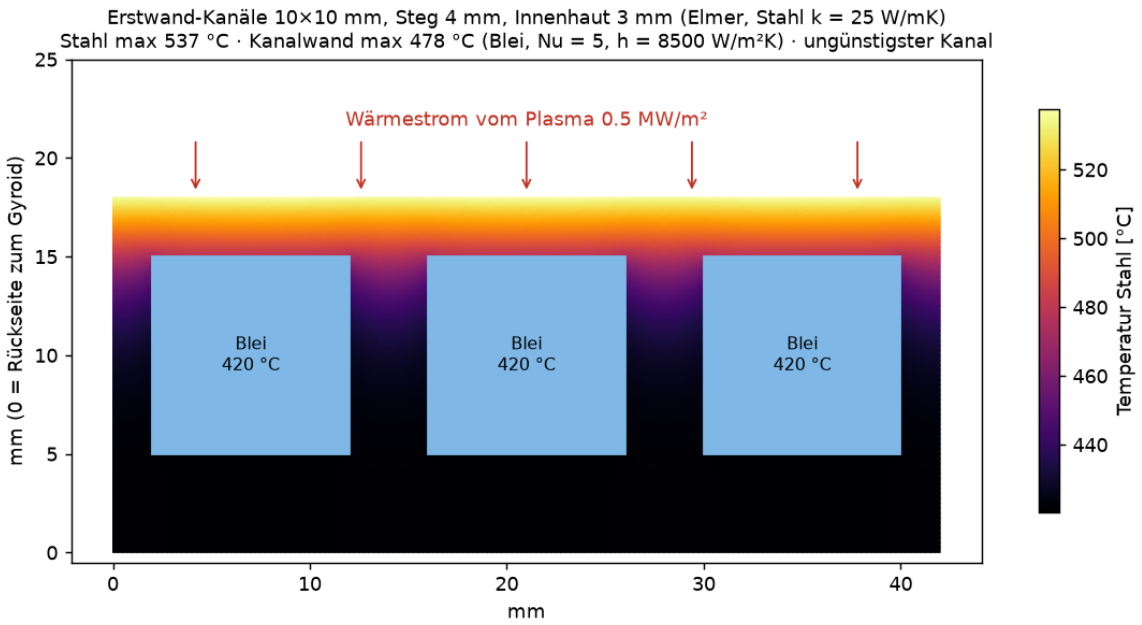

Figure 24: Steel temperature with lead (b5, inner skin 3 mm): at most 537 °C. With a 500 °C outlet and 20 mm channels it would be 711 °C. Labels in German.

Result: **the gyroid is not suitable for liquid metal**, because it keeps turning the flow across the field. Toroidal channels, on the other hand, work well. A sensible layout: lead or lead-lithium in toroidal channels on the plasma side, the dry gyroid behind it for stiffness, and a steam generator outside the coils that passes the heat to the water. At about 420 °C the heat can also be converted to electricity much better.

Not included are the additional losses in bends, manifolds and changes of cross-section (often the largest share), corrosion of steel in lead, volumetric heat from neutrons, and a more accurate heat transfer (conservatively Nu = 5 here). The material properties have not been checked against sources.

10.5 Collaboration

The ideas in this section came from the user: the temperature limit, the double filament, the check against the coils and the gyroid within a gyroid with lead. The assistant calculated them, named the limits and proposed variants. Along the way a small browser tool was built for looking at flow sections: one specifies section planes and hands them to the AI with “send and next section” or “send and finish” (`stellarator/schnitt_werkzeug.py`).

11 Coolant, strength and heat balance up to the turbine

This section follows the cooling from a single segment to the grid. Three questions started it: what does the thinner inner skin gain? Which variant is best when the figures of merit contradict each other? And which pressure and temperature are realistic? The user settled the last one: *reuse an existing turbine from a nuclear power plant and keep everything else as simple as possible.*

11.1 Variant f12: 5 mm inner skin

f12 is f11 with a 5 mm instead of 30 mm inner skin and a 500 mm outlet pipe (rule: outlet always larger than inlet). Flow with FluidX3D as in Section 9.

	f11	f12
Spread around the circumference (gyroid core)	8 %	11 %
Pressure drop	22.0 kPa	19.4 kPa
Flow rate	288 l/s	289 l/s
Velocity on the plasma side	0.10 m/s	0.10 m/s
Temperature drop across the inner skin at 0.5 MW/m^2	600 K	100 K

Table 5: f11 and f12 compared.

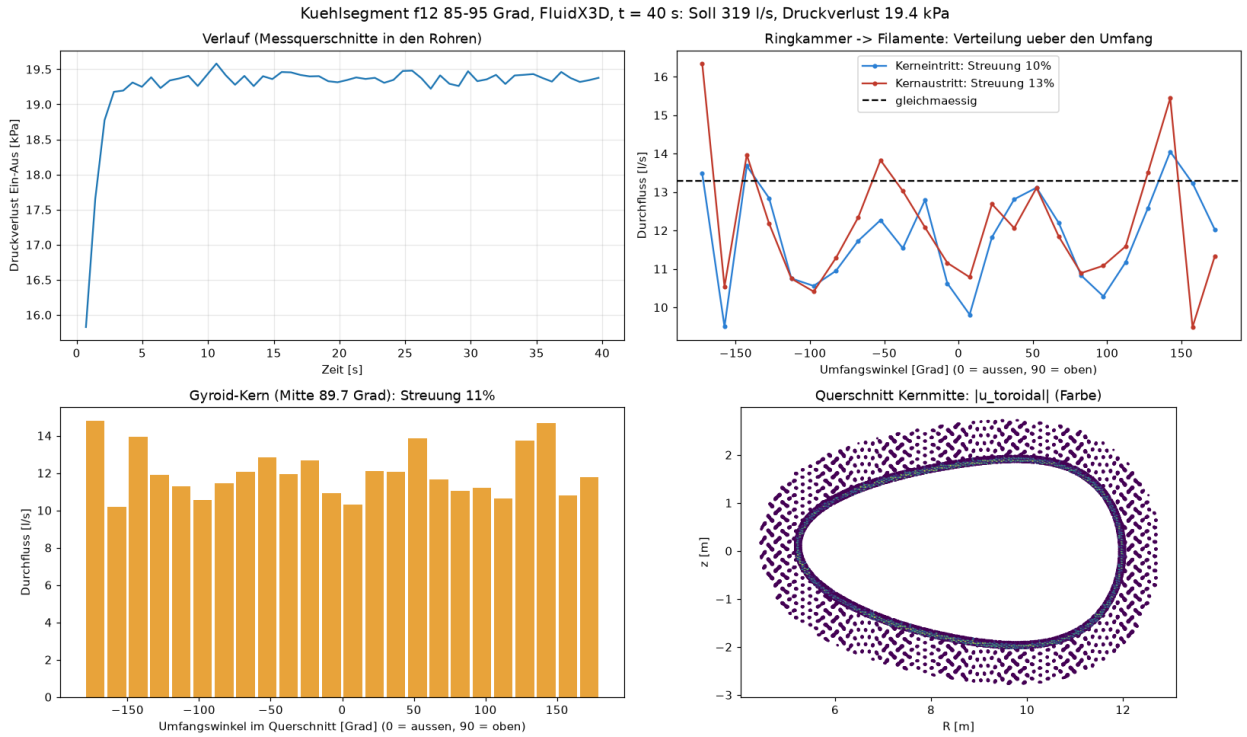

Figure 25: Evaluation of f12: pressure drop over time, distribution around the circumference and flow through the gyroid core.

The wall is no longer the bottleneck. What remains is the transfer into the slow water on the plasma side, about 410 K.

11.2 Ranking by pairwise comparison

Each variant yields several figures of merit that often conflict: f12 beats f11 on temperature and pressure drop but loses on spread. The user suggested deciding how to proceed by pairwise comparison so that resources are used well. Implemented in `stellarator/varianten_vergleich.py`:

1. **Pareto front:** a variant drops out if another one is worse in no criterion and better in at least one. No weights are needed.

2. **Weights by AHP** (Analytic Hierarchy Process): the user compares the criteria in pairs on a scale of 1 to 9. From the comparison matrix A with $a_{ji} = 1/a_{ij}$ the weight vector w is the eigenvector of the largest eigenvalue $\lambda_{\max}$. Consistency is checked by $\mathrm{CR} = \frac{(\lambda_{\max}-n)/(n-1)}{\mathrm{RI}_n}$ with the random index $\mathrm{RI}_4 = 0.90$; below 0.1 the judgements count as consistent.
3. **Rank:** weighted sum of the figures of merit normalised to 0 to 1.

The user's judgements: temperature is five times as important as spread, pressure drop and stagnant water each; pressure drop is three times as important as spread and five times as important as stagnant water. This gives weights of 0.61 for temperature, 0.23 for pressure drop, 0.08 for spread and 0.08 for stagnant water, with CR = 0.10, right at the limit.

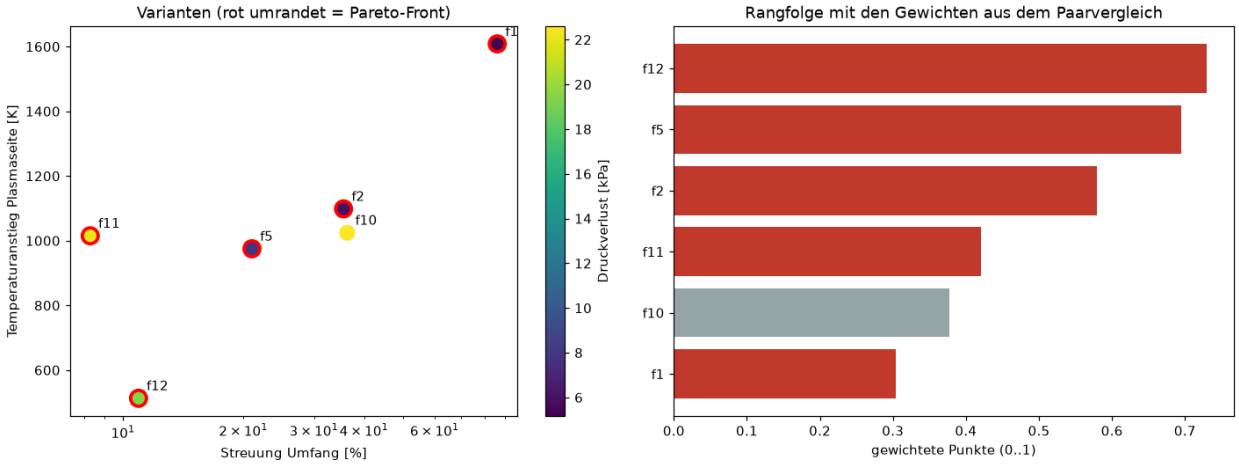

Figure 26: Left: variants by temperature rise and spread, colour pressure drop, red outline the Pareto front. Right: ranking with the user's weights.

Ranking: f12 (0.73) ahead of f5 (0.69), f2 (0.58), f11 (0.42), f10 (0.38) and f1 (0.30). Notably, f11 had been optimised for an even distribution, which the user in the end weighted low. The pairwise comparison thus showed that effort had gone into a secondary goal.

11.3 Does the thin skin hold the pressure?

A thin skin is good thermally but has to carry the coolant pressure. `stellarator/haut_festigkeit.py` computes with Elmer a steel plate of thickness t over 2×2 gyroid cells ($a = 100\,\mathrm{mm}$). It is clamped where the 10 mm gyroid walls touch the skin. In the plane of the skin these are the points with

$$\frac{|f|}{|\nabla f|}\,\frac{a}{2\pi} < \frac{w}{2}, \qquad f = \sin X\cos Y + \sin Y\cos Z + \sin Z\cos X, \quad X = 2\pi x/a.$$

In between, the coolant pushes towards the plasma. As a check, a strip with straight supports at a spacing of $b = a/2 - w = 40\,\mathrm{mm}$ was computed. With both edges clamped, $\sigma = p\,b^2/(2t^2)$: at $t = 5\,\mathrm{mm}$ Elmer gives 34.3 MPa per MPa of pressure, the formula 32.0 MPa.

Skin thickness	Stress per MPa of pressure	Allowable pressure at 150 MPa
5 mm	39 MPa	3.8 MPa
10 mm	14 MPa	10.5 MPa
15 mm	9 MPa	16 MPa

Table 6: Inner skin over the gyroid (Elmer, linear elastic, no thermal stresses).

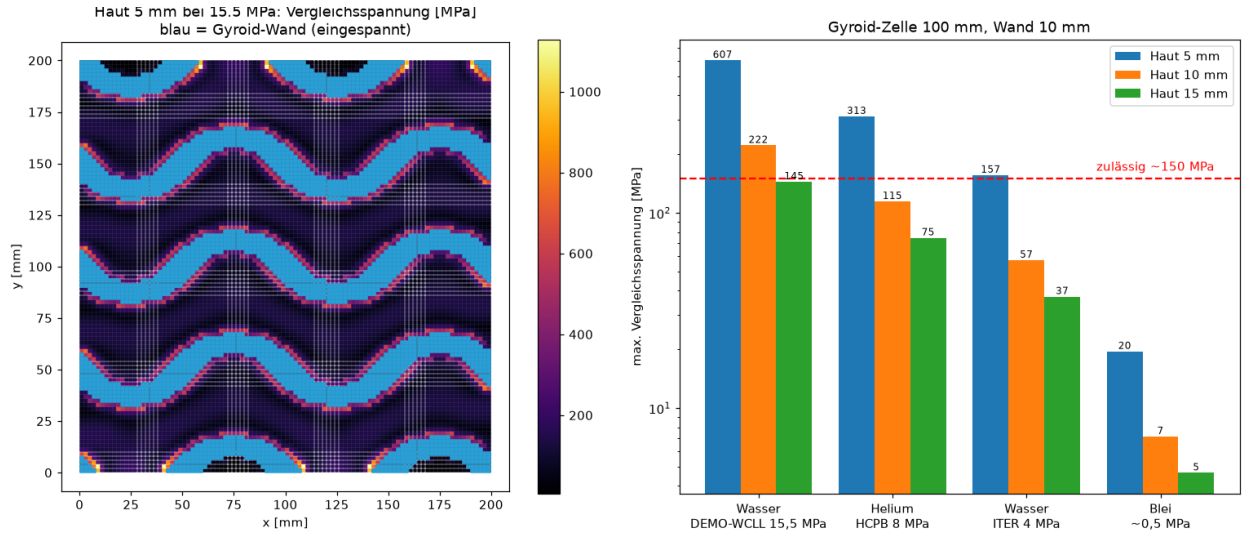

Figure 27: Left: von Mises stress in the 5 mm skin at 15.5 MPa, gyroid walls in blue. Right: peak stress per coolant and skin thickness.

For pressurised water as in a nuclear plant (≈ 16 MPa) a water-filled gyroid is therefore unsuitable. High-pressure water belongs in small channels: a 6 mm channel under a 5 mm skin reaches only about 11 MPa by the same formula. That is the double filament of Section 10.

User's remark: the gyroid only rests on lines. The gyroid walls touch the skin only along narrow lines (26 % of the area in the model). Forces collect there; in the figure above the stress peaks sit exactly at the edges of these joints. In the double filament the water pressure is balanced inside the channel plate (skin and back wall pull on each other through the webs), so the gyroid does not carry it. What remains are thermal expansion between the hot channel plate and the gyroid, eddy-current forces and fatigue under cyclic operation, all acting on these lines. Possible remedies: fillets at the joints, a gyroid that thickens towards the skin, or a layer of ribs as a transition. Not yet calculated.

11.4 Sourced reference values

Coolant data of real plants were researched with the local agent PI. A first attempt with another local agent (Hermes) produced invented values and links. For PI a skill `web-recherche` was therefore written: it actually fetches pages and records each value only with a verbatim quote and its source. A checker confirms that each quote appears in the fetched text. Because the local model has little context, PI researches each plant in a separate call; the user suggested this after a long run had produced nothing. All 29 quotes were confirmed. PI still misread two rows (inlet and outlet swapped at the steam generator); they were corrected on review.

Plant	Coolant	In / out	Pressure	Source
EPR	water	295 / 320 °C	15.5 MPa	Wikipedia
VVER-1000	water	289 / 319 °C	16.0 MPa	Wikipedia
BREST-OD-300	lead	340 / 540 °C	–	Wikipedia
ALFRED	lead	390 / 520 °C	unpressurised	Wikipedia
HTR-PM	helium	250 / 750 °C	7.0 MPa	OSTI report
ITER blanket	water	70 °C / –	4 MPa	iter.org
EU-DEMO HCPB	helium	300 / 500 °C	8 MPa	KIT

Table 7: Sourced coolant data; full quotes in `docs/recherche/pi/ergebnis.md`.

11.5 Decision: pressurised water as in a nuclear plant

With an existing turbine from a pressurised water reactor the steam conditions are fixed, and with them the primary circuit: pressurised water at about 289 °C → 319 °C and 16 MPa as in the VVER-1000. Steam generators, main coolant pumps and pressuriser are off-the-shelf. Lead (freezing at 327 °C, braking in the magnetic field) and helium (8 MPa, blowers) are dropped. The first-wall channels (6×6 mm, three passes) were recomputed for these conditions: steel at most 441 °C, channel wall 340 °C (boiling at 347 °C), pressure

drop 0.24 MPa. The steel thus sits in the working range of EUROFER (about 300 °C to 550 °C); with water at 60 °C to 150 °C it was below it. The margin to boiling, 7 K, is tight (first stated as 5 K with a wrong saturation temperature).

11.6 Derivation of the heat balance

The chain runs backwards from the grid to the plasma, with the same assumptions as the power balance of stage 0 (`invers/stufe0_bilanz.py`), here split by component (`stellarator/waermebilanz.py`).

1. Thermal power from grid power. P_{grid} goes to the grid; the turbine converts the heat P_{th} into electricity with efficiency η. Pumps ($f_p P_{\text{th}}$) and the remaining house load P_{house} (coil refrigeration, tritium plant, buildings) are subtracted:

$$P_{\text{grid}} = \eta P_{\text{th}} - f_p P_{\text{th}} - P_{\text{house}} \quad \Rightarrow \quad P_{\text{th}} = \frac{P_{\text{grid}} + P_{\text{house}}}{\eta - f_p} = \frac{1000 + 60}{0.35 - 0.04}\,\text{MW} = 3419\,\text{MW}.$$

2. Fusion power. One D-T fusion releases 17.59 MeV; the alpha particle carries 3.52 MeV, the neutron 14.07 MeV. The alpha fraction is $f_\alpha = 3.52/17.59 = 0.200$. Nuclear reactions in the blanket multiply the neutron energy by $M = 1.15$:

$$P_{\text{th}} = P_{\text{fus}}\left[(1 - f_\alpha)M + f_\alpha\right] = 1.120\,P_{\text{fus}} \quad \Rightarrow \quad P_{\text{fus}} = 3053\,\text{MW}.$$

3. Split by component. The neutrons pass through the first wall and deposit their energy in the blanket. The alpha power stays in the plasma and leaves it as radiation (onto the whole first wall) or as particle flow along the field lines (onto the divertor). Half each way is assumed ($f_{\text{div}} = 0.5$):

$$P_{\text{blanket}} = (1 - f_\alpha)\,M\,P_{\text{fus}} = 2808\,\text{MW},$$
$$P_{\text{first wall}} = (1 - f_{\text{div}})\,f_\alpha\,P_{\text{fus}} = 305\,\text{MW}, \qquad q_{\text{first wall}} = P_{\text{first wall}}/S = 305\,\text{MW}/1230\,\text{m}^2 = 0.25\,\text{MW/m}^2,$$
$$P_{\text{divertor}} = f_{\text{div}}\,f_\alpha\,P_{\text{fus}} = 305\,\text{MW}, \qquad A_{\text{divertor}} \geq P_{\text{divertor}}/q_{\max} = 305\,\text{MW}/10\,\text{MW/m}^2 = 31\,\text{m}^2.$$

S is the plasma surface of the HTS design (from stage 0: neutron wall load 1.99 MW/m²). The first-wall channels are designed for 0.5 MW/m², a factor of 2 margin for local peaks. 10 MW/m² is the design load of the water-cooled divertor elements of Wendelstein 7-X (Wikipedia).

4. Cooling water. With $\Delta T = 30$ K and $c_p = 5380$ J/(kg K), $\rho = 732$ kg/m³ at about 300 °C:

$$\dot{m} = \frac{P}{c_p\,\Delta T}: \qquad \dot{m}_{\text{blanket}} = 17\,400\,\text{kg/s}\ (23.8\,\text{m}^3/\text{s}), \qquad \dot{m}_{\text{first wall}} = 1890\,\text{kg/s}\ (2.6\,\text{m}^3/\text{s}).$$

The pumping power for the first-wall channels alone is small: $\dot{V}\,\Delta p/\eta_{\text{pump}} = 26.4\,\text{m}^3/\text{s} \cdot 0.24\,\text{MPa}/0.8 \approx 8\,\text{MW}$, even if all the water passed through such channels. The assumed $f_p P_{\text{th}} = 137$ MW leaves room for steam generators, piping and blanket.

5. Turbine. $P_{\text{gross}} = \eta\,P_{\text{th}} = 1197$ MW, waste heat at the condenser $(1 - \eta)P_{\text{th}} = 2223$ MW. Comparison with sourced plants:

- **Isar 2 (Konvoi):** 3950 MW thermal, 1485 MW gross, 1410 MW net, i.e. $\eta_{\text{gross}} = 0.376$; shut down since 15 April 2023. The fusion plant would need 87 % of its heat and 81 % of its turbine output: **one turbine of this class is enough**.
- **VVER-1000:** 3000 MW thermal, net efficiency 31.7 %, about 1 GW electric. The fusion plant delivers 114 % of this heat; a single turbine of this size would be too small.

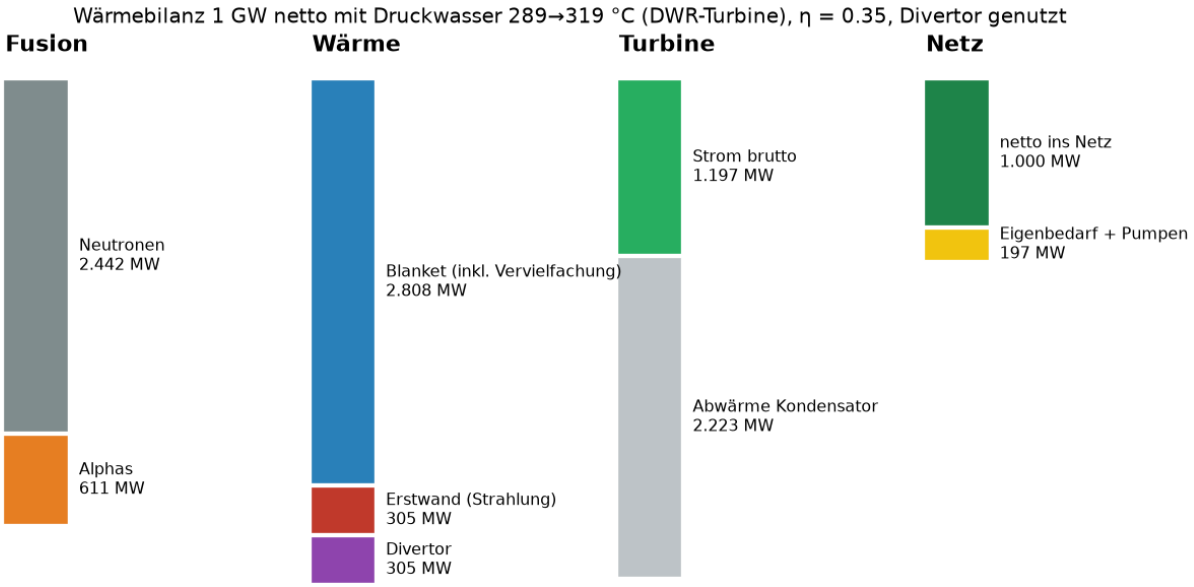

Figure 28: Heat balance for 1 GW net with pressurised water and a turbine from a pressurised water reactor (labels in German).

6. Sensitivity. Efficiency has the largest effect on the required fusion power:

Case	P_{fus}	P_{th}	P_{gross}
$\eta = 0.35$, divertor heat used	3053 MW	3419 MW	1197 MW
$\eta = 0.376$ (as Isar 2)	2817 MW	3155 MW	1186 MW
$\eta = 0.317$ (VVER-1000, net)	3417 MW	3827 MW	1213 MW
$\eta = 0.35$, divertor heat lost	3353 MW	3419 MW	1197 MW

Table 8: Fusion power required for 1 GW net.

The divertor is cooled at low temperature (copper alloy). Its heat can preheat the feedwater. If it is lost, the required fusion power rises by 10 %.

Limits. f_{div} and M are assumptions. Part of the neutron heat is already deposited in the first wall and is counted here with the blanket. The blanket itself is open: it must remove 2.8 GW at pressurised-water conditions and breed tritium at the same time. The EU-DEMO WCLL concept (water as coolant, lead-lithium as breeder) aims at exactly this; sourced values are still missing here because the publishers' sites block automated access. Whether a turbine from a decommissioned plant is actually available was not checked; the reference plants only show the matching size class.

11.7 From plasma to grid: power flow and influences

The heat balance above is derived as a chain of formulas. Figure 29 shows the same numbers as a flow: the width of each band is proportional to the power. Below the nodes are the parameters that set each value. An interactive version, in which every node and every parameter can be clicked, exists as a web page; the figure is produced by `stellarator/einflussgraph.py`.

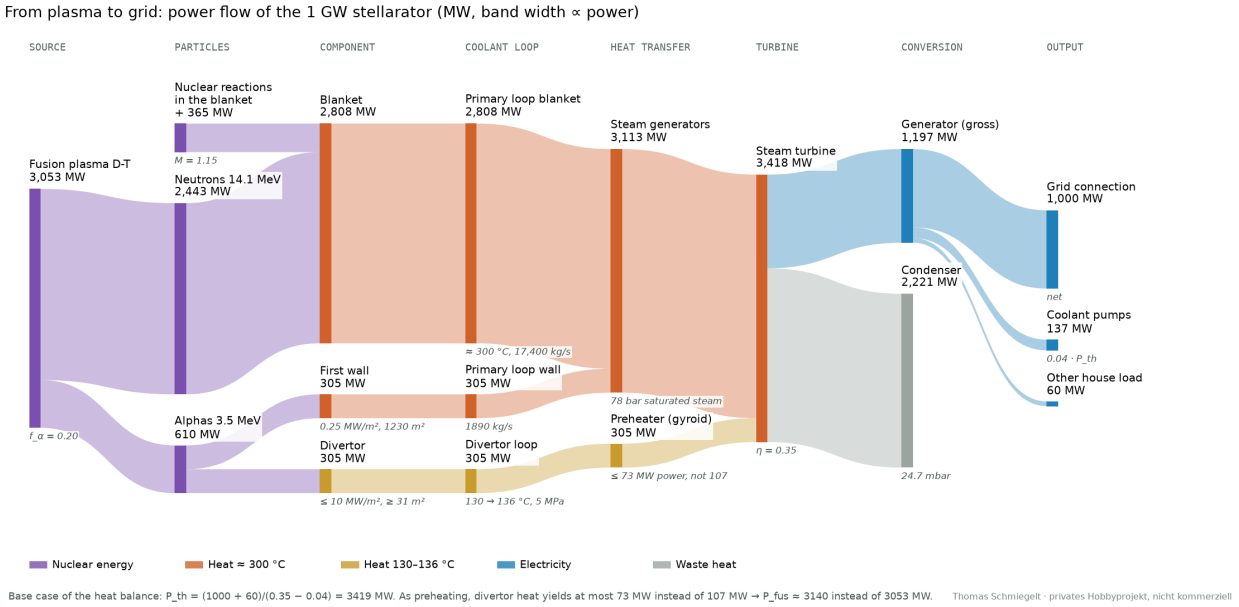

Figure 29: Power flow from the fusion plasma to the grid connection (base case of the heat balance, MW).

How to read it. Of 3053 MW fusion power, 1000 MW reach the grid, just under one third. The largest loss is not in the plasma and magnet part of the plant but in the turbine: the condenser rejects 2223 MW as waste heat. This is true of every steam plant and follows from the live-steam temperature. The chain has three places where the design directly affects the efficiency:

- the **coolant temperature** of blanket and wall (about 300 °C). It sets the live-steam pressure of 78 bar and thus $\eta = 0.35$;
- the **divertor temperature** (136 °C). It turns 305 MW of heat into preheating only, which yields at most 73 MW of electricity (Section 11.13);
- the **pressure losses** in channels and exchangers. Through the pumps (137 MW) they come directly off the net output.

The physics parameters (f_α, M) determine how much fusion power is needed for 1 GW net, not the efficiency downstream.

11.8 The components of the chain

This section explains the components from the plasma to the grid in plain words. General background is referenced; numbers marked "project" are assumptions or results of this design.

1. Fusion plasma and magnet coils. In the plasma, deuterium and tritium (D-T) fuse at about 150 million °C [41]. Gas this hot touches no wall. Magnetic fields hold it in suspension. A stellarator shapes this field with twisted external coils alone and drives no current through the plasma. It therefore runs steadily by nature, but needs stronger magnets than a tokamak [38]. Large machines such as Wendelstein 7-X use superconducting coils cooled to about 4 K with liquid helium [39]. This needs a cryogenic plant. Project: the plasma delivers $P_{\text{fus}} = 3053$ MW (about 3140 MW with the preheater correction). Open: the forces between the coils are about 1000 times those of W7-X, and the support structure is not solved yet.

2. Neutrons and alpha particles. Each D-T fusion releases 17.6 MeV. The neutron carries about 80 %; the helium nucleus (alpha particle) carries the rest [40, 42]. The neutron has no charge. It leaves the magnetic field and deposits its energy in the wall and the blanket. The alpha particle stays in the plasma and heats it [43]. Project: $f_\alpha = 0.2$ (exactly 3.52/17.59). Assumption: the alpha power leaves the plasma half as radiation onto the first wall and half as a particle stream onto the divertor ($f_{\text{div}} = 0.5$). This split is open; it depends on the real plasma.

3. First wall. The first wall is the lining that faces the plasma. It protects the coils and catches heat and radiation [44]. At ITER it sits on detachable panels, and the armour changes from beryllium to tungsten [46]. Tungsten has the highest melting point of all metals, but it is brittle [44]. Project: the wall takes 305 MW of radiation. That is 0.25 MW/m^2 on 1230 m^2. The cooling channels are designed for 0.5 MW/m^2, a factor of 2 in reserve. The cooling water (1890 kg/s) warms by 30 K. Open: neutron damage to the wall material, which is not studied here.

4. Blanket. The blanket sits behind the first wall. It has three jobs: it slows the neutrons and turns their energy into heat; it breeds tritium when neutrons hit lithium; and it shields the coils [45, 46]. Tritium hardly exists in nature, and only the blanket can supply it [45]. Lead or beryllium multiply the neutrons (multiplication) [45]. Project: $M = 1.15$ (assumption). The blanket removes 2808 MW with 17 400 kg/s of pressurised water at about 300 °C and $\Delta T = 30$ K. The EU-DEMO concept WCLL (water, lead-lithium) aims at these conditions [27]. Open: can the blanket remove 2.8 GW and breed enough tritium at the same time?

5. Divertor. The divertor guides particles and heat that escape the plasma along open field lines onto dedicated plates [47]. It removes ash and impurities and keeps the plasma clean [48]. The ash consists mainly of helium, which D-T fusion produces [40]; vacuum pumps remove it. At the ITER divertor the heat load is steadily about 10 MW/m^2, with slow peaks up to 20 MW/m^2 [48]. Project: 305 MW on at least 31 m^2 at 10 MW/m^2. It is cooled cold (water 130 °C in, 136 °C out, 5 MPa), as in EU-DEMO [29]. The project assumes a copper alloy for this. Open: the lifetime under this load.

6. Primary cooling circuits. The cooling circuits of blanket and wall carry the heat out. They use pressurised water, as in a pressurised-water reactor. The high pressure keeps the water liquid although it is far above 100 °C [49]. Project: about 300 °C, $\Delta T = 30$ K, 17 400 kg/s through the blanket and 1890 kg/s through the wall. The pressure of about 16 MPa is an assumption. Choosing these conditions is also a project decision: it allows a turbine from a pressurised-water reactor. Open: the pipes must pass through the narrow coil gaps. The calculation shows ample room (more than a factor of 10 in reserve).

7. Steam generator. The steam generator is a heat exchanger. Hot primary water flows through many thin tubes; on the outside, secondary-loop water boils into steam [50]. The tube walls separate the loops, so primary water does not reach the turbine [50, 49]. The project assumes that the primary water may carry tritium or activation products. Project: live steam at 78 bar, saturated, 2443 kg/s. These are the EPR values [35]. The steam is saturated, as in a pressurised-water reactor; the project assumes $\eta = 0.35$ for it. Open: tritium may migrate through the tube wall; this is not studied.

8. Divertor preheater. The divertor heat arrives at only about 130 °C. That is too cold for steam, but fine for preheating feedwater. This saves extraction steam, which then does work in the turbine. Project: the gyroid exchanger, a printed wall with winding channels, is to transfer the 305 MW. Result of the calculation: the preheater yields at most 73 MW of electricity, and only without temperature approach and without pumps. For comparison, 107 MW would be 35 % of 305 MW. If the heat is lost entirely, the required fusion power rises by 10 %. Open: the correlations come from small-scale experiments.

9. Steam turbine. The turbine lets the steam expand in many stages and turns pressure into rotation [51]. With reheat, the steam returns to a heater after the high-pressure section and then flows on. This yields more work and fewer droplets on the last blades [51]. Project: heat $P_{\mathrm{th}} = 3419$ MW, efficiency $\eta = 0.35$, so 1197 MW gross. For comparison, Isar 2 had 3950 MW thermal and 1485 MW gross; a turbine of this class suffices. Assumption: such a turbine is available. Open: whether a decommissioned turbine can be had.

10. Condenser and cooling tower. The condenser cools the exhaust steam back into water. It works below atmospheric pressure so that the turbine extracts more energy [52]. The heat released must go somewhere: a wet cooling tower dumps it by evaporating part of the water [53]. Project: 2223 MW of waste heat at 24.7 mbar (20.9 °C). That is almost twice the electrical output, as in every steam power plant. Open: site and water demand of the cooling tower; the project has not determined them.

11. Generator. The generator turns the rotation of the turbine shaft into electricity. It uses electromagnetic induction: a moving magnetic field induces a voltage in windings [54]. Power plants use synchronous generators, which must run in step with the grid [54, 51]. Project: it produces 1197 MW gross of alternating

current. Of this, 197 MW go to the plant's own demand. The balance uses only the overall efficiency η; there is no separate generator efficiency in it. Open: selection and size.

12. Plant's own demand. The plant uses part of its electricity itself. Project: the coolant pumps need 137 MW (0.04 P_{th}, assumption). A further 60 MW covers the rest: the coil cryogenic plant, the tritium plant and buildings. The cryogenic plant keeps the coils at a few kelvin [39]. The tritium plant is to recover tritium from the blanket and exhaust gas and feed it back to the plasma (project description, not referenced). The value for the rest is an assumption; for the first-wall channels alone the project finds only about 8 MW of pumping power. Open: the total of 197 MW is not built up from single items.

13. Transformer and grid connection. A transformer raises the generator voltage. At high voltage the same power needs less current, and the lines lose less heat [55]. The high-voltage grid then takes the power. Project: 1000 MW net go to the grid. This is the target from which the whole chain was computed backwards. Transformer losses are not included in the balance. Open: grid connection point and design.

11.9 Do the pipes fit through the gaps between the coils?

All the cooling water has to go from the vessel through the coils to the steam generators and back. On the inboard side the winding packs of the HTS design leave only 0.58 m (stage 3). This question could rule out the pressurised-water solution, so it was checked before anything else (`stellarator/rohr_durchfuehrung.py`).

Demand. From the heat balance: $\dot{V} = (\dot{m}_{\mathrm{blanket}} + \dot{m}_{\mathrm{first\ wall}})/\rho = 19\,290\,\mathrm{kg/s} \,/\, 732\,\mathrm{kg/m^3} = 26.4\,\mathrm{m^3/s}$ in each direction, so twice that pipe area for supply and return. With velocity v and inner diameter d_i one needs

$$n = \frac{2\,\dot{V}}{v\,\pi d_i^2/4} \quad \text{pipes, e.g. } n = \frac{2 \cdot 26.4}{10 \cdot \pi \cdot 0.8^2/4} = 11 \text{ at } d_i = 0.8\,\mathrm{m},\ v = 10\,\mathrm{m/s}.$$

The pipe wall follows from the boiler formula $s = p\,d_i/(2\sigma)$: at 16 MPa and 150 MPa, $s = 0.053\,d_i$, so the outer diameter is $D = d_i(1 + p/\sigma) = 1.11\,d_i$.

Space. From each point of the vessel's outer surface (plasma boundary plus 1.05 m) a straight pipe runs 4 m outward along the surface normal. The smallest distance a_{min} of this path to winding pack, potting, casing and connectors of the surrogate model ($\mathrm{R}_0 = 10.4$ m) limits the pipe diameter:

$$D_{\mathrm{max}} = 2\,(a_{\mathrm{min}} - a_{\mathrm{shield}}), \qquad a_{\mathrm{shield}} = 0.15\,\mathrm{m}.$$

a_{shield} is room for a gap and a thermal shield: the pipe is at 300 °C, the superconducting coils are cryogenic. Pipes must be at least $D + 0.2$ m apart; counting was greedy, freest spots first.

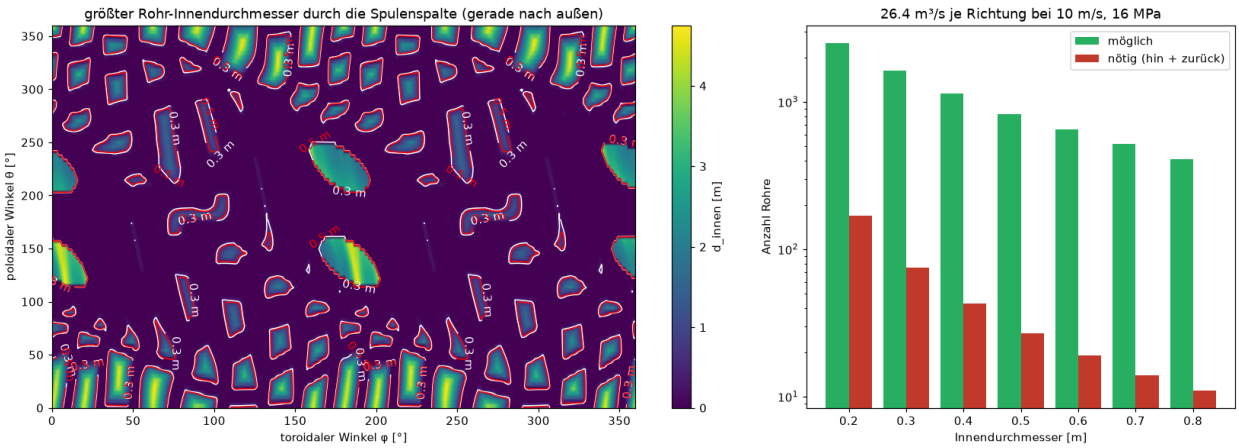

Figure 30: Left: largest inner diameter of a straight pipe over the unrolled vessel surface ($\theta = 0$ outboard, 180° inboard); the windows between coils and connectors are clearly visible. Right: possible and required pipes per size (labels in German).

d_i	required ($v = 10\,\mathrm{m/s}$)	required ($v = 5\,\mathrm{m/s}$, shield 0.3 m)	possible (default / cautious)
0.3 m	75	150	1634 / 1386
0.5 m	27	54	824 / 732
0.8 m	11	22	410 / 340

Table 9: Pipes through the gaps between the coils.

Result. The pipes fit with plenty of room, by more than a factor of 10 even in the cautious case. On 19 % of the vessel surface a straight path for 0.8 m pipes is free. The windows lie mainly outboard (θ near 0) and at top and bottom. Inboard (θ around 180°) almost everything is covered by coils: the inboard blanket has to be fed through headers from above and below. The initial concern about 0.58 m applied only to the narrowest inboard gap.

Limits. The support structure between the coils is still open (forces about 1000 times those of W7-X). It will occupy part of the windows, as will access for heating, diagnostics, vacuum pumps and maintenance. Only straight pipes along the normal were considered; with bends more room would be found.

11.10 Blanket: comparison with the EU-DEMO WCLL concept

The blanket takes 82 % of the heat and has to breed tritium at the same time. The reference is the water-cooled lead-lithium blanket (WCLL) of the European DEMO [27]. The user downloaded the open-access article; all values below are quoted verbatim from it (`docs/recherche/wcll/wcll_werte.md`).

	EU-DEMO WCLL [27]	this plant
Cooling water	295 → 328 °C, 15.5 MPa	289 → 319 °C, 16 MPa
Circuits	first wall and breeding zone separate	first wall and blanket separate
EUROFER steel	not above 550 °C	first wall at most 441 °C
First-wall heat flux	$0.32\,\mathrm{MW/m^2}$	$0.25\,\mathrm{MW/m^2}$ (mean), channels for 0.5
Fusion power / net	1998 MW / 500 MW	3053 MW / 1000 MW
Neutron wall load	$1.04\,\mathrm{MW/m^2}$	$1.99\,\mathrm{MW/m^2}$
Breeder	PbLi, Li enriched to 90 % ^{6}Li	open
Tritium breeding ratio	1.14 (target 1.15)	open

Table 10: WCLL and the pressurised-water design of this plant.

Consequences.

1. **The choice of pressurised water is confirmed.** DEMO uses practically the same conditions for WCLL as a pressurised water reactor, with two separate circuits for first wall and breeding zone. The heat balance here is split the same way.
2. **Lead returns, but as breeder, not as coolant.** In WCLL the PbLi flows only slowly through the blanket and carries the bred tritium out. The heat is taken up by water in double-walled tubes. The objections to lead as a coolant (fast flow across the field) hardly apply.
3. **The neutron wall load is twice that of DEMO.** In stage 0 the HTS design was shrunk up to the limit of $2\,\mathrm{MW/m^2}$. That means roughly twice the heat density in the blanket and twice the radiation damage to the steel, hence a shorter life of the blanket modules, and more shielding in front of the coils. The Nb_3Sn design would be at $0.59\,\mathrm{MW/m^2}$, but is almost twice as large.
4. **The tritium breeding ratio is tight.** Even DEMO misses the target of 1.15 slightly. A stellarator has more openings and a more complex shape, losing blanket area. This is a risk of its own and would need a neutronics calculation.

11.11 Graded structure: flat contact at the skin, branching outward

The user suggested solving the line-contact problem differently: “On the blanket side one would need a flat support, then it becomes a balancing act. On the one hand the metal strands must not be so dense that they obstruct the water too much, on the other hand they must be thick enough for heat conduction. Further out the material can branch heavily and become thin.” Three design rules follow (pressurised water $p = 16\,\mathrm{MPa}$, skin $t = 5\,\mathrm{mm}$, allowable $\sigma = 150\,\mathrm{MPa}$):

1. Support spacing at the skin. With the factor 1.22 from the Elmer calculation over the gyroid (39.2 instead of 32.0 MPa per MPa):

$$\sigma = 1.22\,\frac{p\,b^2}{2t^2} \le \sigma_{\text{allow}} \quad \Rightarrow \quad b \le t\sqrt{\frac{2\sigma_{\text{allow}}}{1.22\,p}} = 5\,\text{mm}\cdot\sqrt{\frac{300}{19.5}} \approx 20\,\text{mm}.$$

At 4 MPa it would be 39 mm. With supports every 20 mm the thin skin holds pressurised water as in a nuclear plant.

2. How deep the strands conduct heat. A strand of thickness d in water acts as a cooling fin with $m = \sqrt{4h/(k\,d)}$; roughly the length $2/m$ is effective (steel $k = 25\,\text{W}/(\text{m}\,\text{K})$):

d	$h = 1200\,\text{W}/(\text{m}^2\text{K})$ (slow)	$h = 10\,000\,\text{W}/(\text{m}^2\text{K})$ (fast)
3 mm	8 mm	3 mm
10 mm	14 mm	5 mm

Table 11: Effective fin depth $2/m$.

The heat is in the water within 3 mm to 15 mm. Only this first layer needs thick strands; behind it the strands only carry load.

3. How thin it may get further out. The pressure pushes inner and outer skin apart; the strands hold them together in tension. Their area fraction φ must everywhere satisfy

$$\varphi \ge \frac{p}{\sigma_{\text{allow}}} = \frac{16}{150} \approx 11\,\%.$$

Branching is therefore allowed as long as the sum of the cross-sections remains. For a gyroid of cell a and wall thickness t_g, $\varphi \approx 3.09\,t_g/a$.

Implementation in PicoGK. `picogk/GradientWand`: 5 mm skin, then four layers whose cell doubles from layer to layer. Between two layers the distance fields of the gyroids with cell a and $2a$ are blended smoothly, so the walls branch. The joints are filleted with 1.5 mm.

Layer	Cell	Wall	to depth	Steel fraction
0	20 mm	2.5 mm	30 mm	39 %
1	40 mm	3 mm	80 mm	23 %
2	80 mm	4.5 mm	180 mm	17 %
3	160 mm	8 mm	570 mm	15 %

Table 12: Layers of wall piece g1 ($240 \times 240 \times 600$ mm, voxel 1 mm); 84 % water overall.

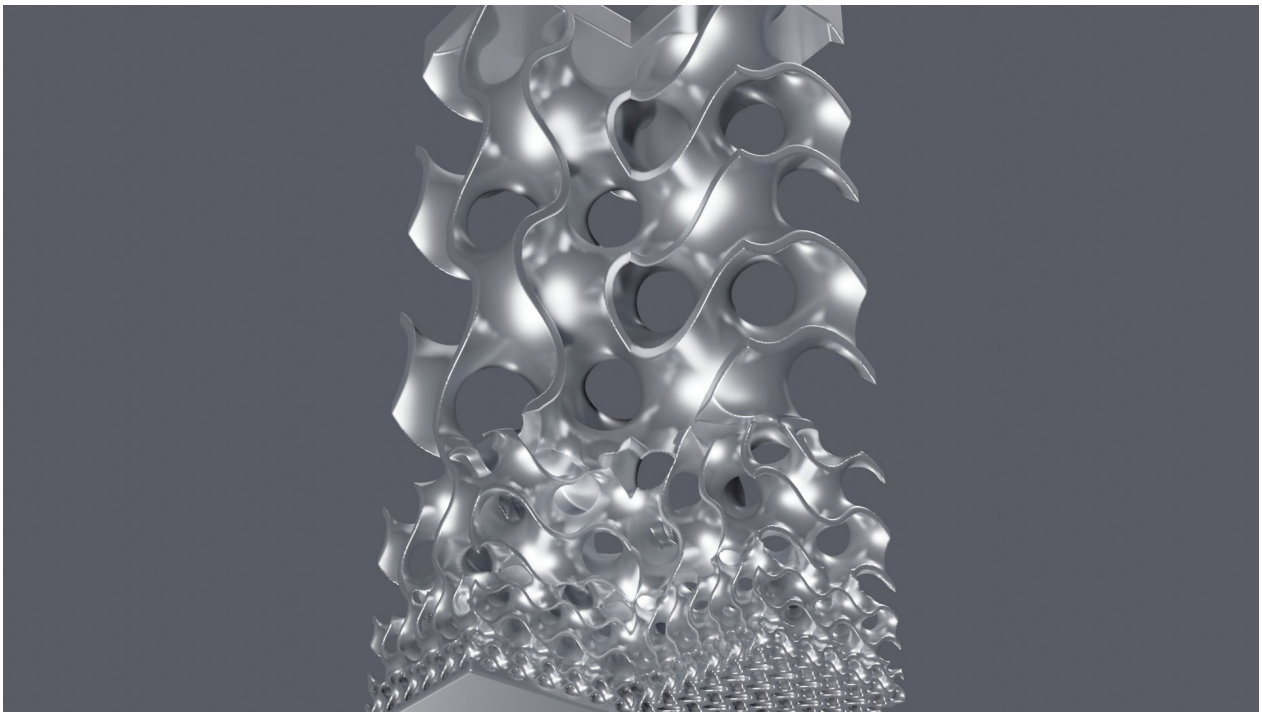

Figure 31: Graded wall g1 seen from the plasma side, one quarter cut away: small cells at the skin at the bottom, the cell doubles with each layer further out.

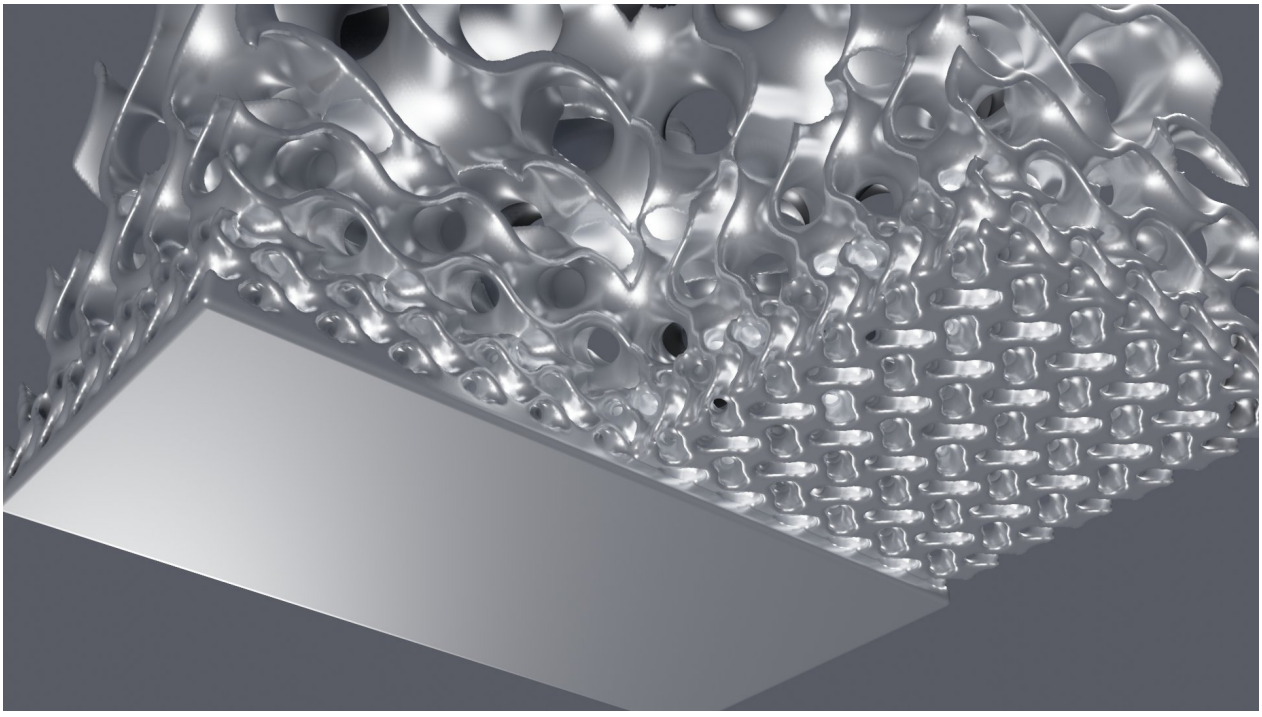

Figure 32: Close-up of the skin with a window: the fine layer supports the skin almost over its full area, support spacing about 10 mm.

Still open are the flow test (FluidX3D: does the water pass through the dense skin layer or bypass it further out?) and the strength test (Elmer: skin and branches at 16 MPa). The wall piece also writes out the water volume for this.

11.12 How does the heat spread sideways and outward?

The user's question: how do we spread the heat sideways and so lead the heat flow outward?

Estimate. A plate of thickness t and conductivity k, cooled on one side with heat transfer coefficient h, spreads a hot spot sideways only over the spreading length

$$\lambda = \sqrt{k\,t/h}. \tag{1}$$

For 5 mm steel ($k \approx 25\,\mathrm{W/(m\,K)}$) and $h = 10^4\,\mathrm{W/(m^2\,K)}$ that is 3.5 mm. A gyroid wall of thickness t acts as a fin cooled on both sides. Its effective depth is

$$1/m = \sqrt{k\,t/(2h)}, \tag{2}$$

which is 1.8 mm for 2.5 mm steel. Both values are small compared with the 20 mm cell. Expectation: in steel the heat stays where it arrives and enters the water directly behind the skin.

Calculation. Elmer HeatSolve (steady state) on a cut-out of wall g1:

- Cut-out $80 \times 80 \times 100$ mm, one full period of the first three layers. Same layers and blending as `GradientWand`; fillets from morphological closing.
- Voxel hexahedra of 0.5 mm, 1.4 to 1.5 million elements.
- Plasma side: $q = 0.5\,\mathrm{MW/m^2}$.
- Wetted surfaces: $h = 10^4\,\mathrm{W/(m^2\,K)}$ to water at 304 °C. The voxel staircase area exceeds the smooth area by the factor $|n_x| + |n_y| + |n_z|$, so h is divided by this factor on each face.
- Sides and back: adiabatic.
- Energy balance: input and output agree to 0.1 %.

Five cases were computed:

- uniform load;
- hot spot (radius 10 mm, $2q$);
- the same spot with 3 mm tungsten in front of the skin ($k = 130$);
- skin and structure of copper alloy ($k = 320$, like CuCrZr in the divertor);
- doubled h (faster water).

The k values are rough assumptions, not sourced.

Case	plasma side max	wet wall max	removed within 10 mm	removed under spot
steel, uniform	469 °C	364 °C	95 %	–
steel, hot spot	588 °C	399 °C	95 %	6.9 %
tungsten + steel, spot	535 °C	377 °C	95 %	5.6 %
copper alloy, spot	351 °C	340 °C	62 %	5.8 %
steel, spot, h doubled	559 °C	363 °C	98 %	7.0 %

Table 13: Heat test on wall g1. Saturation temperature at 16 MPa: 347.4 °C. “Removed within 10 mm”: share of heat entering the water within 10 mm behind the plasma side of the steel skin. “Removed under spot”: share leaving within the spot radius. The spot receives 9.4 % of the heat; any shortfall was spread sideways.

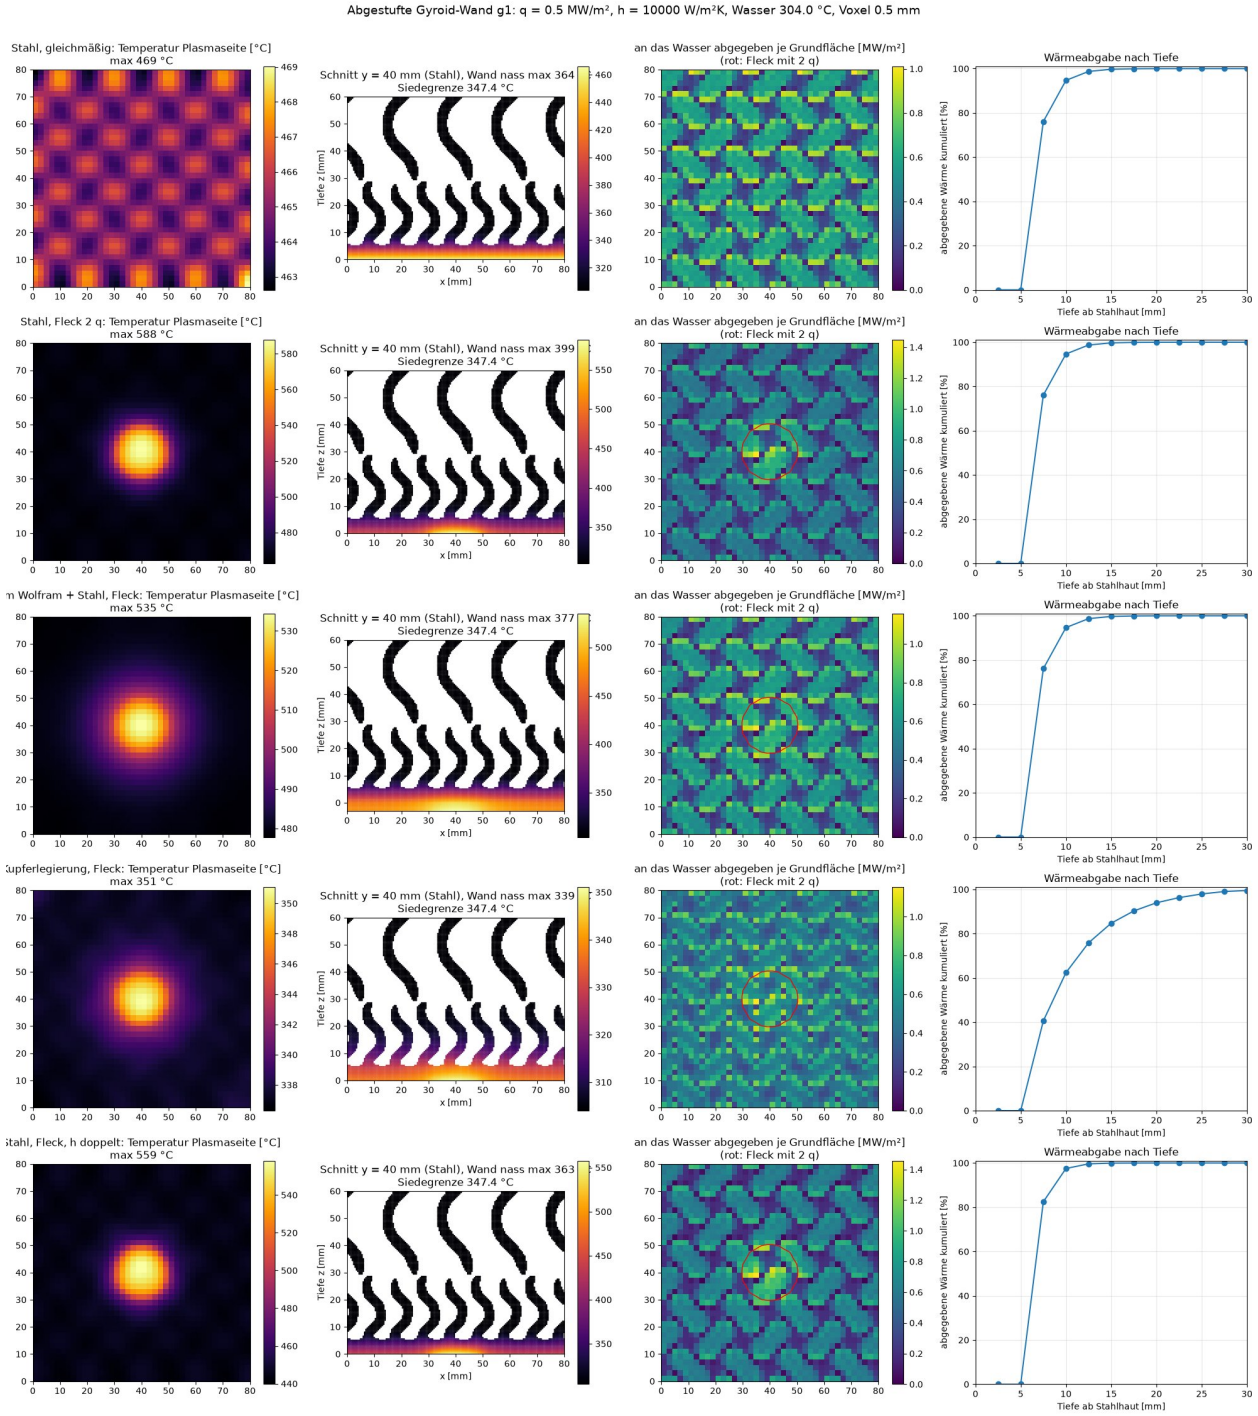

Figure 33: Heat test of wall g1. Columns from left: plasma-side temperature; section through the middle; heat given to the water per footprint area (red: spot); cumulative heat removal over depth.

Result.

- **In steel the heat does not travel into the depth.** 95 % leaves the steel within the first 5 mm behind the skin, through the back of the skin and the roots of the struts. The struts themselves stay almost at water temperature. The wetted area is about 15 times the footprint but is hardly used. In steel the graded structure is a load-bearing structure, not a heat exchanger. This confirms the estimate.
- **The saturation temperature is exceeded even under uniform load.** 3 % of the wet wall is above 347.4 °C, where the struts join the skin and less area cools locally. With the hot spot the wall reaches 399 °C.

- **Tungsten spreads the hot spot sideways**: 53 K less at the hottest point. Because the extra 3 mm of tungsten must also conduct the heat, the plasma side is 12 K warmer on average.
- **Faster water** lowers the wall temperature most for the same material, but barely helps the spot.
- **Copper alloy is the only case that carries heat into the depth.** Only 62 % leaves within the first 5 mm; the rest leaves further back through the struts. The plasma side stays below 351 °C and the wet wall below saturation. The fins now work because $1/m = \sqrt{kt/(2h)}$ grows to 6 mm with $k = 320$.

Conclusions.

- Only a well-conducting material spreads heat sideways and into the depth. Options are a thin tungsten or copper layer at the skin, or copper-alloy struts in the first 10 mm to 20 mm. Copper is problematic under neutrons (embrittlement, activation) but common in divertors. The structure further back can stay steel; it then only carries load.
- With steel alone the water must flow fast directly at the skin. The skin layer therefore needs its own supply, as the rules above already require.
- In the blanket the heat is generated throughout the volume by neutrons. There, branching well-conducting paths to the water make sense, as constructal theory describes for cooling a heat-generating volume [28]. That is the next step, together with flow (FluidX3D) and strength.

Limits of the calculation:

- h is uniform because no flow calculation is coupled.
- There is no volumetric neutron heating.
- k is constant.
- The sides are adiabatic instead of periodic.

Script: `stellarator/gradient_waerme.py`.

11.13 Divertor preheater: which design yields the most electricity?

The divertor heat arrives at low temperature. In EU-DEMO the target coolant enters at 130 °C and 5 MPa and heats up by 6 K. The authors plan to use this heat for preheating [29]. In this design it amounts to 305 MW (see heat balance).

Question and hypothesis. Which preheater design extracts the most net electricity from these 305 MW? The candidates are a tube bundle, a plate exchanger, gyroid (uniform and graded) and diamond.

$$P_{\text{net}} = \Delta P_{\text{turbine}} - P_{\text{pump, divertor side}} - P_{\text{pump, condensate side}}.$$

Hypothesis H1: at equal pumping effort, a TPMS structure reaches a smaller approach temperature and therefore more net electricity.

1. Plant side. The divertor water preheats condensate, which saves bleed steam. The saved steam still does work on its way to the condenser. For heating by $\mathrm{d}T$:

$$\mathrm{d}W = \dot{m}_k \, c_p \, \mathrm{d}T \, \frac{h_A - h_{\text{exh}}}{h_A - h'_A}, \qquad p_A = p_s(T + \Delta T_{\text{TTD}}).$$

Here h_A is the bleed steam enthalpy, h'_A the enthalpy of saturated condensate at bleed pressure and h_{exh} the enthalpy at the turbine exhaust. ΔT_{TTD} is the terminal temperature difference of the heater. All states come from CoolProp [33].

The expansion line starts with saturated steam at 78 bar, as in the EPR [34]; English Wikipedia gives 77.2 bar at the turbine inlet [35]. It ends at a condenser at 24.7 mbar (20.9 °C). The condensate flow is the EPR steam flow (2443 kg/s) scaled to our 3419 MW thermal, i.e. 1820 kg/s. The condensate then heats up by $Q/(\dot{m}_k c_p) = 39.8$ K. It can reach at most 136 °C.

Without approach temperature and without pumping this gives an **upper limit of** 73.3 MW, i.e. 24 % of the divertor heat.

2. Heat exchanger side. Counterflow with $UA = Q/\Delta T_{\mathrm{lm}}$. The approach at the hot end is $\Delta T_1 = 136\,^{\circ}\mathrm{C} - T_{k,\mathrm{out}}$. At the cold end $\Delta T_2 = 130\,^{\circ}\mathrm{C} - T_{k,\mathrm{in}} = \Delta T_1 + 33.8\,\mathrm{K}$. The divertor water acts almost like a source at fixed temperature. ΔT_{lm} therefore stays large even for a small ΔT_1, e.g. 7 K at $\Delta T_1 = 0.3\,\mathrm{K}$.

Both sides share the wall area A. With $d_h = 4V_f/A$ the fluid volume per side is $V_f = A\,d_h/4$. For equal path length this gives

$$\frac{Re_{\mathrm{hot}}}{Re_{\mathrm{cold}}} = \frac{\dot{m}_{\mathrm{hot}}\,\mu_{\mathrm{cold}}}{\dot{m}_{\mathrm{cold}}\,\mu_{\mathrm{hot}}} \approx 7,$$

independent of geometry. The reason is the unequal mass flows: 12 000 kg/s of divertor water against 1820 kg/s of condensate. A symmetric TPMS structure therefore cannot keep both sides inside the valid range of the correlation ($2961 < Re < 18254$). Several passes on the condensate side (n) solve this, because the condensate path becomes n times longer:

$$A_{f,\mathrm{hot}} = n\,A_{f,\mathrm{cold}}\,\frac{d_{h,\mathrm{hot}}}{d_{h,\mathrm{cold}}}, \qquad \Delta p = 4f\,\frac{L}{d_h}\,\frac{\rho v^2}{2}, \qquad P_{\mathrm{pump}} = \frac{\Delta p\,\dot{V}}{\eta_P}.$$

Core volume: $V = A\,(d_{h,\mathrm{hot}} + d_{h,\mathrm{cold}})/4 + A\,t$, without headers and manifolds.

For each design the scan covers 40 approaches (0.3 K to 25 K), 12 hydraulic diameters, 13 ratios of the two diameters, 24 velocities and $n = 1$ to 8. For each volume the best valid design counts (Pareto front).

Design	Heat transfer and friction (Fanning)	Band
Diamond	$Nu = 0.2644\,Re^{0.69}Pr^{1/3}$, $f = 1.850\,Re^{-0.17}$, CFD [30]	Li 2020, Li 2022 [31]
Gyroid uniform	as diamond, $f \times 1.145$ (12.3–16.7 % more friction [30])	Li 2020, Li 2022
Gyroid graded	$j/f^{1/3}$ 28 % better (measured [32]), here $Nu\times$ 1.28	as above
Plate (stopgap)	as uniform gyroid; measured gyroid is "slightly lower" [32]	up to +10 %
Tube bundle	Gnielinski and Petukhov, shell side as tube side (weakest assumption)	–

Table 14: Correlations; verbatim quotes in `docs/recherche/vorwaermer/`.

Hydraulic diameters: tube 10 mm to 25 mm, plate 3 mm to 10 mm, TPMS 2 mm to 15 mm. Assumptions: wall 0.5 mm stainless steel (printed as in [32]), terminal temperature difference 3 K, isentropic turbine efficiency 0.85 (HP) and 0.88 (LP), reheat at 11 bar to 270 °C, pumps 0.8, condensate pressure 2 MPa.

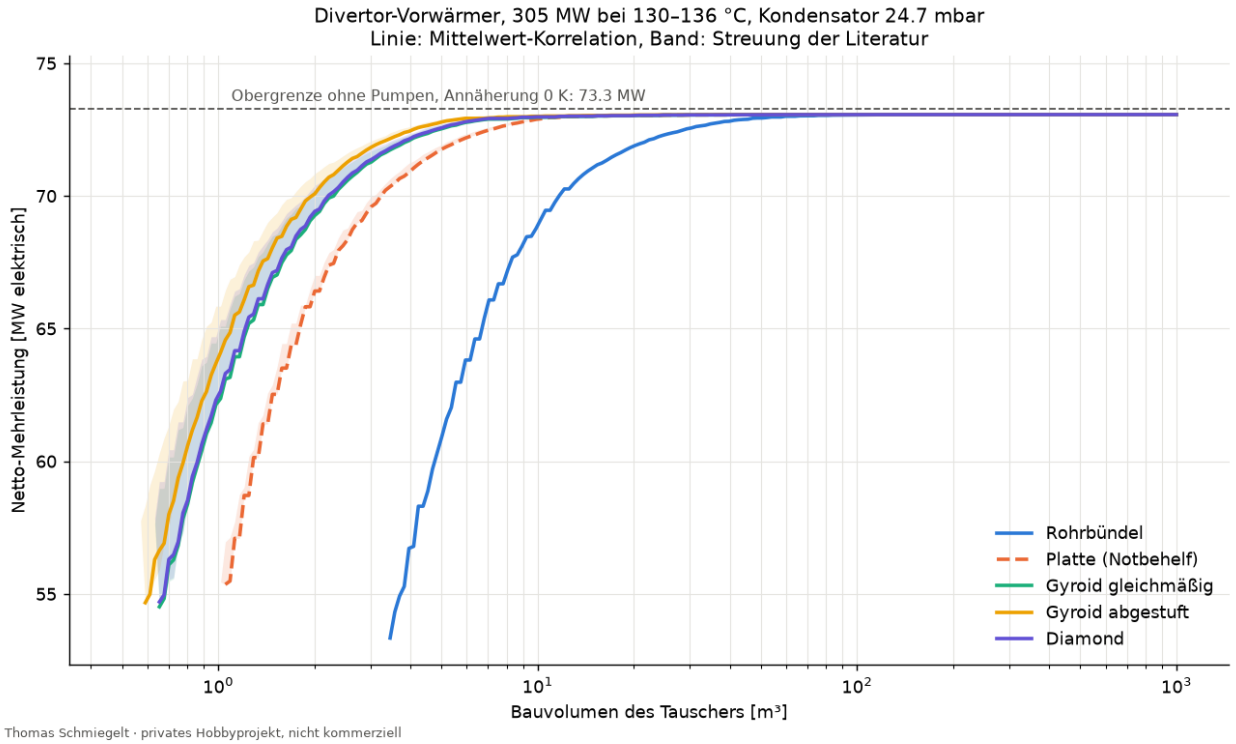

Figure 34: Net extra power versus core volume (labels in German). The line is the central correlation, the band the literature spread. Dashed: plate (stopgap).

Design	Volume for 95 %	Volume for 99 %	Net up to 50 m³
Tube bundle	11 m³	30 m³	72.9 MW
Plate (stopgap)	2.8 m³ to 3.0 m³	6.7 m³ to 7.3 m³	73.1 MW
Gyroid uniform	1.8 m³ to 2.1 m³	4.5 m³ to 5.0 m³	73.1 MW
Gyroid graded	1.6 m³ to 1.8 m³	3.8 m³ to 4.2 m³	73.1 MW
Diamond	1.7 m³ to 2.1 m³	4.4 m³ to 4.9 m³	73.1 MW

Table 15: Core volume for 95 and 99 % of the upper limit. The range comes from the correlations.

3. Findings.

- **Efficiency does not decide the design.** From about 30 m³ on, every design reaches 72.9 MW to 73.1 MW. Pumping power is then 0.01 MW to 0.14 MW. H1 holds only at small volume. Given enough size, all designs exhaust the upper limit.
- **The designs differ in volume.** For 99 % a TPMS exchanger needs 4 m³ to 5 m³, a plate about 7 m³ and a tube bundle about 30 m³. The graded gyroid is smallest, diamond close behind. The choice must therefore be made on volume, cost, manufacturing, fouling and strength, not on efficiency.
- **Correction of the heat balance.** The balance above counted the divertor heat at the full efficiency of 0.35, i.e. 107 MW of electricity. As preheating it yields at most 73 MW. For 1 GW net the required fusion power thus rises by about 3 %, from 3053 MW to roughly 3140 MW.
- **Mass flows.** Divertor water and condensate flow at a ratio of 7:1. A symmetric exchanger therefore needs several passes on the condensate side. The best designs use $n = 8$ and sit at the lower limit $Re \approx 3000$. Slower flow would be even better, but verified correlations for it are missing.
- **Bigger levers than the design.** A condenser at 50 mbar (river cooling) lowers the upper limit to 63.5 MW. A condensate flow of 75 % lowers it to 68.7 MW. In EU-DEMO the cassette body also removes heat at 180 °C to 210 °C (4.2 MW per cassette) [29]. Heat at higher temperature is worth more. That is the next calculation step.

Limits of the calculation:

- The correlations scatter widely, and the plate curve is not a model of its own.
- The expansion line is simplified: one efficiency per section, bleeds without feedback on the steam flow.
- Headers, manifolds and flow distribution are missing. The best TPMS designs are only about 1.3 m long and need a large frontal area for that.
- Whether a 0.5 mm printed wall holds 5 MPa against 2 MPa has not been checked. That is a question for strength experts.

Script: `stellarator/vorwaermer.py`, results in `stellarator/daten/vorwaermer.json`.

11.14 Exchanger modules built and simulated with FluidX3D

The choice of preheater design so far rests on literature correlations whose range of validity (Pr = 3–5) does not cover divertor water with Pr ≈ 1.3–1.8. The modules were therefore built and the flow computed directly. The available computer cannot handle OpenFOAM; the lattice Boltzmann solver FluidX3D runs on the graphics card [17, 20].

1. Modules. PicoGK builds one graded two-fluid block each ($48 \times 48 \times 96$ mm, wall 0.5 mm, shell 1.5 mm): gyroid with cell size 4.33 mm → 5.8 mm and diamond with 5.37 mm → 7.2 mm. Both give the same hydraulic diameter $d_h = 2 \to 2.9$ mm. The cell size grows logarithmically along the block, with phase $p_z = 2\pi L/(a_1 - a_0)\,\ln(a/a_0)$. End caps close the cold side, side windows the warm side. At 0.2 mm voxel size the block is 39.7 % steel and 30 % each warm and cold side. The plates (flat and chevron, $\beta = \pm 60°$) are built parametrically in FreeCAD.

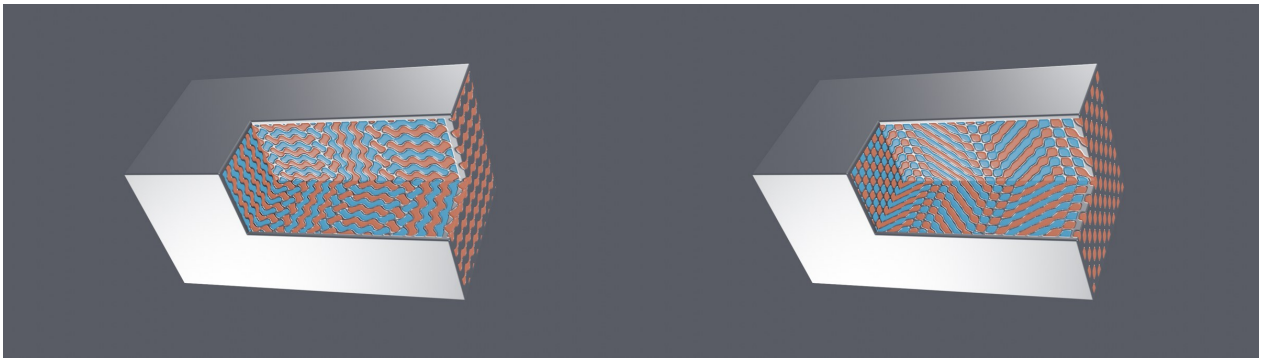

Figure 35: Graded two-fluid modules, cut open. Left gyroid, right diamond. Red is the warm (divertor) side, blue the cold (condensate) side, grey the steel wall.

2. Domain. A domain with inlet and outlet fails at Re ≈ 3000: the required pressure drop needs more than 25 % density difference in the lattice Boltzmann method, far outside its validity. The domain is therefore a section of one side, periodic in all directions. A body force f_z drives the flow, a uniform heat source q heats the fluid, and the wall holds $T_w = 0$. Then

$$f = \frac{f_z\, d_h}{2\, u^2}, \qquad h = \frac{q\, V_f}{A\, T_m}, \qquad \mathrm{Nu} = \frac{h\, d_h}{D}, \qquad T_m = \frac{\sum u_z T}{\sum u_z}.$$

f_z is iterated until Re matches within 15 %.

3. Changes to FluidX3D. Five points had to be changed (sources with licence note in `fluidx3d/`):

- The temperature model (D3Q7) has $c_s^2 = 1/4$. The actual diffusivity is $D = \alpha/2$; a pure diffusion test confirmed this (measured $D/\alpha = 0.44$–0.53). The code therefore passes $\alpha = 2D$.
- The wall sat one cell too far into the fluid when the wall temperature was set in fluid cells. This gave a Nu about three times too high. Solid cells now use anti-bounce-back: $g_{\mathrm{in}} = -g_{\mathrm{out}} + \frac{1}{4}(T_w - 1)$.
- At Re = 3000 the cell Péclet number is about 185. A turbulent diffusivity $D_t = \nu_t/0.85$ from the solver's Smagorinsky model is therefore added.

- The simple BGK collision for temperature became abruptly unstable at 96 cells per d_h after 172 000 steps (temperature to 10^{20} within 500 steps). The relaxation time $\tau \approx 0.504$ is close to the stability limit 0.5. It was replaced by a TRT collision after Ginzburg [37]: the antisymmetric part relaxes at the rate that sets D, the symmetric part with the "magic" combination $\Lambda = (1/\omega^+ - \frac{1}{2})(1/\omega^- - \frac{1}{2}) = 1/4$.
- The heat source is either uniform or proportional to u_z. The second form is the usual formulation for fully developed periodic flow with axially uniform heat flux (Patankar, Liu and Sparrow 1977 [36]). A uniform source raises Nu even in the laminar gap from 8.235 to 10.

4. Verification. For the laminar plate channel the values are known exactly: $f\,\mathrm{Re} = 24$, and $\mathrm{Nu} = 10$ (uniform source) or 8.235 (source $\propto u$). The simulation gives $f\,\mathrm{Re} = 29.7 \to 25.7 \to 24.7$ at 12, 24, 48 cells per channel height. For Nu it gives $10.6 \to 10.3$ (BGK, uniform) and $8.9 \to 8.5$ (TRT, $\propto u$). The error falls with order $p \approx 1.3$. The choice of Λ (1/4 or 3/16) did not change the test case.

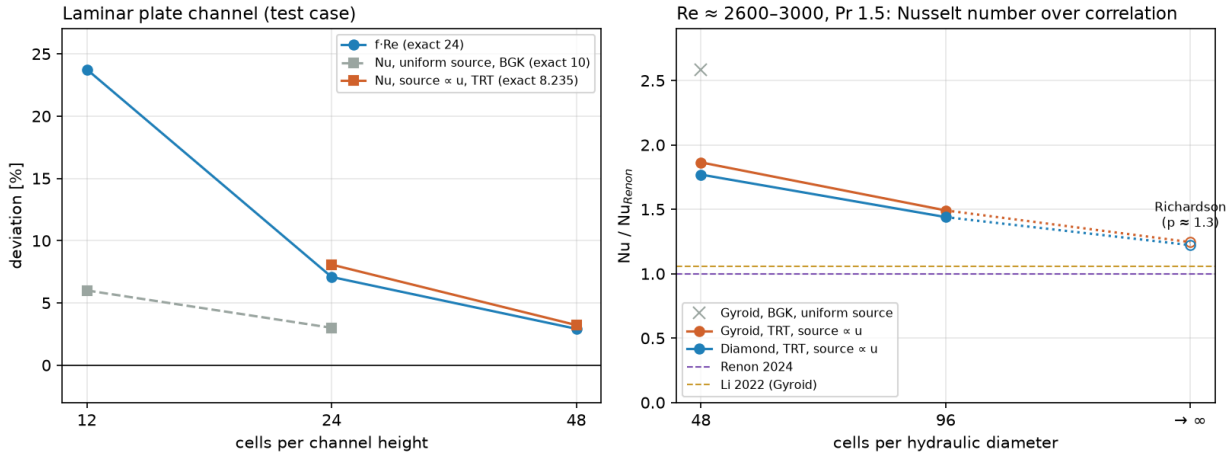

Figure 36: Left: test case laminar plate channel, deviation from the exact value. Right: Nusselt number relative to Renon at the same Re, over grid resolution; open circles: Richardson estimate.

5. Comparison of designs. All runs with TRT and source $\propto u_z$, $\mathrm{Pr} = 1.5$. The driving force is adjusted until Re matches within 15 %; at 96 cells gyroid and diamond stayed somewhat below. $j/f^{1/3}$ with $j = \mathrm{Nu}/(\mathrm{Re}\,\mathrm{Pr}^{1/3})$ is the area goodness: heat transfer per pumping effort.

Design	Grid	Re	f	Nu	$j/f^{1/3}$	$\mathrm{Nu}/\mathrm{Nu}_{\mathrm{Renon}}$
Gyroid	48	2952	0.805	139.8	0.0445	1.86
Gyroid	96	2583	0.732	102.0	0.0383	1.49
Diamond	48	2974	0.668	133.4	0.0448	1.77
Diamond	96	2676	0.596	100.9	0.0391	1.44
Chevron ±60°	48	2908	1.006	121.5	0.0364	–
Chevron ±60°	96	2921	0.863	103.5	0.0325	–

Table 16: FluidX3D results per side (periodic section). Renon applies to TPMS only.

Finding.

- **Hydraulics:** f lies within the band of the correlations (Renon 0.48–0.56, Li 2022 0.81–0.83 [30, 31]). At equal d_h diamond has about 20 % less friction than gyroid; Renon gives 12–17 %.
- **Heat transfer:** Nu lies above the correlations, but the gap shrinks with the grid: factor 1.8 at 48 and 1.5 at 96 cells. With the order from the test case, a Richardson estimate gives about 1.2–1.25. This is a rough estimate from two levels. The third level (144 cells, 62.5 million cells) gave 1.36 (gyroid) and 1.30 (diamond) and shows that the series is not yet asymptotic; extrapolated, Nu lands on Renon (Section 11.15).
- **Collision model:** BGK had overestimated Nu by about 25 % even without blowing up (194 instead of 154 at 48 cells, same source).

- **Designs:** Gyroid and diamond are level in area goodness (0.038 and 0.039 at 96 cells). Chevron plates are about 15 % lower: they reach similar Nu only with more friction. This supports the ranking of Section 11.13 (TPMS more compact than plates). The simulation shows no advantage of gyroid over diamond.
- **Limits:** uniform cell size (the grading was not simulated), a single Re, Pr = 1.5, isothermal wall without conduction in the steel, turbulence via an LES model at only moderate resolution.

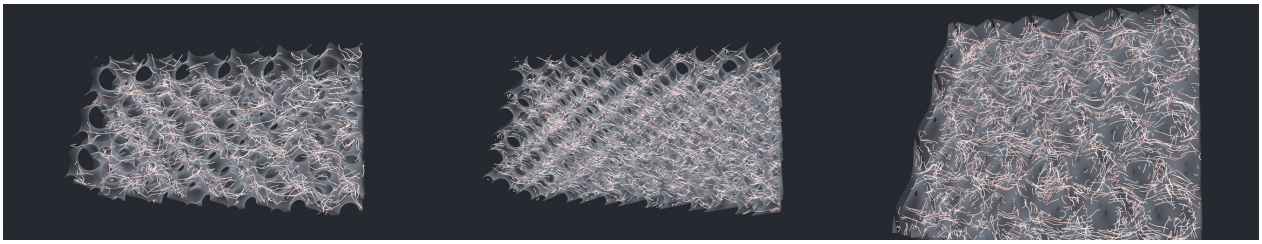

Figure 37: Particle paths in gyroid, diamond and chevron plates (from left), Re 3000, 48 cells per d_h. The wall is shown in X-ray view. A video was made from each set of paths. The colour of the paths comes from the runs with BGK collision and is for illustration only.

These and 30 further runs went into the surrogate model of Section 11.15.

11.15 Surrogate model from the flow simulations and the question of grading

Why a surrogate? For every design, the preheater calculation needs Nu and f on both sides. So far there were three ways to get them, each with a catch:

1. **Correlations from the literature** are fast but scatter widely. At Re 3000, Renon, Li 2020 and Li 2022 give $f = 0.54$, 0.15 and 0.81 for the same gyroid, more than a factor of 5, and Nu = 75, 63 and 79 [30, 31]. They also apply only to the cells and walls they were measured or computed for.
2. **Our own flow simulations** match our geometry exactly but cost 1–70 minutes per point. The segment model below evaluates about 440,000 combinations of approach, velocity, passes, cell and segment per variant. At two minutes each, that would take almost two years.
3. **Fitting our own correlation** (power law $\mathrm{Nu} = c\,\mathrm{Re}^m\,\mathrm{Pr}^n$) to the computed points would be fast, but it fixes the form in advance and does not say how reliable it is between the points.

A *surrogate* combines the advantages. It is trained once on a few expensive runs and afterwards delivers values for any input in microseconds, together with a measure of how reliable they are. The concrete question here is whether grading the cells along the flow path pays off. The literature only gives a blanket factor for that (Nu × 1.28 [32]). To check it ourselves, we need to know how Nu and f depend on cell size, and that is exactly what the surrogate learns from our own runs.

How a Gaussian process works. The surrogate is a Gaussian process (kriging). It assumes no formula, only that similar inputs give similar results. A kernel describes this similarity:

$$k(x, x') = \sigma^2 \exp\Big(- \sum_i \frac{(x_i - x'_i)^2}{2\ell_i^2} \Big) + \sigma_n^2\, \delta_{xx'}.$$

The length scale ℓ_i says over what distance in input i the result changes noticeably. A large ℓ_i means input i hardly matters. σ_n^2 absorbs the noise of the runs, such as the time fluctuation of the mixed-mean temperature. All parameters are determined from the data (maximum marginal likelihood). At a new point the process returns a mean, a weighted blend of the known points, and a standard deviation. The latter is small near computed points and grows in gaps, so one can see where the model is guessing. The length scales are a result in their own right: they show which input really matters.

Why multi-level? Fine grids are accurate but expensive (n144: 30–70 minutes, 8 GB of graphics memory); coarse grids are cheap (n48: 1–2 minutes) but systematically too high. A multi-level surrogate (co-kriging [56]) uses both. Many coarse runs describe the *shape* of the dependence on Re, Pr and cell. A few fine runs determine the *shift* caused by the grid. This works well if the grid error varies slowly and smoothly across the input space, and that is checked below.

Design of experiments. For each of gyroid and diamond, 14 points were chosen by Latin hypercube [57]: Re 500–10,000 (logarithmic), Pr 1.2–2.0, cell size $a = 0.75$–1.6 times the standard cell, wall fixed at 0.5 mm. These points ran on the coarse grid (48 cells per d_h). Three control points ran on the medium grid (96), and one run per type on the fine grid (144 cells per d_h, 62.5 million cells). Together with the runs of Section 11.14 that makes 40 runs. The series ran unattended from a job queue that could be paused at any time. A second AI agent watched memory, disk and divergence. No run failed.

Structure. The data are referred to the correlation of Renon [30], so that the Gaussian process only has to learn the deviation:

$$y_{\mathrm{Nu}} = \ln \frac{\mathrm{Nu}}{\mathrm{Nu}_{\mathrm{Renon}}(\mathrm{Re}, \mathrm{Pr})}, \qquad y_f = \ln \frac{f}{f_{\mathrm{Renon}}(\mathrm{Re})}, \qquad x = (\ln \mathrm{Re},\ \mathrm{Pr},\ a).$$

1. **Coarse level:** per type, a Gaussian process with an anisotropic RBF kernel and a noise term over the 15 points on 48 cells. Leave-one-out gives a mean error of 2.8 % (gyroid) and 0.9 % (diamond) for Nu. For f it is 9.5 % and 1.1 %: the gyroid friction scatters noticeably.

2. **Grid correction:** at each control point, $d_{96} = y_{96} - y_{48}(x)$. Over Re this is a straight line in $\ln \mathrm{Re}$ (residual 0.03 in ln). This is co-kriging after Kennedy and O'Hagan [56] with $\rho = 1$ in the logarithmic quantity.

3. **Finest level:** the series 48/96/144 at Re 3000 is *not* in the asymptotic range. We get $d_{96}/d_{144} = 0.58$, while a grid error $\propto n^{-p}$ needs more than $\ln 2/\ln 3 = 0.63$; no positive order fits. There are therefore two upper levels. *n144* is computed: $d_{96}(x)$ is scaled by $d_{144}/d_{96} \approx 1.7$. *Grid* $\to \infty$ is estimated, with Richardson over the last step only and the order from the plate test: $d_\infty = d_{144} + (d_{144} - d_{96})/(1.5^{1.33} - 1)$. Half the distance between the two levels serves as the uncertainty.

		$\mathrm{Nu}/\mathrm{Nu}_{\mathrm{Renon}}$		f/f_{Renon}	
Type	Re	n144	grid $\to \infty$	n144	grid $\to \infty$
Gyroid	1000	1.20	1.03 ± 11 %	0.94 ± 15 %	(uncertain)
Gyroid	3000	1.29	1.03 ± 14 %	1.06 ± 15 %	(uncertain)
Gyroid	8000	1.49	1.14 ± 16 %	1.35 ± 15 %	(uncertain)
Diamond	1000	1.14	0.97 ± 12 %	0.98 ± 6 %	(uncertain)
Diamond	3000	1.23	1.00 ± 14 %	1.01 ± 6 %	(uncertain)
Diamond	8000	1.40	1.08 ± 16 %	1.06 ± 6 %	(uncertain)

Table 17: Surrogate at Pr 1.5 and standard cell. The extrapolated friction cannot be used: it would drop below Renon, because f does not converge monotonically between the grids.

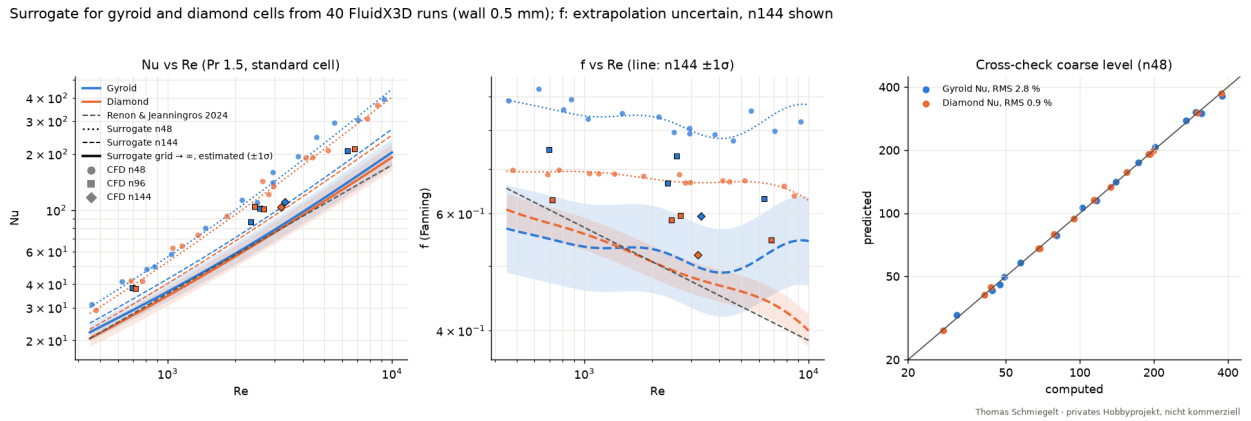

Figure 38: Surrogate for gyroid (blue) and diamond (orange). Left, Nu: dotted the coarse level, dashed n144, solid with band grid $\to \infty$, grey dashed Renon. Centre, f (n144 with band). Right, the cross-check of the coarse level.

What the surrogate says.

- Extrapolated to an infinitely fine grid, Nu lands on the Renon correlation (0.97–1.14). The excess of 1.5 at 96 cells was therefore grid error, not physics. The estimate of 1.2 from only two levels was too high.
- At equal d_h, Nu and f hardly depend on the cell size: for the gyroid, the kernel length scale for a sits at its upper bound. A smaller cell gives a smaller d_h but hardly changes the dimensionless numbers.
- Gyroid and diamond are level for heat transfer. For pressure drop, the gyroid is higher at high Re.

Application: does grading pay off? The preheater of Section 11.13 was set up again as a segment model (`stellarator/vorwaermer_stufung.py`):

- Counterflow, two-fluid TPMS; both channel networks share the cell, so d_h is the same on both sides.
- 8 segments of equal heat duty, each with local temperatures, properties, Pr and LMTD.
- Nu and f from the surrogate (level n144; Nu additionally at grid $\to \infty$ as a band).
- *Uniform:* one cell for all segments. *Graded:* each segment chooses its own cell so that $P_i + \lambda V_i$ is minimal, with λ sampled.
- Also free: approach, condensate velocity, and number of passes on the warm side.

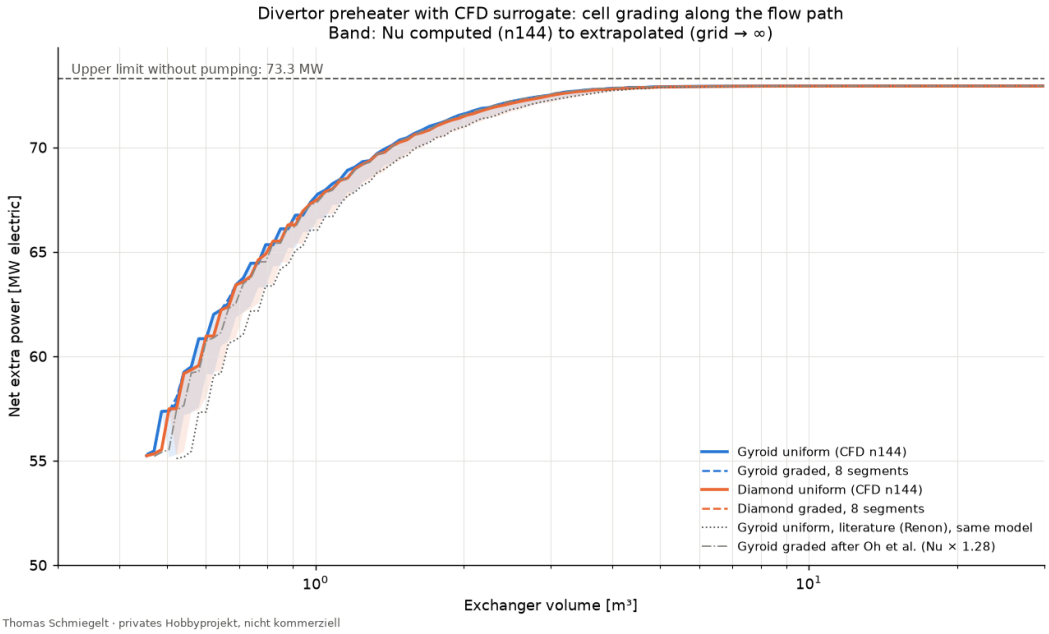

Figure 39: Net extra power over exchanger volume. Graded (dashed) and uniform coincide. For comparison, Renon and Renon with Nu × 1.28 (grading after Oh et al.) in the same model.

Result.

- **Grading brings nothing.** In every segment the smallest allowed cell wins ($a = 0.75\,a_{\mathrm{std}}$, $d_h = 1.33\,\mathrm{mm}$). The graded design is therefore identical to the uniform one. The reason is that pumping power is small here (below 1 MW). Power is lost mainly through the approach at the warm end, and the cheapest way to reduce that is more area per volume, i.e. fine cells everywhere.
- **Volume:** the preheater reaches 95 % of the 73.3 MW upper limit with $1.3\,\mathrm{m}^3$ (n144) to $1.5\,\mathrm{m}^3$ (grid $\to \infty$). At $1\,\mathrm{m}^3$, CFD and Renon give 66.0–67.6 MW, a small difference.
- The literature factor for grading (Nu × 1.28 [32]) here matches almost exactly the gap between Renon and the n144 surrogate. The improvement reported by Oh et al. cannot be reproduced at the cell level. It must come from effects this model does not contain, such as the distribution in the headers.
- **Limits:** the best cell lies at the edge of the surrogate range; finer cells have not been computed. Whether the 0.5 mm wall carries the 3 MPa pressure difference has not been checked. Manufacturing such fine cells is an open question. Grading may pay off where pumping power matters, for example on the steam side of the condenser.

11.16 Collaboration

From the user came: the thinner wall with a larger outlet (f12), pairwise comparison as a tool for decisions and optimisation loops together with the judgements, the point that the walls of the pressure space must withstand the pressure, the point about the small context of the local model, and the decision for an existing turbine. The assistant did the calculations, built the tools, checked the research and named the consequences.

Code and data

Software: ElmerFEM, gmsh, FreeCAD, simsopt, DESC, PyVista, ParaView, Blender, FluidX3D 3.8 [17], PicoGK [25]. `https://github.com/ThomasSchmiegelt/stellarator-design` (MIT licence): pipeline code, findings log, figures.

Citations: only numbers, formulas and short verbatim quotes with source are taken over (German copyright act, §§ 51, 63 UrhG); no figures from other works are reproduced. Source licences are listed in the evidence files under `docs/recherche/`.

References

[1] H.-S. Bosch, R. C. Wolf, T. Andreeva et al.: Technical challenges in the construction of the steady-state stellarator Wendelstein 7-X. *Nucl. Fusion* 53 (2013) 126001. doi:10.1088/0029-5515/53/12/126001

[2] Th. Rummel, K. Rieße, H. Ehmler: Manufacture and test of the non-planar coils for WENDELSTEIN 7-X. *Fusion Eng. Des.* 75–79 (2005) 117–121. doi:10.1016/j.fusengdes.2005.06.118

[3] R. K. Maix, V. Bagnato, M. Fricke et al.: Design, production and QA test results of the NbTi CIC conductors for the W7-X magnet system. *J. Phys.: Conf. Ser.* 43 (2006) 753–758. doi:10.1088/1742-6596/43/1/184

[4] C. Bachmann, M. Siccinio, A. Ciula et al.: Re-design of EU DEMO with a low aspect ratio. *Fusion Eng. Des.* 204 (2024) 114518.

[5] P. F. Gil, V. Brack, T. Schuler, P. Huslage, E. V. Stenson: Manufacturing tolerances of non-planar coils for an optimized tabletop stellarator. arXiv:2507.22516 (2025).

[6] D. Schmeling, M. Russo, B. T. Gebreamlak et al.: Experimental results from early non-planar NI-HTS magnet prototypes for the Columbia Stellarator eXperiment (CSX). *IEEE Trans. Plasma Sci.* (2026), doi:10.1109/TPS.2026.3681798; arXiv:2510.07042.

[7] M. Landreman, B. Medasani, F. Wechsung, A. Giuliani, R. Jorge, C. Zhu: SIMSOPT: A flexible framework for stellarator optimization. *J. Open Source Softw.* 6(65) (2021) 3525. doi:10.21105/joss.03525. W7-X coil data: `simsopt/configs/W7-X.dat`, coil set `w7x_v001`.

[8] D. W. Dudt, E. Kolemen: DESC: A stellarator equilibrium solver. *Phys. Plasmas* 27 (2020) 102513.

[9] C. Geuzaine, J.-F. Remacle: Gmsh: A 3-D finite element mesh generator with built-in pre- and post-processing facilities. *Int. J. Numer. Meth. Eng.* 79 (2009) 1309–1331.

[10] ElmerFEM, version 9.0. CSC – IT Center for Science, Finland. `https://www.elmerfem.org`

[11] FreeCAD, version 1.1. `https://www.freecad.org`

[12] C. B. Sullivan, A. Kaszynski: PyVista: 3D plotting and mesh analysis through a streamlined interface for the Visualization Toolkit (VTK). *J. Open Source Softw.* 4(37) (2019) 1450.

[13] J. Ahrens, B. Geveci, C. Law: ParaView: An end-user tool for large data visualization. In: *Visualization Handbook*, Elsevier (2005).

[14] Blender, version 4.0. Blender Foundation. `https://www.blender.org`

[15] Nous Research: Hermes Agent. `https://github.com/NousResearch/hermes-agent` (MIT licence).

[16] Anthropic: Claude Code. `https://www.anthropic.com/claude-code`

[17] M. Lehmann: FluidX3D, version 3.8. `https://github.com/ProjectPhysX/FluidX3D`. Licence: free for non-commercial use only, source-available, not open source. Altered version of `setup.cpp` and `defines.hpp` in the repository folder `fluidx3d/`.

[18] M. Lehmann: Computational study of microplastic transport at the water-air interface with a memory-optimized lattice Boltzmann method. PhD thesis, University of Bayreuth (2023). doi:10.15495/EPub_UBT_00006977

[19] M. Lehmann: Esoteric Pull and Esoteric Push: Two simple in-place streaming schemes for the lattice Boltzmann method on GPUs. *Computation* 10 (2022) 92. doi:10.3390/computation10060092

[20] M. Lehmann, M. Krause, G. Amati, M. Sega, J. Harting, S. Gekle: Accuracy and performance of the lattice Boltzmann method with 64-bit, 32-bit, and customized 16-bit number formats. *Phys. Rev. E* 106 (2022) 015308.

[21] M. Lehmann: Combined scientific CFD simulation and interactive raytracing with OpenCL. *IWOCL'22: International Workshop on OpenCL* (2022) 3, 1–2.

[22] M. Lehmann, L. M. Oehlschlägel, F. Häusl, A. Held, S. Gekle: Ejection of marine microplastics by raindrops: a computational and experimental study. *Microplastics and Nanoplastics* 1 (2021) 18. doi:10.1186/s43591-021-00018-8

[23] M. Lehmann: High performance free surface LBM on GPUs. Master's thesis, University of Bayreuth (2019). doi:10.15495/EPub_UBT_00005400

[24] M. Lehmann, S. Gekle: Analytic solution to the piecewise linear interface construction problem and its application in curvature calculation for volume-of-fluid simulation codes. *Computation* 10 (2022) 21. doi:10.3390/computation10020021

[25] LEAP 71: PicoGK. https://github.com/LEAP71/PicoGK (Apache 2.0 licence).

[26] sketch-to-prompt: sketches on 3D models as input for AI assistants. https://github.com/ThomasSchmiegelt/sketch-to-prompt (MIT licence).

[27] P. Arena, G. Bongiovì, I. Catanzaro et al.: Design and Integration of the EU-DEMO Water-Cooled Lead Lithium Breeding Blanket. *Energies* 16 (2023) 2069. doi:10.3390/en16042069

[28] A. Bejan: Constructal-theory network of conducting paths for cooling a heat generating volume. *Int. J. Heat Mass Transfer* 40 (1997) 799–816. doi:10.1016/0017-9310(96)00175-5

[29] J. H. You, G. Mazzone, E. Visca et al.: Divertor of the European DEMO: Engineering and technologies for power exhaust. *Fusion Eng. Des.* 175 (2022) 113010. doi:10.1016/j.fusengdes.2022.113010; open copy: doi:10.5445/IR/1000142224

[30] C. Renon, X. Jeanningros: A numerical investigation of heat transfer and pressure drop correlations in Gyroid and Diamond TPMS-based heat exchanger channels. *Int. J. Heat Mass Transfer* (2024) 126599. doi:10.1016/j.ijheatmasstransfer.2024.126599; manuscript: HAL cea-04896410

[31] M. Al-Safadi, F. Ejaz, S. Z. Shuja, S. M. Zubair: Quantitative synthesis and correlations of heat transfer and friction factor in triply periodic minimal surface (TPMS) structures. *Results Eng.* 29 (2026) 109402. doi:10.1016/j.rineng.2026.109402

[32] S.-H. Oh, J. E. Kim, C. H. Jang, J. Kim, C. Y. Park, K. Park: Multifunctional gradations of TPMS architected heat exchanger for enhancements in flow and heat exchange performances. *Sci. Rep.* 15 (2025) 19931. doi:10.1038/s41598-025-04940-2

[33] I. H. Bell, J. Wronski, S. Quoilin, V. Lemort: Pure and pseudo-pure fluid thermophysical property evaluation and the open-source thermophysical property library CoolProp. *Ind. Eng. Chem. Res.* 53 (2014) 2498–2508. doi:10.1021/ie4033999

[34] Wikipedia: EPR (Kernkraftwerk). https://de.wikipedia.org/wiki/EPR_(Kernkraftwerk), accessed 2026-10-31.

[35] Wikipedia: EPR (nuclear reactor). https://en.wikipedia.org/wiki/EPR_(nuclear_reactor), accessed 2026-10-31.

[36] S. V. Patankar, C. H. Liu, E. M. Sparrow: Fully developed flow and heat transfer in ducts having streamwise-periodic variations of cross-sectional area. J. Heat Transfer 99(2) (1977) 180–186. https://doi.org/10.1115/1.3450666

[37] I. Ginzburg: Equilibrium-type and link-type lattice Boltzmann models for generic advection and anisotropic-dispersion equation. Adv. Water Resour. 28(11) (2005) 1171–1195. https://doi.org/10.1016/j.advwatres.2005.03.004

[38] Wikipedia: Stellarator. https://en.wikipedia.org/wiki/Stellarator, abgerufen 2026-10-31.

[39] Wikipedia: Wendelstein 7-X. https://en.wikipedia.org/wiki/Wendelstein_7-X, abgerufen 2026-10-31.

[40] Wikipedia: Deuterium–tritium fusion. https://en.wikipedia.org/wiki/Deuterium%E2%80%93tritium_fusion, abgerufen 2026-10-31.

[41] ITER Organization: What is fusion? https://www.iter.org/sci/whatisfusion, abgerufen 2026-10-31.

[42] Wikipedia: Fusion power. https://en.wikipedia.org/wiki/Fusion_power, abgerufen 2026-10-31.

[43] Wikipedia: Tokamak. https://en.wikipedia.org/wiki/Tokamak, abgerufen 2026-10-31.

[44] Wikipedia: First wall. https://en.wikipedia.org/wiki/First_wall, abgerufen 2026-10-31.

[45] Wikipedia: Breeding blanket. https://en.wikipedia.org/wiki/Breeding_blanket, abgerufen 2026-10-31.

[46] ITER Organization: Blanket. https://www.iter.org/machine/blanket, abgerufen 2026-10-31.

[47] Wikipedia: Divertor. https://en.wikipedia.org/wiki/Divertor, abgerufen 2026-10-31.

[48] ITER Organization: Divertor. https://www.iter.org/machine/divertor, abgerufen 2026-10-31.

[49] Wikipedia: Pressurized water reactor. https://en.wikipedia.org/wiki/Pressurized_water_reactor, abgerufen 2026-10-31.

[50] Wikipedia: Steam generator (nuclear power). https://en.wikipedia.org/wiki/Steam_generator_(nuclear_power), abgerufen 2026-10-31.

[51] Wikipedia: Steam turbine. https://en.wikipedia.org/wiki/Steam_turbine, abgerufen 2026-10-31.

[52] Wikipedia: Surface condenser. https://en.wikipedia.org/wiki/Surface_condenser, abgerufen 2026-10-31.

[53] Wikipedia: Cooling tower. https://en.wikipedia.org/wiki/Cooling_tower, abgerufen 2026-10-31.

[54] Wikipedia: Electric generator. https://en.wikipedia.org/wiki/Electric_generator, abgerufen 2026-10-31.

[55] Wikipedia: Transformer. https://en.wikipedia.org/wiki/Transformer, abgerufen 2026-10-31.

[56] M. C. Kennedy, A. O'Hagan: Predicting the output from a complex computer code when fast approximations are available. Biometrika 87 (2000) 1–13. doi:10.1093/biomet/87.1.1

[57] M. D. McKay, R. J. Beckman, W. J. Conover: A comparison of three methods for selecting values of input variables in the analysis of output from a computer code. Technometrics 21 (1979) 239–245. doi:10.2307/1268522